

Entropy Collapse Has a Direction: Correcting the Visitation Channel in Token-Level Entropy Control for RLVR

Anonymous ACL submission

Abstract

Reinforcement learning with verifiable rewards (RLVR) reliably drives a policy’s entropy toward zero, and a growing family of methods intervenes by re-weighting the policy loss with a per-token estimate of how much each token’s update costs in entropy. We show that every such method estimates only half of the relevant quantity. The gradient of the *trajectory* entropy splits exactly into two terms: a *local* channel, which asks how the conditional distribution at a visited state flattens or sharpens, and a *visitation* channel, which asks how the probability of *reaching* that state changes. Existing token-level estimators capture the first and are structurally blind to the second. We show the blindness is worse than an omission: the local estimator is exactly zero at a uniform branch point, that zero is the correct first-order answer, and the visitation term’s prefactor is simultaneously maximal there — the two channels are *anti-correlated*, so no refinement of the local term can recover the missing signal. We then derive an estimator for the visitation channel. The reward-to-go implied by the decomposition is the remaining conditional entropy H_{togo} , which we forecast with parallel multi-token-prediction heads at the cost of one extra forward pass — though we also measure that forecast to be close to a scalar multiple of a quantity the trainer already computes, and report the head-free control that follows; the sibling-prefix baseline that keeps the estimator unbiased is supplied by a tree-structured rollout in which several samples share a prefix and diverge at scheduled depths. The resulting advantage A_H sums to zero within a sibling set by construction, so the method cannot change the average amount of damping — it can only decide which of two siblings is damped. We give the full derivation, the four approximations it requires and what each discards, and a five-arm design that isolates the contribution of the forecast’s *value* from that of the apparatus carrying it.

1 Introduction

Policy-gradient fine-tuning with verifiable rewards sharpens a language model’s output distribution. Measured as the mean per-token entropy of the policy, this sharpening is fast, monotone, and, past a point, coincident with the end of useful learning. The standard reading is that the collapse is the problem, and the standard remedy is to spend part of the update budget slowing it down: an entropy bonus, a KL anchor, a clip on the tokens that lose the most entropy, or — in the most recent line of work — a per-token re-weighting of the policy loss by a first-order estimate of the entropy each token’s update destroys.

We argue that the accounting behind that remedy is incomplete, and that the incompleteness is structural rather than a matter of estimator quality.

Write $\mathcal{H}(\pi_\theta \mid x)$ for the entropy of the distribution the policy induces over whole responses to a prompt x . The autoregressive factorization turns it into an expectation of a sum of per-token conditional entropies, and the gradient of that expression splits, by the product rule, into exactly two terms. The first holds the visited states fixed and asks how the conditional distribution at each one changes: this is what a per-token entropy estimate measures. The second holds the per-state entropies fixed and asks how the probability of *visiting* each state changes. An update that concentrates mass on one continuation makes the states downstream of that continuation more likely and the states downstream of its abandoned siblings less likely; the diversity those siblings carried leaves the trajectory distribution even though no conditional distribution anywhere became more peaked.

Our first contribution is to show that the local channel is not merely blind to the second term but *anti-correlated* with it. The local estimator factors into the advantage, a saturation factor, and the deviation of the sampled token’s surprisal from the

mean surprisal — and that last factor has mean zero by the definition of entropy. At a branch point where several continuations are close to equiprobable it vanishes identically. The vanishing is not an estimator failure: a uniform distribution maximizes entropy, so the first-order change is genuinely zero and the damage is second order (Section 4.4). Meanwhile the visitation term’s prefactor is $1 - \pi_\theta(y_t | s_t)$, which is maximal precisely at those flat states. The local channel therefore spends its budget where the conditional entropy is fragile — rare tokens sampled from confident distributions — and spends nothing where the trajectory entropy is most at risk.

Our second contribution is an estimator for the missing channel that requires no additional rollouts. Eliminating the non-causal cross terms turns the visitation term into a REINFORCE estimator whose reward-to-go is the sum of the *remaining* conditional entropies, H_{togo} ; subtracting any function of the state leaves it unbiased. We estimate H_{togo} with parallel multi-token-prediction (MTP) heads read off the hidden state the policy already computes, and we obtain the state-only baseline from rollouts that share a prefix. Plain group-relative rollouts almost never share prefixes, so we schedule the sharing: a tree-structured sampler branches a single root at fixed depths, and the resulting sibling sets give the baseline an actual support (Section 9).

The resulting per-token score A_H is a deviation from a sibling mean and therefore sums to zero over each sibling set. This has a consequence we state up front because it determines what evidence can count: the method *cannot* raise the average weight at branch points, only spread the weights of siblings apart. “Total entropy went up” is therefore not a prediction of this method, and an experiment looking for it is testing a different hypothesis (Section 9.3).

We make the following contributions.

- An exact two-channel decomposition of the trajectory-entropy gradient, and an argument that the channel captured by current token-level entropy-control methods is anti-correlated with the channel they miss (Sections 4–4.4).
- A derivation of the missing channel into a per-token advantage A_H with no approximation up to the point of estimating H_{togo} , together

with an explicit ledger of the approximations that estimation requires, the sign of each one’s error, and the mechanism that absorbs it (Sections 5–10).

- A tree-structured rollout that supplies the sibling-prefix baseline, and a bound showing what fraction of positions can carry a non-zero A_H under *any* sampler (Section 9).
- A five-arm experimental design that separates the contribution of the forecast’s value from the contribution of the apparatus that delivers it, by running the full pipeline twice with the forecast’s value destroyed in two different ways (Section 11).

2 Related Work

Entropy collapse in RLVR. Entropy decay under policy-gradient fine-tuning is well documented, and a range of interventions add an explicit entropy term to the objective or constrain the update so that high-entropy positions are preserved. Recent analyses report that naive entropy regularization does not prevent the collapse and that continuing to train under it degrades accuracy (Anonymous, 2025), which is consistent with our reading: the regularizer acts on the local channel, and the local channel is not where trajectory diversity is being lost.

Assumption. Both Hao et al. (2026) and this work rely on one structural assumption, which we restate so that our approximation ledger can refer to it.

Assumption 1 (parameter-independent softmax; Hao et al., 2026). For any state s , each action a carries an independent logit $z_{s,a}(\theta)$, so that a gradient step on the sampled token does not substantially move the logits of the other tokens at that state.

It is false for a network with shared parameters. We inherit it rather than introduce it, but our construction adds a second estimator that rests on it.

Token-level entropy control. The method we build on re-weights the clipped surrogate loss by a per-token first-order estimate of the entropy change that token’s update induces (Hao et al., 2026). That work observes that earlier interventions are “indirect” and that their effect is therefore limited and

can fail. Our decomposition supplies a mechanism for that observation: the estimate is a first-order estimate of the local channel only, and it is exactly zero at the states where the visitation channel is largest.

Which tokens carry the gradient. Wang et al. (2025d) show that restricting policy-gradient updates to a small minority of high-entropy “forking” tokens matches full-gradient training. That result establishes that the identity of the tokens, not the aggregate entropy, is what matters, and it motivates a per-token treatment; it does not, however, distinguish the two channels, since forking tokens are identified by their local entropy. Our analysis predicts that a local criterion will select a systematically different set from the one the visitation channel would select.

Local stochasticity without global diversity. Wu et al. (2025a) report that token-level entropy can remain high while the diversity of final answers falls. This is precisely the signature our decomposition predicts when the visitation channel is unmanaged: per-state entropies are preserved — the local channel is doing its job — while the trajectory distribution concentrates.

Multi-token prediction. We use Medusa-style parallel prediction heads (Cai et al., 2024) not for their intended purpose (speculative decoding) but as an entropy forecaster: we need the entropy of the distribution over a token k steps ahead, not a good sample of that token. Section 6 shows that this choice of estimator has a signed, monotone error which we characterize and correct.

3 Preliminaries

Notation. A prompt is x , a response is $y = y_{1:T}$, and the state before emitting token t is $s_t = (x, y_{<t})$. We write $H(\pi_\theta(\cdot | s_t))$ for the entropy of the conditional distribution at a single position (lowercase, per-token) and $\mathcal{H}(\pi_\theta | x)$ for the entropy of the distribution over whole responses (uppercase, per-trajectory). The surprisal of a token a is $I_a := -\log \pi_\theta(a | s_t)$, so that $H = \mathbb{E}_{a \sim \pi_\theta}[I_a]$.

Group-relative policy optimization. We optimize with GRPO. For each prompt, n responses are sampled and scored, and each response receives the group-standardized advantage $A_i = (r_i - \text{mean}_j r_j) / \text{std}_j r_j$, broadcast to every token of that response. Replacing a learned value function

with the group mean is what makes GRPO cheaper than PPO; it also means that a group whose members all receive the same reward contributes no gradient.

Token-level entropy control. Hao et al. (2026) leave the advantage untouched and instead multiply each token’s contribution to the clipped surrogate by a weight $\alpha_t \in [\lambda_{\min}, 1]$ derived from a first-order estimate of the entropy change that token’s update induces. Their Theorem 1 gives that estimate in closed form: under a parameter-independent softmax (their Assumption 1; our 2), the change in the conditional entropy at a *fixed* state s after one GRPO step is, to first order,

$$\Omega(s) = -\frac{\eta}{L} \mathbb{E}_a \left[\frac{\mathbb{I}_{\text{clip}}(s, a) A(s, a)}{\pi_{\text{old}}(a | s)} \cdot \pi_\theta(a | s) (1 - \pi_\theta(a | s)) \cdot (\log \pi_\theta(a | s) + H) \right], \quad (1)$$

where η is the learning rate, L the total decoded length in the update, and $\mathbb{I}_{\text{clip}}$ the PPO clipping indicator. The implementation evaluates the bracket at the sampled token rather than in expectation. The policy loss becomes $\sum_u \alpha_u \cdot \text{PPO-clip}(A_u, r_u)$, with

$$\alpha_u = \exp\left(-\alpha \frac{|\Omega_u|}{\max_v |\Omega_v|}\right), \quad \lambda_{\min} = e^{-\alpha}, \quad (2)$$

a batch-normalized exponential map into $[\lambda_{\min}, 1]$, so the weight can only attenuate. The factor $1/\pi_{\text{old}}$ in Eq. (1) will matter later: it makes Ω heavy-tailed, which is why the two channels must be normalized before they are combined (Section 8).

4 Two Channels

4.1 Step 1: the chain rule

The object we would like to protect is the entropy of the distribution the policy induces over *whole responses*,

$$\mathcal{H}(\pi_\theta | x) := -\sum_y \pi_\theta(y | x) \log \pi_\theta(y | x), \quad (3)$$

a sum over an exponentially large set. The autoregressive factorization $\pi_\theta(y | x) = \prod_t \pi_\theta(y_t | s_t)$ turns the logarithm into a sum and lets us exchange

the order of summation:

$$\begin{aligned}\mathcal{H} &= - \sum_y \pi_\theta(y | x) \sum_{t=1}^T \log \pi_\theta(y_t | s_t) \\ &= \sum_{t=1}^T \mathbb{E}_{y \sim \pi_\theta} [-\log \pi_\theta(y_t | s_t)].\end{aligned}\quad (4)$$

Each summand is still an expectation over whole responses, but the integrand depends on y only through $y_{\leq t}$. Conditioning on the prefix and applying the tower rule replaces the inner expectation with a per-position entropy,

$$\begin{aligned}\mathbb{E}_y [-\log \pi_\theta(y_t | s_t)] \\ = \mathbb{E}_{y_{<t}} \left[\underbrace{\mathbb{E}_{y_t \sim \pi_\theta(\cdot | s_t)} [-\log \pi_\theta(y_t | s_t)]}_{= H(\pi_\theta(\cdot | s_t))} \right],\end{aligned}$$

because the inner expectation is, by definition, the entropy of the conditional distribution at s_t . Collecting the terms back under a single expectation gives

$$\mathcal{H}(\pi_\theta | x) = \mathbb{E}_{y \sim \pi_\theta} \left[\sum_{t=1}^T H(\pi_\theta(\cdot | s_t)) \right]. \quad (5)$$

Equation (5) is an identity, not an approximation. Its left-hand side is what we want to protect; its right-hand side is what instrumentation can measure, since $H(\pi_\theta(\cdot | s_t))$ is read directly off the logits at each position.

The entire argument of this paper follows from one feature of Eq. (5) that is easy to pass over: *the measure $\mathbb{E}_{y \sim \pi_\theta}$ depends on θ as well*. The right-hand side is not a sum of θ -dependent terms over a fixed set of positions — it is a θ -dependent average of them over a θ -dependent set of trajectories.

4.2 Step 2: the product rule

Write the per-trajectory total

$$f_\theta(y) := \sum_{t=1}^T H(\pi_\theta(\cdot | s_t(y))), \quad (6)$$

so that Eq. (5) reads $\mathcal{H} = \sum_y \pi_\theta(y | x) f_\theta(y)$. Both factors of every summand carry θ , so differentiating requires the product rule:

$$\begin{aligned}\nabla_\theta \mathcal{H} &= \sum_y [\nabla_\theta \pi_\theta(y | x)] f_\theta(y) \\ &\quad + \sum_y \pi_\theta(y | x) \nabla_\theta f_\theta(y).\end{aligned}\quad (7)$$

The second sum is already an expectation. The first is not, because $\nabla_\theta \pi_\theta$ is not a probability distribution; the score-function identity $\nabla_\theta \pi_\theta = \pi_\theta \nabla_\theta \log \pi_\theta$ restores it to one. Substituting and naming the two terms,

$$\begin{aligned}\nabla_\theta \mathcal{H} &= \underbrace{\mathbb{E}[\nabla_\theta f_\theta(y)]}_{\text{(L) local}} \\ &\quad + \underbrace{\mathbb{E}[f_\theta(y) \nabla_\theta \log \pi_\theta(y | x)]}_{\text{(V) visitation}}.\end{aligned}\quad (8)$$

Equations (3)–(8) contain no approximation of any kind. The split into two channels is forced by the product rule, not chosen.

4.3 Reading the two terms

(L), the local channel, holds the visited states fixed and asks how the conditional distribution at each one changes. Equation (1) is a first-order estimate of exactly this.

(V), the visitation channel, holds the per-state entropies fixed and asks how the probability of reaching each state changes. An update that concentrates mass on one continuation makes the subtree below it more likely and the sub-trees below its siblings less likely. Nothing in Eq. (1) can see this: Ω reads the conditional distribution at a position and never reads the probability of arriving there.

The two terms can have opposite signs. A method that drives (L) to zero has not thereby bounded $\nabla_\theta \mathcal{H}$.

4.4 Why the local channel cannot be repaired

A natural objection is that the two channels might coincide in practice: branch points are positions where the policy is uncertain, so they are high-entropy positions, so Ω should be large there and the local term should protect them automatically. The premise is correct — in warm-up rollouts on our model, branch diversity correlates with local entropy at Spearman $\rho = 0.65$, and with the entropy-to-go at $\rho = 0.82$. The inference from it fails, in three steps.

Step 1: Ω reads a deviation, not a level.

Collecting the constants of Eq. (1) into $c := \eta \mathbb{I}_{\text{clip}} A / (L \pi_{\text{old}})$ and rewriting in terms of surprisal gives, for the contribution of a single action a ,

$$\Omega \propto c \pi_a (1 - \pi_a) (H - I_a), \quad (9)$$

and $H = \mathbb{E}_{a \sim \pi} [I_a]$ by definition. The factor $H - I_a$ is therefore a mean-zero fluctuation: what sets the

magnitude of $|\Omega|$ is the *spread* of surprisal, not the *height* of the entropy. Nothing in c depends on the shape of $\pi(\cdot | s)$, so the argument below is unaffected by the clipping indicator or the importance weight.

Step 2: at a uniform branch point $\Omega \equiv 0$. If the distribution is uniform over k candidates then $I_a = \log k = H$ for every a , so Ω vanishes identically. The positions that are most branch-like — several continuations at comparable probability — are exactly where the local term contributes nothing.

Step 3: and that zero is the correct answer. This is the step that closes the argument, and it follows from Hao et al.’s own derivation. Differentiating the conditional entropy with respect to a single logit gives (Hao et al., 2026, Eq. 20)

$$\frac{\partial H(\pi_\theta | s)}{\partial z_{s,a}} = -\pi_\theta(a | s)(\log \pi_\theta(a | s) + H(\pi_\theta | s)), \quad (10)$$

and at a uniform distribution $\log \pi_\theta(a | s) = -\log k = -H$ makes the bracket vanish for every a . So $\partial H / \partial z_a = 0$ identically: Ω is not wrong, it is right. Perturbing one logit of a uniform distribution by δ and writing $S := e^\delta + k - 1$ gives the closed form

$$H(\delta) = \log S - \frac{\delta e^\delta}{S}, \quad H''(0) = -\frac{k-1}{k^2}, \quad (11)$$

with $H'(\delta)$ of sign $-\text{sign}(\delta)$: any push in any direction lowers the entropy, but the loss is entirely second order. For a 10-way uniform distribution and $\delta = 0.1$ the actual entropy loss is 4.7×10^{-4} , roughly two orders of magnitude below that of a peaked distribution under the same perturbation. As far as local entropy is concerned, uniform branch points are genuinely safe, and a local estimator that allocates them no protection is allocating correctly for its own objective. No refinement of the local term recovers the missing signal, because the missing signal is not local.

The two channels are anti-correlated. The visitation term’s prefactor, derived in Section 7, is $1 - \pi_a$, which approaches 1 as the distribution flattens — and the probability that a group of n samples actually diverges at a position approaches 1 in the same limit. Table 1 evaluates both channels on a family of distributions that interpolates

distribution	H	$P(\text{split})$	$\mathbb{E} \Omega / A $	$\mathbb{E}[1-\pi_a]$
$q=.99$	0.10	0.08	0.092	0.020
$q=.9, \text{ tail}$	1.25	0.57	1.129	0.190
$q=.6, m=3$	1.11	0.98	0.457	0.587
unif. $k=10$	2.30	1.00	0.000	0.900
unif. $k=100$	4.61	1.00	0.000	0.990

Table 1: The local channel and the visitation channel are ordered oppositely. Distribution family: top token has mass q , the remaining $1 - q$ is spread uniformly over m tokens. $P(\text{split})$ is the probability that $n=8$ samples do not all agree. $\mathbb{E}|\Omega|$ is taken over the token actually sampled at that position.

between confident and flat. $|\Omega|$ is dominated by positions where the model is fairly confident but a tail token was sampled; at those positions siblings diverge only about half the time. Where divergence is certain, the local term is identically zero and the visitation prefactor is maximal.

The statement we can therefore make is stronger than “the local term is incomplete”: *the local term is anti-correlated with the term it omits*, and the omission cannot be fixed within the local channel.

5 Estimating the Visitation Channel

5.1 Causal elimination gives a reward-to-go

The visitation term of Eq. (8) is not yet in a usable form: the factor $f_\theta(y)$ multiplying the score is a property of the *whole* trajectory, so a naive estimator would credit every token with entropy that appeared before it. Causality removes those terms exactly.

Expanding both factors — f_θ by Eq. (6) and the score by $\nabla_\theta \log \pi_\theta(y | x) = \sum_u \nabla_\theta \log \pi_\theta(y_u | s_u)$ — turns (V) into a double sum:

$$(V) = \sum_t \sum_u \mathbb{E}[H(\pi_\theta(\cdot | s_t)) \cdot \nabla_\theta \log \pi_\theta(y_u | s_u)]. \quad (12)$$

Consider a term with $u \geq t$. The entropy $H(\pi_\theta(\cdot | s_t))$ depends on y only through $y_{<t}$, and $y_{<t} \subseteq y_{<u}$, so it is measurable with respect to $y_{<u}$ and may be pulled out of a conditional expectation given $y_{<u}$. What remains is the expected score at a single position, which vanishes for any distribution:

$$\begin{aligned} & \mathbb{E}[\nabla_\theta \log \pi_\theta(y_u | s_u) | y_{<u}] \\ &= \sum_a \pi_\theta(a | s_u) \frac{\nabla_\theta \pi_\theta(a | s_u)}{\pi_\theta(a | s_u)} \\ &= \nabla_\theta \sum_a \pi_\theta(a | s_u) = \nabla_\theta 1 = 0. \end{aligned} \quad (13)$$

By the tower rule every (t, u) term with $u \geq t$ is therefore *exactly* zero — not small, zero. This is the same identity that licenses a baseline in policy gradients, used here to enforce causality rather than to reduce variance.

Only $u < t$ survives. Swapping the order of summation collects, for each token u , the entropies at all later positions:

$$(V) = \mathbb{E} \left[\sum_u \nabla_{\theta} \log \pi_{\theta}(y_u \mid s_u) H_{\text{togo}}(s_u \oplus y_u) \right], \quad (14)$$

with $H_{\text{togo}}(s_u \oplus y_u) := \sum_{t>u} H(\pi_{\theta}(\cdot \mid s_t))$.

We stress that H_{togo} was not posited. Equation (14) is a REINFORCE estimator whose reward is “the conditional entropy remaining after this token”, and H_{togo} is that reward-to-go.

5.2 The sibling-prefix baseline

Any function $b(s_u)$ of the state alone leaves Eq. (14) unbiased, by the same argument. Subtracting one gives the per-token advantage

$$A_H(s_u, y_u) := H_{\text{togo}}(s_u \oplus y_u) - \bar{H}_{\text{togo}}(s_u), \quad (15)$$

and up to this point *no approximation has been made*.

The baseline we use is the mean of H_{togo} over the rollouts in the same prompt group whose prefixes $y_{<u}$ agree exactly. Those rollouts share the state s_u by definition, so their mean is a function of s_u alone and the baseline is admissible. The condition is agreement on $y_{<u}$ and not on $y_{\leq u}$: the rollouts must diverge *at* u for the difference to be a branch score.

Two properties follow immediately and matter later. First, the implementation includes a rollout in its own sibling set, so b depends weakly on y_u and the estimator is not strictly leave-one-out; the resulting bias is $O(1/m)$ in the number m of surviving siblings, the same property GRPO’s group-mean advantage has. Second, when $m = 1$ the baseline equals the rollout’s own value and $A_H \equiv 0$, so a position with no siblings is silent rather than noisy.

6 Four Approximations

Equation (15) is exact but not computable: H_{togo} sums conditional entropies at states the update has not visited. Three approximations make it computable and a fourth bounds its influence.

6.1 A1: horizon truncation and discounting

We truncate the sum to κ steps and discount:

$$H_{\text{togo}}(s_u \oplus y_u) \approx \sum_{k=1}^{\kappa} \gamma_H^k H(\pi(\cdot \mid s_{u+k})). \quad (16)$$

The discount serves two purposes: distant terms carry more forecast error (A2) and therefore deserve less trust, and the sum stays bounded independently of κ .

What this discards — the tail $\sum_{k>\kappa}$ and a fraction of the retained terms — is largely harmless for a reason specific to our use. Equation (15) consumes only the *difference* between siblings. Any bias common to rollouts sharing s_u cancels in that difference, and truncation bias is common to first order.

6.2 A2: forecasting the unrealized future

Even Eq. (16) requires the future states s_{u+k} . At update time they do not exist: rollouts came from π_{old} , each state has exactly one sampled continuation, and re-rolling is prohibitive. We replace the realized future with a parallel forecast from the current hidden state h_u :

$$H(\pi(\cdot \mid s_{u+k})) \longrightarrow H(p_{\text{MTP}}^{(k)}(\cdot \mid h_u)). \quad (17)$$

Head k is a two-layer residual MLP applied to h_u , followed by the policy’s own unembedding; the heads own no vocabulary parameters. Their last layer is zero-initialized, so before any training every head reproduces the policy’s own next-token distribution and the forecast starts from a sane value. In the RL path the logits are materialized in chunks and reduced to an entropy immediately, so the $[\kappa, B, T, V]$ tensor never exists.

The error of this substitution has a known sign. Conditioning on the intervening tokens is exactly what the forecast omits. Fixing S_u , the k -th term of Eq. (16) has expectation $H(Y_{u+k} \mid Y_{u+1:u+k-1}, S_u)$, whereas the MTP head models the *marginal* $p(Y_{u+k} \mid S_u)$. Since conditioning reduces entropy,

$$\begin{aligned} H(Y_{u+k} \mid S_u) - \mathbb{E}[H(\pi(\cdot \mid S_{u+k-1}))] \\ = I(Y_{u+k}; Y_{u+1:u+k-1} \mid S_u) \\ \geq 0. \end{aligned} \quad (18)$$

The forecast *always over-estimates*, and the excess is precisely the mutual information with the intervening tokens, which grows monotonically with k .

Equation (18) is not merely a bound: it is the reason the marginal is the quantity we want. It is the fan-out of the continuation — how many ways the next k tokens can unfold — rather than the residual uncertainty along the one path that happened to be sampled.

6.3 A3: per-head affine calibration

Equation (18) predicts a per-head, monotone inflation, so we absorb it with a per-head affine correction fitted by least squares against oracle entropies on warm-up rollouts:

$$\hat{H}_k = a_k H(p_{\text{MTP}}^{(k)}(\cdot | h_u)) + b_k. \quad (19)$$

The fitted coefficients confirm the prediction. On our model the first head has $a_1 = 1.03$ — a one-step forecast has no intervening tokens, so the mutual information in Eq. (18) is zero and the approximation is exact — while a_k drops to 0.10–0.19 for $k \geq 2$. Calibration switches the unreliable heads off on its own.

Combining A1–A3 gives the forecast that is actually computed:

$$\hat{H}_{\text{togo}}(s_u \oplus y_u) = \max\left(0, \sum_{k=1}^{\kappa} \gamma_H^k (a_k H_k(h_u) + b_k)\right), \quad (20)$$

the clamp encoding the fact that an entropy cannot be negative. With $\kappa = 2$ and $\gamma_H = 0.7$ this instantiates to $\hat{H}_{\text{togo}} = 0.72 H_1 + 0.06 H_2 + 0.11$: the second head contributes about 8%, so the effective horizon is close to one step. We regard this as a consequence of Eq. (18) rather than a tuning artifact, and note it as a limitation — a forecaster whose error did not grow with k would leave a longer horizon usable.

Alignment. One detail decides whether the signal is a branch score at all. The hidden state at position u is conditioned on s_u and predicts y_u ; it does not know y_u . The quantity we need, $H_{\text{togo}}(s_u \oplus y_u)$, must therefore be read from position $u + 1$ and shifted back by one. Omitting the shift measures the value of the state *before* the branch rather than the value of the branch taken.

7 From Gradient to Token Weights

Equation (15) is a term to add to a gradient, but methods in this family do not add gradients — they re-weight a loss. What is needed is a first-order prediction of how much one update, acting through token u , changes the trajectory entropy.

How far one step moves a logit. We reuse the logit-space update of Hao et al. (2026, Eq. 23) rather than rederiving it. With $w_u = \text{clip}(r_u, 1 - \epsilon_{\text{lo}}, 1 + \epsilon_{\text{hi}}) A_u$ the coefficient the policy loss actually applies, Assumption 2 makes a single step move only the sampled token’s logit, $\Delta z_{u,a} = \eta w_u \delta_{a,y_u} + O(\eta^2)$. The softmax Jacobian $\partial \log \pi_a / \partial z_a = 1 - \pi_a$ then gives

$$\widehat{\Delta \log \pi_u} = \eta w_u (1 - \pi_u), \quad (21)$$

which vanishes as $\pi_u \rightarrow 1$: a token that already holds the mass cannot be given more. This is the prefactor whose anti-correlation with $|\Omega|$ Table 1 records.

The visitation-channel prediction. Because $\nabla_{\theta} \log \pi_u \cdot \Delta \theta = \Delta \log \pi_u$, substituting into Eq. (15) gives the per-token quantity

$$\text{visit}_u := \widehat{\Delta \log \pi_u} \cdot \text{clip}(A_H(s_u, y_u), -c, +c). \quad (22)$$

A4 is the clip, and it is applied to A_H *before* the product rather than after, so that one token’s influence is bounded by $|\eta w_u| \cdot c$. Clipping the product instead would let forecast error be amplified without bound at tokens with large w_u .

Sign quadrants. Equation (22) defines the collapse channel precisely. When $\widehat{\Delta \log \pi_u} > 0$ and $A_H > 0$, mass moves toward a branch richer than its siblings and trajectory entropy rises. When $\widehat{\Delta \log \pi_u} > 0$ and $A_H < 0$, mass moves into a dead end and the diversity carried by the abandoned siblings is destroyed — this is the collapse path that (L) is structurally unable to see. Negative $\widehat{\Delta \log \pi_u}$ reverses both readings.

8 Combining the Channels

The visitation term could be added to Ω before the weight map, but measured on real rollouts that map destroys the signal: Ω carries a factor $1/\pi_{\text{old}}$, so $\max |\Omega|/\text{median}|\Omega|$ reaches 3×10^4 and the band is spent on one or two outliers while the bulk lands on a point. We therefore apply the correction *after* the map, as a bounded adjustment of the weights the base method has already produced.

Let $\mathcal{S} = \{u : \text{visit}_u \neq 0\}$ be the support — by Eq. (24), the branch points — and let $b = \lambda_{\text{max}} - \lambda_{\text{min}}$ be the width of the weight band. For $u \in \mathcal{S}$,

$$\alpha_u^{\text{f}} = \text{clamp}\left(\alpha_u + \lambda b \tanh(\text{visit}_u / \rho), \lambda_{\text{min}}, \lambda_{\text{max}}\right), \quad (23)$$

with $\rho = \text{rms}\{\text{visit}_v : v \in \mathcal{S}\}$; $\alpha_u^f = \alpha_u$ elsewhere.

Three properties matter. The tanh is odd, monotone and bounded, so it preserves the sign of the visitation term, caps one token’s movement at λb , and does not depend on the absolute scale of visit. The RMS is taken over $\mathcal{S}$ alone rather than over the whole response: normalizing by the whole response would shrink the correction in proportion to how sparse branch points are, which is a property of the sampler and not of the signal. And at $\lambda = 0$ the correction vanishes identically, so the pipeline is bitwise identical to the base method — a regression test enforces this.

Because the map runs first and the correction second, the step size η of Eq. (21) enters only through visit_u/ρ , where it cancels. We fix $\eta = 1$ and sweep λ alone.

The support is part of the treatment. $\mathcal{S}$ is defined by $\text{visit}_u \neq 0$ rather than by a threshold, deliberately: it selects both where the correction lands and, through ρ , how large every correction is, so a tighter test would change what is trained rather than what is logged. One consequence is that $\mathcal{S}$ admits float32 round-off — A_H is a deviation from a mean over token-identical siblings, and $\sum_j g_j/m$ does not reproduce g_i bitwise, leaving $|A_H| \sim 6 \times 10^{-8}$ where it is interpretively zero. Those residuals are $\sim 10^{-6}$ relative to a real branch score and cannot move a weight, but they are counted in the support statistics of Section 12.6, which should be read as an upper bound on the branch-point fraction.

9 Tree-Structured Rollouts

9.1 Why plain rollouts cannot support A_H

$\hat{H}_{\text{togo}}$ is a deterministic function of the causal prefix, so two rollouts sharing $y_{\leq u}$ produce identical forecasts. Combined with Eq. (15) this gives a sharp characterization:

$$A_H[i, u] \neq 0 \iff \begin{array}{l} \text{siblings share } y_{<u} \\ \text{but diverge at } y_u, \end{array} \quad (24)$$

that is, u is a branch point of the group. Under independent sampling from a sharpened policy, group members diverge at their first token and never share a prefix again, so $m = 1$ almost everywhere and $A_H \equiv 0$: the visitation channel is not merely weak, it is identically absent.

We therefore schedule the sharing. A tree sampler expands a single root and cuts it at fixed depths $d_1 < d_2 < d_3$ with branching factor 2, yielding $n = 8$ leaves whose sibling structure is known in advance. The cuts create the only positions at which Eq. (24) can fire.

9.2 A ceiling on the support

It is tempting to treat the fraction of positions with $A_H \neq 0$ as a quantity to be maximized. It is not, and the reason is combinatorial rather than empirical. Two quantities behave very differently under the tree:

- the fraction of positions where *some* other rollout shares $y_{<u}$ is moved by the cuts, and equals d_L/T by construction;
- the fraction where $A_H \neq 0$ is *not*, because by Eq. (24) it counts branch points, and n sequences form a trie with at most $n - 1$ internal branching nodes.

Since the one-position alignment shift widens each branch node into two columns, the fraction of positions carrying a non-zero A_H is bounded by $2(n - 1)/T$ for *any* sampler whatsoever. With $n = 8$ and responses of $\sim 10^3$ tokens this ceiling is on the order of a percent. The tree’s role is to make the branch points exist and to place them at controlled depths, not to make them numerous.

9.3 The correction cannot move the mean

A_H is a deviation from a sibling mean, so on each sibling set

$$\sum_{j \in \text{sib}(s_u)} A_H(s_u, y_u^{(j)}) = 0. \quad (25)$$

The correction in Eq. (23) therefore cannot change the average weight assigned to branch tokens. It can only spread siblings apart. This is not a defect to be engineered away: an unbiased estimator requires a baseline, and a baseline forces the mean to zero.

The consequence for evaluation is direct, and we state it before presenting any experiment. “Mean entropy at branch points increased” is *not* a prediction of this method and cannot serve as evidence that it works; the corresponding prediction is that the *variance across siblings* widens. A method designed to raise the mean is a different and weaker hypothesis — it does not need a forecast at all, only a support — and we include it below as a control arm.

10 The Approximation Ledger

Table 2 collects every departure from Eq. (8).

The strongest assumptions are A2 and B2. A2 has a known sign and a known monotonicity (Eq. 18), which is what makes the affine correction in A3 an appropriate remedy rather than a curve fit. B2 — that updating one token does not change the conditional distribution at other positions — is false for a network with shared parameters and has no corrective mechanism. It is not a risk this work introduces, since the base method already relies on it, but adding a second channel adds a second copy of its error.

11 Experimental Setup

11.1 Model, data, and optimization

Within a backbone, all runs fine-tune the same base model on the same data with the same optimizer settings, and differ only in the intervention. Our primary backbone is a 1.5B-parameter mathematics-specialized model; the full five-arm decomposition and every ablation are run there. Three further backbones (Table 3) carry the headline three-way contrast only, which is what a second model can answer: not *how* the effect decomposes—that is a question about the method, and it is settled once—but whether its *sign* survives a change of scale and of pre-training family. All runs train on the same 17k-problem competition-mathematics set. With a training batch of 512 prompts, 110 optimization steps correspond to about 3.2 epochs.

11.2 Validation set, and why it is not the same everywhere

The mathematics-specialized backbones validate on AIME 2024 (30 problems, 32 samples each). AIME 2024 is small: with a per-problem accuracy spread of .1542 measured across our runs, the standard error of its mean is $.1542/\sqrt{30} \approx .0282$. For a backbone that scores near zero on it, every arm lies inside one standard error, so the benchmark cannot separate them.

The sharper problem is not the reported number but the *selection*: we choose each run’s checkpoint by its best validation accuracy, so a benchmark that cannot distinguish arms also cannot distinguish checkpoints, and the selection becomes a draw from noise. The non-mathematics backbones therefore validate and select on MATH500 (500 problems, one sample each), whose standard error is roughly four times smaller.

This buys a measurable comparison at the cost of a protocol that is not identical across backbones. We state it here rather than in a footnote because it constrains what may be read off the results: every contrast we report is *within* a backbone, and no claim compares absolute accuracy between rows of Table 3.

Base or instruction-tuned. The same measurement forces a second protocol difference. The Llama backbone is the instruction-tuned checkpoint; the other three are base models. We began with the base one and scored it under this protocol before training it: it reached 0.020 on MATH500 and 0.024 on grade-school arithmetic. The cause is mechanical rather than mathematical—that checkpoint ships a chat template, which our data pipeline applies, so a model that was never instruction-tuned receives the instruction format and continues the prompt instead of answering it.

The instruction-tuned checkpoint of the same family and the same parameter count, scored under the identical protocol, reaches 0.340 and 0.803 on those two benchmarks. Holding family and scale fixed and changing only the checkpoint moves MATH500 by a factor of seventeen, which is what identifies the failure as one of format rather than of capacity: this backbone can do the arithmetic, and the base checkpoint was not being asked for it in a form it could answer.

What this costs is not a weak row but an absent one. A group of $n = 8$ responses receives zero advantage when its members score alike, so at that solve rate $(1 - 0.020)^8 \approx 0.85$ of groups carry no gradient at all, on a training set harder than the benchmark just quoted. Neither our method nor the baselines can learn from the remainder, and the row would compare two untrained models. We therefore gate a backbone on the fraction of non-degenerate groups on the *training* set, measured against a backbone known to train, rather than on a validation score: a validation set says whether a backbone can be measured, not whether it can learn. On that gate the instruction-tuned checkpoint leaves — of groups informative against — for the 1.5B mathematics backbone, a ratio of —.

Optimization uses GRPO with $n = 8$ responses per prompt, a constant learning rate of $1e-6$ with no warm-up (so that the schedule is independent of the declared total step count), a maximum response length of 3072 tokens, no entropy bonus, and no KL loss. The clipped surrogate is asym-

	approximation	what it discards	sign / bound	mitigation
—	Eqs. (5)–(15)	nothing — identities	—	—
A1	horizon κ , discount γ_H	the tail $\sum_{k>\kappa}$, and a fraction of the retained terms	under-estimate	the common component cancels in the sibling difference
A2	realized future $\rightarrow$ MTP marginal	mutual information with intervening tokens, Eq. (18)	over-estimate, monotone in k	A3
A3	per-head affine calibration	non-affine residual of A2	least-squares residual	oracle-forecast control arm
A4	clip A_H to $\pm c$	magnitude of extreme branch scores	bounded by c	saturation rate is logged
B1	first order in η	$O(\eta^2)$	—	small learning rate
B2	only the sampled logit moves	cross-position effects of weight sharing	—	<i>inherited from the base method</i>
B3	rollouts come from π_{old}	mismatch with the π_θ expectation	—	importance ratio restores r
B4	self included in sibling set	$O(1/m)$ relative to leave-one-out	—	same property as the GRPO advantage

Table 2: Every approximation between the exact decomposition and the implemented weight. A1–A4 are introduced by this work; B1–B4 are shared with the token-level entropy-control method we extend.

Backbone	Family	Math-spec.	Validation
1.5B	Qwen2.5-Math	yes	AIME 2024
7B	Qwen2.5-Math	yes	AIME 2024
3B*	Llama 3.2	no	MATH500
7B	Mistral	no	MATH500

Table 3: Backbones. The validation set differs by backbone, deliberately; Section 11.2 says why, and the analysis never compares absolute accuracy across rows. *This row uses the instruction-tuned checkpoint; the other three are base models, and Section 11.2 says why this one cannot be.

metric, $\varepsilon_{\text{low}} = 0.2$ and $\varepsilon_{\text{high}} = 0.28$, with the dual-clip lower bound at $c = 10$; we state this because “GRPO” with a symmetric $\varepsilon = 0.2$ and a KL term is a different baseline, and because the follow-up arm of Section 12.10 is defined by moving these two numbers. Every arm, the baseline included, uses the same three values: the only configuration difference between the GRPO arm and STEER-F is the loss mode and the method hyperparameters listed next. Rollouts are sampled at temperature 1.0 with no nucleus truncation; validation uses temperature 1.0 and top- $p = 0.7$. Every arm uses the same data seed and the same initial checkpoint, which we verify by requiring the step-0 validation scores to agree.

11.3 Method hyperparameters

The forecaster uses $\kappa = 2$ heads with discount $\gamma_H = 0.7$ and the per-head calibration of Eq. (19). The combination uses $\lambda = 0.25$, RMS normalization, clip $c = 1.0$, and the sibling-prefix baseline. The weight map is the base method’s exponential

map of Eq. (2), taken un-rescaled, at $\lambda_{\min} = 0.7$, so weights land in $[0.7, 1]$ and can only attenuate. The released configuration of the base method sets $\lambda_{\min} = 0.8$; every other value in that launcher matches it. We state plainly how ours was chosen, because it bears on which contrasts can be read: 0.7 was selected on *our* method, not on the base method, so the baseline runs at a damping strength picked for the treatment and 1.5 times the maximum attenuation its own release uses. Section 12.10 reports the base method at $\lambda_{\min} = 0.8$ for this reason.¹ The tree sampler uses a single root, cuts at depths $\{64, 192, 384\}$, and branching factor 2 at each cut, giving $n = 8$ leaves.

11.4 Five arms

The design isolates one question: does the *value* of the forecast matter, or only the apparatus that carries it? A comparison between the base method and the full method changes two things at once (λ and the sampler), so it cannot answer this. We therefore add two arms that keep the entire apparatus — the tree, the forecaster, the support, the weight mapping — and destroy only the information in A_H .

UNIFORM applies the same attenuation to every position in the support, discarding A_H ’s value but keeping its support — this is the implementation of the “raise the mean” hypothesis that Section 9.3 distinguishes from ours. PERMUTED keeps the multiset of A_H values within each sibling set and shuf-

¹Our implementation exposes an optional reshaping of $\tilde{\Omega}$ before this map (winsorizing or rank-transforming it); every run reported here uses the identity setting, so the base method’s map sees exactly what it would have seen.

arm	λ	rollout	signal in the weight
GRPO	—	plain	none
STEER	0	plain	local Ω only
uniform	.25	tree	support only
permuted	.25	tree	A_H shuffled
STEER-F	.25	tree	A_H as derived

Table 4: The five arms. UNIFORM and PERMUTED run the full pipeline with the forecast’s value destroyed in two different ways, so any gain they show is attributable to the apparatus rather than to the visitation signal.

fles the assignment, so the support, the magnitudes, and Eq. (25) all survive while the pairing between a branch and its score is destroyed. Together they bracket the claim: if the visitation signal is doing the work, both controls should behave like the no-forecast arms.

11.5 Metrics

We report acc@32 (mean accuracy over 32 samples) and maj@32 (accuracy of the majority vote over the same 32 samples), and their difference

$$\text{uplift} := \text{maj@32} - \text{acc@32},$$

which measures how much better a set of 32 samples is than its average member — that is, whether correct answers concentrate while incorrect ones scatter. We use uplift rather than pass@32 because the latter does not separate the arms.

Because single validation points are noisy on a 30-problem set, the primary statistic is the mean over the converged window, steps 40–110 (8 validation points per arm). We quote differences in units of the across-problem standard error, $\text{SE} = \text{std@32}/\sqrt{30} \approx 0.027$, which is the correct scale for generalization to new problems. We separately note that the pure re-sampling noise of the protocol — estimated from two validations of identical weights — is about 0.004 in accuracy, an order of magnitude smaller; the two scales answer different questions and we are explicit about which is in use.

One instrumentation caveat propagates into every entropy number we report. The logged `actor/entropy` is aggregated with a token-mean, so it is $\mathcal{H}$ divided by the mean response length, not $\mathcal{H}$. Since response lengths differ by under 5% across arms, we report the raw quantity and, where the distinction could matter, the length-corrected product as well.

step	GRPO	STEER	UNIF	PERM	STEER-F
0	.039/.035	.043/.040	.039/.035	.039/.035	.039/.035
10	.075/.127	.064/.109	.070/.124	.062/.126	.078/.148
20	.120/.176	.114/.195	.102/.171	.087/.155	.103/.181
30	.108/.195	.135/.215	.121/.223	.099/.178	.113/.211
40	.144/.236	.115/.187	.119/.201	.116/.197	.141/.230
50	.120/.204	.125/.196	.121/.216	.153/.224	.145/.258
60	.126/.178	.139/.202	.125/.203	.135/.190	.145/.237
70	.130/.200	.128/.200	.129/.197	.132/.227	.147/.257
80	.135/.182	.135/.202	.140/.204	.134/.197	.144/.206
90	.136/.191	.144/.212	.142/.178	.130/.210	.166/.239
100	.132/.171	.146/.207	.134/.192	.144/.212	.157/.222
110	.157/.208	.132/.181	.140/.171	.152/.218	.151/.230

Table 5: AIME24 $\text{acc@32}/\text{maj@32}$ at every validation point. Rows below the rule are the converged window used for all subsequent statistics.

12 Results

All numbers below come from the runs described in Section 11: one benchmark, 110 optimization steps per arm, and the same initial checkpoint, whose step-0 accuracy is identical across arms. Every figure is recomputed from a training log released with this paper. We report our own measurements only; no figure in this section is taken from prior work.

The campaign is not finished. Each arm below rests on 1 seed, and where an arm has more we hold it out of the tables rather than average arms over different seed sets — a reader subtracting two rows of Table 6 would otherwise get a different number from Table 7. Those extra runs are not wasted: Section 12.8 uses them for the one thing a single seed cannot give, a direct measurement of how far the same configuration moves when it is run again.

12.1 Validation trajectory

Table 5 gives AIME24 accuracy and majority-vote accuracy at every validation point. The four arms that share the tree pipeline start from an identical step-0 score; STEER starts at .043 rather than .039, a gap we attribute to re-sampling noise rather than to a different checkpoint, since two validations of the *same* base weights in our logs differ by .004 in accuracy.

12.2 Converged window

Table 6 averages each arm over steps 40–110.

Two orderings are worth stating before the contrasts. On accuracy, the two arms with no forecast (STEER, UNIFORM) sit *below* plain GRPO, and

arm	n	acc@32	maj@32	uplift	entropy
GRPO	1	.1350	.1962	.0612	.1491
STEER	1	.1330	.1984	.0654	.1214
UNIFORM	1	.1313	.1953	.0640	.1374
PERMUTED	1	.1370	.2094	.0724	.1190
STEER-F	1	.1495	.2349	.0854	.1210

Table 6: Means over the converged window (steps 40–110, eight validation points per run), averaged over the n seeds *common to all five arms*, so that no row moves merely because one arm is a seed ahead of the others. Each contrast in Table 7 pairs on the seeds *its own two arms* share, which can be more, so the difference of two rows here need not equal the contrast there; both seed sets are listed in the released `per_seed.tsv`. *uplift* is $\text{maj@32} - \text{acc@32}$. Entropy is the token-mean quantity discussed in Section 11, i.e. $\mathcal{H}$ divided by mean length.

contrast	n	Δacc	t	Δmaj	t
STEER-F – GRPO	1	+.0145	–	+.0386	–
STEER-F – STEER	2	+.0135	+4.50	+.0211	+1.36
STEER-F – UNIF	2	+.0114	+1.65	+.0249	+1.69
STEER-F – PERM	1	+.0125	–	+.0255	–
STEER – GRPO	3	–.0101	–2.33	–.0067	–1.11
UNIF – STEER	3	+.0036	+1.34	–.0031	–3.71

Table 7: Contrasts paired *by seed*: each arm is averaged over its converged window, arms are paired within a seed, and t is the paired t on the n seed-level differences ($\text{df} = n - 1$). Pairing within a seed cancels any per-seed effect common to all five arms. A contrast’s n is the smaller of its two arms’ seed counts.

only the arm carrying A_H is clearly above it. On entropy, GRPO retains the most and STEER-F among the least, so the two columns are neither aligned nor reversed. Section 12.5 shows that the entropy column tracks the damping apparatus rather than A_H , and why ranking five arm means is the wrong instrument for either question at this seed count.

12.3 Does the forecast’s value matter?

Table 7 averages each arm over its converged window, pairs the arms within a seed, and reports the mean difference together with a paired t on those seed-level differences. The step-by-step version of this statistic, which the single-seed draft reported, is in Appendix D; it answers a question about one trajectory rather than about re-running.

The block structure is what the design was built to expose. Every contrast that changes the *value* of A_H is positive; the one that leaves it alone is not distinguishable from zero. UNIFORM runs the entire apparatus — tree sampler, forecaster, support, weight map — and differs from STEER by +.0036 ($t = +1.34$ over 3 seeds): turning the machinery on and feeding it a constant buys nothing.

PERMUTED keeps the machinery *and* the magnitudes of A_H and destroys only the pairing between a branch and its score, and STEER-F still exceeds it by +.0125.

We state the signs, not significance. At the seed counts reached so far — the headline contrast rests on 1 — no accuracy contrast in Table 7 clears a conventional threshold, and the two rows with $n = 1$ carry no t at all. What the table shows is that the sign is the same wherever A_H carries information and absent where it does not; Section 12.8 states what the completed campaign is powered to detect, and the pre-registration fixes the test before the remaining seeds land.

This is the contrast the design was built for. A comparison of STEER-F against STEER alone changes two things at once, and cannot separate the forecast from the sampler that delivers it; the two destroyed-forecast arms do separate them.

12.4 Where the gain comes from

The majority-vote gain decomposes into a part explained by per-sample accuracy and a part explained by the sample set’s structure:

$$\underbrace{+.0386}_{\Delta\text{maj}} = \underbrace{+.0145}_{\Delta\text{acc}} + \underbrace{+.0241}_{\Delta\text{uplift}},$$

against GRPO; against STEER the uplift term is +.0076 ($t = +0.61$ over 2 seeds). Against GRPO roughly two thirds of the improvement is not that individual samples are more often right, but that the right ones agree with each other more and the wrong ones agree less. The two backbones of that split are measured on one seed so far, and the split itself is the claim to re-check as the remaining seeds arrive: against STEER the same term is smaller and well inside its own spread.

Where that structural change comes from is a separate question, and the arm that answers it finished after this section was first written. UNIFORM runs the sampler at $\lambda = .25$ with the forecast’s value destroyed and shows no uplift gain, which reads as though the structure belonged to A_H . It does not: the $\lambda = 0$ tree arm — the sampler with the future channel off, STEER’s local damping still in place — reaches an uplift of .0870, against .0654 for STEER and .0854 for STEER-F. On this evidence most of the structural gain is the sampler’s, and Section 12.10 reports it that way. The decomposition above is unaffected — it is measured against GRPO, which has no tree — but the attribution is not what the five-arm design alone could see.

contrast	n	$\Delta\text{ent.}$	t	Δacc	t
<i>the apparatus changes, A_H does not</i>					
STEER – GRPO	3	−.0487	−3.92	−.0101	−2.33
UNIF – STEER	3	+.0176	+21.67	+.0036	+1.34
<i>A_H changes, the apparatus does not</i>					
STEER-F – UNIF	2	−.0071	−0.76	+.0114	+1.65
STEER-F – PERM	1	+.0020	−	+.0125	−
<i>both change</i>					
STEER-F – STEER	2	+.0101	+0.96	+.0135	+4.50

Table 8: Converged aggregate entropy and accuracy, paired within a seed and grouped by what the contrast changes. Turning the damping apparatus on or off moves the aggregate entropy at $|t| \geq 4$; changing what A_H contains, with the apparatus held fixed, does not. Accuracy runs the other way. A row with $n = 1$ has no t .

12.5 Entropy: the aggregate moves with the apparatus, not with A_H

What the method predicts here, fixed in advance.

This method does not set out to hold the aggregate entropy up, and two facts already established above say it cannot. Section 9 bounds the positions at which A_H can be non-zero at $2(n-1)/T$, about one in eighty; Section 9.3 shows A_H sums to zero over a sibling set, so at the positions it does reach it redistributes weight among siblings rather than adding any. An intervention that is this sparse and mean-preserving by construction should leave an average over every response token close to where it found it. Both statements were written into the manuscript before the campaign was launched, Section 9.3 explicitly so (“we state it before presenting any experiment”), and the artifact history dates them an hour ahead of the pre-registration discussed below.

What the seeds show. Table 8 groups the contrasts by what each one changes, and the grouping separates them cleanly. The two contrasts that change the apparatus move the aggregate: turning the damping on costs −.0487 at $t = -3.92$, and adding the tree with damping applied uniformly returns +.0176 at $t = +21.67$ — the latter reproducing across all 3 seeds to within a thousandth. The contrasts that hold the apparatus fixed and change only what A_H contains move it by −.0071 ($t = -0.76$) when its value is discarded and +.0020 when its pairing is destroyed. Accuracy runs the other way: the apparatus contrasts give +.0036 and −.0101, while the two A_H contrasts give +.0114 and +.0125. The quantity that

moves the aggregate entropy is the damping apparatus; the quantity that moves accuracy is A_H . That is the predicted pattern, and it is the sense in which this paper is not a method for suppressing entropy.

A pre-registered condition we retire rather than score. The pre-registration (Section C) also fixed a condition on this section: if aggregate entropy orders the arms the same way accuracy does, a dissociation argument fails. We retire that condition instead of reporting it, and say plainly why, since retiring a registered test is exactly the move a reader should be suspicious of. First, it tests the aggregate as a signal about A_H , which the two paragraphs above — and Section 9.3, committed an hour before the registration was — say it cannot be. Second, an ordering over five arm means is not stable at this seed count: the rank correlation between converged entropy and converged accuracy comes out with a different sign depending on whether each arm is averaged over the seed common to all five (Table 6) or over its own, which differ in number while the campaign runs. Both seed sets are in the released `per_seed.tsv`. Third, and because of the first two, *we claim nothing from it in either direction*: the condition is not reported as passed, and the theory is not reported as damaged by it. It was the wrong instrument to have registered, and the honest correction is to say so and to register a better one.

The measurement that is actually owed. Section 9.3 states the prediction that a sum-zero correction does make: not that entropy at branch points rises, but that the *spread across siblings* widens. No quantity we currently log tests it. The trainer does log an entropy split (`steerf/branch_entropy`), and an earlier draft of this section read it as entropy at branch points; it is not. That monitor takes a fixed top decile of A_H — a tenth of all response tokens — while the positions A_H can reach are the one-in-eighty that Section 12.6 measures, an eighth as many. Since A_H is exactly zero at almost every position, the remainder of that decile is filled by an arbitrary tie-break rather than by branching. Section 12.6 reports what that monitor does and does not select. The replacement restricts the split to $A_H \neq 0$, the same predicate the correction itself uses, and is instrumented in the release; it reports − for the contrast that isolates the pairing, and the cells stay empty until runs carry it. That test is

registered in Section C, before those runs, and we state here that it is registered after this section’s original condition was retired rather than before.

What the numbers above cannot carry. The A_H contrasts rest on 2 and 1 seeds. A $|t|$ near one on two seeds is absence of evidence, not evidence of absence: we are entitled to say the aggregate does not move detectably when A_H changes, and not entitled to say it is unchanged. What carries weight is the contrast between the groups — the apparatus rows reach $|t| \geq 4$ on the same or larger seed counts — rather than any single row.

12.6 Mechanism diagnostics

Section 9 bounds the fraction of positions that can carry a non-zero A_H at $2(n-1)/T$ for any sampler. With $n = 8$ and a mean response length of 984 tokens the bound is .0142. Measured over training, the three tree arms sit at .0136 (UNIFORM), .0129 (PERMUTED) and .0120 (STEER-F), i.e. at 84–96% of a ceiling that no sampler can exceed. A plain-rollout run with the same λ reaches .003. As Section 8 notes, this count admits float32 round-off and is an upper bound; the runs reported here predate the round-off-free counter, so we quote the loose figure and label it as such.

Two readings follow. The tree sampler is close to saturating what the combinatorics allow, so there is little headroom in scheduling more cuts; and the intervention is extremely sparse — about one position in eighty — which makes the size of the effect in Table 7 the surprising part rather than the small support.

That sparsity also disqualifies one of our own monitors, which we record here rather than leave for a reader to find. The logged entropy split (`steerf/branch_entropy` and its gap) selects the top decile of A_H by rank, a tenth of all response tokens, against a support an eighth that size. A rank-based cut was chosen deliberately — A_H is exactly zero wherever a rollout has become its own only sibling, so a value-based threshold selects almost everything — but with a support this small the decile is filled out by the tie-break among those zeros, and the tie-break is not a random sample of them. A top-decile cut does recall about half of the true divergence points, five times what selecting at random would give, so the bucket is genuinely enriched; it is nonetheless mostly not branch positions, and the mean entropy it reports cannot be attributed to them. Section 12.5 states

what replaces it.

12.7 Is the future term a local term?

Section 6 reports a calibration that suppresses every head past the first, leaving $\hat{H}_{\text{togo}}$ dominated by one step ahead. That invites the sharpest objection to the paper: if the forecast reaches one step, is A_H anything a purely local quantity could not have produced? The question is answerable without training anything. We apply the *same* operator — the same sibling baseline, mask and grouping — to two inputs: the calibrated forecast, and the policy’s own next-token entropy read at the same offset, the state after the branch token. Any difference between the two is the forecaster’s contribution and nothing else. We report it on the sibling support, where A_H can be non-zero at all; outside it both are identically zero and a statistic over all positions measures the support rather than the forecast.

The answer is that the forecaster is not earning its cost. Over 464 branch points from 4 validation prompts, the two advantages correlate at $r = +0.917$ (Spearman $+0.937$), the local twin explains $R^2 = 0.840$ of the treatment, and the two agree in sign 0.921 of the time — and the sign is what the token weights are monotone in. The levels correlate less than the advantages ($+0.881$), which is the wrong way round for an artefact of scale.

The mechanism is one the paper had already measured without connecting it. With $\kappa = 2$ and $\gamma_H = 0.7$ the calibration instantiates to $\hat{H}_{\text{togo}} = 0.72 H_1 + 0.06 H_2 + 0.11$, the sibling baseline removes the constant, and the first head is built to reproduce the policy’s own next-token distribution — its final layer is zero-initialised and its fitted $a_1 = 1.03$ says it stayed there. So A_H should be close to $\gamma_H a_1 = 0.72$ times the local twin, and the measured magnitudes are 0.3445 against 0.4647, a ratio of 0.741. A_H is very nearly a scalar multiple of a quantity the trainer already computes, and Eq. (23) divides the scale out through $\rho = \text{rms}\{\text{visit}\}$.

Two things this does *not* say. It does not touch Eq. (8), which is an identity, nor Section 4.4: the local twin is not Ω . Ω carries no sibling baseline and no tree, vanishes at a uniform branch point, and the twin exists only because the tree supplies siblings to difference against. What the measurement indicts is the *forecasting apparatus* — the MTP heads, their warm-up and their calibration — not the visitation channel. Read alongside Section 12.10, where the sampler alone recovers .1439

against STEER’s .1330 while the forecast adds the remainder, two independent measurements now point the same way.

The caveats are the ones a single diagnostic run carries: 4 prompts at one validation step of the base policy, on independently sampled rollouts, whose sibling support therefore sits early in the response rather than at the scheduled branch depths. The head-free arm of Table 10 is the direct test, and it is cheaper than the arm it would replace: ORACLE reads H_{togo} from the policy’s realised entropy, which verl computes in the same pass, so it carries none of the overhead in Section 12.13.

12.8 Seeds, not steps

Every statistic above treats the seed as the unit of analysis: each run is first collapsed to one number by averaging its converged window, the arms are then paired within a seed, and the test is a one-sample t on the paired differences with $\text{df} = n - 1$. Pairing within a seed matters because a seed carries effects common to all five arms — the data order, and on a split campaign the machine a run landed on — and those cancel exactly in a difference. A contrast’s n is the smaller of its two arms’ seed counts, which is why the table reports n per row rather than once.

The alternative statistic, a paired t over the eight window steps *within* one run, measures whether a gap is stable along a single trajectory. That is a different question from whether it survives re-running, and we report it only in Appendix D so the two are not confused.

One run is excluded, and the pre-registration says which. STEER-F at seed 4 stopped at training step 91 when its machine lost a GPU, so its window covers six validation points where every other run covers eight. A mean over six points is not a shorter version of a mean over eight; paired against the others it is a different statistic. The pre-registration (Appendix C) excludes a run for mechanical failure and relaunches it with the same seed, so that is what happens here, and the release reports the run with an `partial` flag rather than deleting it. It is counted in the spread below, for the reason given there.

Run-to-run spread, and it is the number that governs this section. Every arm has now finished more than once. GRPO’s converged-window accuracy moves from .1350 to .1560 across its 3 seeds, a spread of .0210; STEER-F moves from .1392 to

arm	AIME24	AIME25	AMC23	MATH500	Minerva	Olympiad	mean
							—

Table 9: Six-benchmark evaluation of each arm’s selected checkpoint, averaged over seeds. The rightmost column is the unweighted mean, the headline figure of the post-hoc evaluation.

.1497 across its 3 seeds, a spread of .0106; PERMUTED moves only .0014 across 3. Both extremes are measured on the same benchmark with the same protocol, and they differ by more than an order of magnitude.

Set the headline gap of +.0145 against those. **It is smaller than GRPO’s own spread and comparable to STEER-F’s**, and the two arms it compares are each observed a handful of times. That is the regime in which a single paired difference cannot decide the question, and we do not claim it does. A range between two or three draws is also not a standard error: three such ranges differing by a factor of fifteen do not bound each other, so no contrast here should be read against a noise floor quoted to one significant figure. For scale at the other end, two validations of the *same* weights on these 30 questions differ by about .004, so seed variation is the larger term by roughly an order of magnitude and any comparison quoted at single-run precision is quoting noise.

We therefore state what this experiment establishes and what it does not. The mechanism contrasts — what changes when A_H ’s value is discarded or its pairing destroyed, with the apparatus held fixed — are the claim. The effect size is what the remaining seeds are for, and on the seeds in hand the headline gap is not separated from the spread.

12.9 Six benchmarks

A single 30-question benchmark cannot separate a method from a lucky draw, so each run’s selected checkpoint is re-evaluated on six mathematics benchmarks (Table 9) under the protocol of Section 11: AIME24, AIME25 and AMC23 at `avg@32`, MATH500, Minerva and OlympiadBench at `avg@1`.

The quantity fixed in advance as the headline here is not any single cell but *direction consistency*: in how many of the six benchmarks the ordering observed on AIME24 is preserved. Four or better of six we declared support, three equivocal, two or fewer evidence against. The observed count

arm	what changes	acc@32	maj@32
<i>the damping coefficient</i>			
$\lambda = 0.1$	weaker	.1365	.2360
$\lambda = 0.25$	as used	.1495	.2349
$\lambda = 0.5$	stronger	.1444	.2265
<i>the sampler, with the forecast switched off</i>			
$\lambda = 0$, tree	tree only	.1439	.2309
STEER	neither	.1330	.1984
<i>the forecaster, replaced by what the trainer already has</i>			
ORACLE	no heads	—	—
<i>the base method at its released damping strength</i>			
STEER, $\lambda_{\min}=0.8$	weaker	—	—
<i>relaxed ratio clipping ($\epsilon_{hi}=5$, $\epsilon_{lo}=0.99$)</i>			
STEER-F	—	.1516	.2304
STEER	—	.1393	.2063
<i>other advantage estimators</i>			
RLOO, STEER-F	—	—	—
RLOO, STEER	—	—	—
OPO, STEER-F	—	—	—
OPO, STEER	—	—	—

Table 10: Follow-up ablations, seed 1 only, converged window. Each block isolates one element of the design.

is — of six for STEER-F > UNIFORM, and — of six for STEER-F > PERMUTED. This is the test that distinguishes a method from an artifact of one question set, and it is reported whatever it returns.

12.10 How much of the design is load-bearing

Table 10 varies one element at a time, at a single seed. Every comparison in it is therefore $n = 1$: these rows describe a mechanism, and none of them carries an error bar or licenses an effect size.

The coefficient. λ scales the visitation term before the two channels are combined in Eq. (23). The value used throughout was not tuned on these runs, and the sweep exists to say whether the result rests on a knife edge. A flat response across 0.1–0.5 would say the method is insensitive to it; a peak at 0.25 with both neighbours below would say it is a tuned quantity and must be reported as one.

The damping strength, and what it does and does not reach. We measure STEER below GRPO by $-.0101$ over 3 seeds, which the base method’s own results contradict. The explanation is not a difference in the code — we read the released weighting function against ours and it agrees line for line, as do every other value in the two launchers — but in $\lambda_{\min}$, which we set to 0.7 by selecting it on STEER-F while the released script uses 0.8. The baseline is therefore running at a damping strength chosen for the method it is the control for, and one that attenuates half again as hard as its own default. Section 4.4 predicts the

sign of that: if the local channel is blind at branch points and anti-correlated away from them, more of it should cost accuracy. This arm is stock STEER ($\lambda = 0$) at the released value, and separates a failure to reproduce from an over-damped control.

The reach of the problem is bounded, and by the design rather than by luck. All four damped arms share $\lambda_{\min} = 0.7$ exactly, so it cancels in any contrast between two of them. STEER-F — UNIFORM and STEER-F — PERMUTED — the pair that carries this paper’s central claim, that the forecast’s value is what matters — hold $\lambda_{\min}$, λ , the sampler and the weight map all fixed and vary only what A_H contains. Those two are unaffected by how $\lambda_{\min}$ was chosen. What the choice does reach is the two contrasts against an arm that does not share it: STEER-F — STEER, where only the treatment’s side was tuned, and STEER — GRPO, where GRPO has no such knob at all. We report those two with this stated rather than resolved.

The sampler. This is the cell no other arm supplies. Tree rollouts make a group share prefixes, which changes the GRPO advantage estimate itself, quite apart from any damping. UNIFORM and PERMUTED both run at $\lambda = .25$ and so cannot separate the sampler from the correction. The $\lambda = 0$ tree arm switches the forecast off while leaving the sampler in place, and is the direct answer to “is tree rollout simply a better sampler”. Read against STEER, which has neither, its difference is the sampler’s own contribution; read against STEER-F, what remains is the forecast’s.

The answer is not the one the rest of this paper was written expecting. The sampler alone reaches .1439 against STEER’s .1330, and .2309 against .1984 — about two thirds of STEER-F’s accuracy margin over STEER and the large majority of its majority-vote margin. Its uplift, .0870, is not below STEER-F’s .0854. What remains for the forecast is the accuracy gap between .1439 and .1495.

Two things keep this from being decisive, and neither rescues the original reading. Every cell here is one run, and the spread across this block is about twice the run-to-run spread measured in Section 12.8 — suggestive, not settled. And the λ sweep is not monotone: $\lambda = .1$ sits below $\lambda = 0$ on accuracy, which is what noise looks like at this sample size. We report the block as measured and do not reconcile it. The five-seed campaign of Appendix C decides whether the sampler or the forecast carries the structural term; until then, UNI-

arm	steps	wall clock	acc@32
STEER-F	110	—	.1495
GRPO, step-matched	110	—	.1350
GRPO, compute-matched	—	—	—

Table 11: Equal steps versus equal compute. The matched step is read from the integrated per-step time of the long GRPO run, on the same hardware as its STEER-F reference; times measured on a different machine are not comparable and are not mixed in.

FORM and PERMUTED establish that the forecast’s *value* matters at fixed λ , and nothing in this paper establishes that the tree does not.

Clipping and the advantage estimator. The relaxed-clipping setting is where entropy interventions are supposed to matter most, because ratio clipping otherwise removes much of the update a token weight would have scaled. The RLOO and OPO rows test whether the construction depends on GRPO specifically: the sibling-prefix baseline of Eq. (15) depends only on the group sharing a prefix, which those estimators also provide, so the method should transfer unchanged. Each is reported beside its own base-method control, so every comparison stays within one estimator.

12.11 Compute-matched comparison

STEER-F costs more per step than GRPO, almost entirely in the MTP forward pass, so an equal-step comparison hands it more wall clock. The honest comparison holds compute fixed instead. We run GRPO to — steps, integrate its measured per-step time, and read its accuracy at the last step whose cumulative time fits inside what STEER-F spent reaching step 110 — rather than assuming a cost ratio.

Both readings are reported. If the gain survives the compute-matched comparison it is a gain; if it does not, the method buys accuracy with time and should be described that way.

12.12 Does the sign survive a change of backbone?

Everything above is one model. A reviewer’s first question is whether the effect is a property of the method or of that model, and the honest form of the answer is narrow: on three further backbones we run only GRPO, the base method, and ours, and we ask whether the *ordering* holds. Absolute accuracies are not comparable across rows—the backbones differ in scale, in pre-training fam-

	GRPO	base	ours	Δ GRPO	Δ base
1.5B Qwen-Math	.1350	.1330	.1495	+.0145	+.0135
7B Qwen-Math	—	—	—	—	—
3B Llama	—	—	—	—	—
7B Mistral	—	—	—	—	—

Table 12: The headline contrast on four backbones, seed 1, converged window. The two rightmost columns are within-backbone differences; columns are not averaged, because the rows are not on a common scale. The Llama row uses the instruction-tuned checkpoint (Section 11.2); the other three are base models.

ily, in whether mathematics was specialized for, and (Section 11.2) in which benchmark selects the checkpoint—so the table reports each backbone’s own three numbers and the within-backbone differences, and nothing is averaged down a column.

Two questions are separable here, and the design answers them one at a time. The 7B model shares its tokenizer, architecture and pre-training corpus with the 1.5B, so a difference between those two rows is about *scale*. The Llama and Mistral backbones share none of that, so a difference there is about *family*, and in particular about whether the effect needs a base policy that is already sharp on mathematics. Our account predicts that the visitation channel matters wherever the rollout distribution has branch points that differ in what they lead to, which is a property of the task rather than of the pre-training corpus; a backbone on which the sign reverses would be evidence against that account, not merely a weaker result.

The statistic we report is the sign of Δ_{GRPO} , and how many of the four backbones share it. With one seed per cell, a single backbone’s difference carries about the same uncertainty as a single seed on the primary backbone (Section 12.8), so no individual row is decisive; the claim rests on agreement across rows. We report the count, — of 4, and treat a split result as a limit on generality rather than as noise to be averaged away.

12.13 Cost

Within the runs executed on identical hardware, STEER takes 833 s per optimization step and STEER-F takes 1516 s, an overhead of 82%. The UNIFORM and PERMUTED arms cost 1480 and 1492 s, confirming that the cost is the tree rollout plus the MTP forward pass and not the A_H computation, which is a subtraction over a mask. We do not report a GRPO timing because that arm ran on a different machine.

12.14 A note on how these numbers were selected

We report window means rather than the best checkpoint, and the difference is not cosmetic. Selecting the maximum over the twelve validation points inflates every arm by a similar amount: $+0.013$ for STEER, $+0.011$ for UNIFORM, $+0.016$ for PERMUTED and $+0.017$ for STEER-F, relative to the window mean. Since a single AIME24 measurement carries a standard error of $.0282$ across problems, a maximum over a dozen noisy points is a biased statistic whose bias grows with the number of checkpoints compared. The arm ordering happens to be the same under both rules here, but we regard the window mean as the defensible one and use it throughout.

13 Conclusion

The gradient of trajectory entropy has two terms, and the token-level entropy controls now in use estimate one of them. We showed that the omission is structural: the estimated term vanishes exactly where the omitted term is largest, and it vanishes for a correct reason, so the gap cannot be closed by improving the estimator. We then showed that the omitted term is a REINFORCE estimator whose reward-to-go is the remaining conditional entropy, that a parallel forecaster estimates it for the price of one forward pass with an error whose sign and monotonicity are known in closed form, and that a tree-structured rollout supplies the baseline it requires. The resulting correction is mean-preserving by construction: it does not slow entropy collapse, it redirects it.

Limitations

A single validation benchmark, and what our t -statistics do not cover. Every claim rests on AIME 2024, which has 30 problems. A single measurement of one arm carries an across-problem standard error of $.0282$, so the gaps in Table 7 are a fraction of one problem in absolute terms. The paired t we report is computed over the eight validation steps of one run: it establishes that a gap is stable across the converged window, and nothing more. The statistic that would speak to generalization — the across-problem standard error of the *paired* difference, in which per-problem difficulty cancels between two arms evaluated on the same 30 questions — is not recoverable from our logs, which record per-benchmark means and standard

deviations rather than per-problem scores. We state this rather than substitute a number we cannot compute. A multi-benchmark evaluation, and logging per-problem scores so the paired error can be estimated, are the two things that would settle it.

A single seed, and a gap the same order as the noise on it. Every arm rests on one data seed, so no contrast in this paper carries a seed-level test. This is the limitation that bounds all the others. We can now put a number on it rather than an argument, though not yet a stable one: of the two arms run twice, GRPO moved $.0210$ in converged-window accuracy between its seeds and PERMUTED moved $.0014$, a sixfold disagreement between two estimates of the same quantity. Against the larger, and a headline gap of $+0.0145$, the gap is the bigger of the two, and on majority vote by more than a factor of two, but not by the margin that would let a single draw from each arm settle it — and the sixfold spread between the two estimates means the floor itself is not yet measured. The across-problem standard error we quote elsewhere is a different quantity and does not bound this one. Rollout sampling is not seeded across processes, so a re-run of the same configuration is a second draw rather than a restoration — which is what makes the repeat measurable, and what makes it necessary. The pre-registered five-seed campaign in Appendix C is designed to settle the effect size; what this paper establishes without it is the ordering of the arms and the mechanism that produces it, since the three controls sit at zero while only the arm carrying A_H moves.

The forecaster does not earn its cost. The calibration in Eq. (19) suppresses heads beyond the first, so $\hat{H}_{\text{togo}}$ is dominated by a single step ahead. This is the predicted behavior of Eq. (18) rather than a bug, but Section 12.7 measures the consequence and it is worse than a short horizon: put through the same sibling baseline, the policy’s own next-token entropy reproduces A_H at $r = +0.917$ and $R^2 = 0.840$, with the magnitudes in the ratio 0.741 that γ_{Ha1} predicts. On that evidence the MTP heads, their warm-up and their calibration are machinery around a quantity the trainer already computes, and they carry the overhead of Section 12.13 to do it. We report the ORACLE arm of Table 10 as the direct test. What the measurement does not reach is the visitation channel itself: the local twin is not Ω , and it exists only because the tree supplies siblings to difference against. A

forecaster whose error did not grow with k would test the idea more sharply; a sequential forecaster is the obvious comparison and we have not run it.

One hyperparameter was selected on the method and applied to its baseline. $\lambda_{\min}$, the floor of the weight map, is 0.7 throughout, chosen because it was best for STEER-F; the base method’s released script uses 0.8. The base method therefore appears here at a damping strength picked for the method it controls for, and at 1.5 times the maximum attenuation of its own default, which is the most likely reason we measure it below GRPO where its authors do not. We did not re-tune $\lambda_{\min}$ per arm, and a reader who wants each method at its own best setting will not find it here.

What this does not reach is the claim the paper makes. All four damped arms carry the same $\lambda_{\min}$, so it cancels between any two of them, and the contrasts that separate the forecast’s value from its apparatus — STEER-F against UNIFORM and against PERMUTED — hold it fixed along with everything else. The two contrasts it does reach, STEER-F — STEER and STEER — GRPO, are reported with the confound stated; Section 12.10 adds the base method at its released value so the reproduction question can be read on its own.

The support is small by construction. Section 9 bounds the fraction of positions carrying a non-zero A_H at $2(n - 1)/T$, on the order of a percent for our settings. Whether the effect scales with n , or with a branching schedule chosen adaptively rather than at fixed depths, is untested.

An assumption with no remedy. Approximation B2 — that a token’s update does not change the conditional distribution at other positions — is false in a network with shared parameters. It is inherited from the base method rather than introduced here, but our construction adds a second estimator that relies on it.

One prediction of the analysis is unmeasured. Section 4.4 argues that the two channels are anti-correlated. The prefactor half of that claim is analytic and the A_H half is inferred from a correlation with branch diversity; we have not measured their *product* at real sibling-divergence positions in a live rollout. Until that measurement exists, Table 1 is a prediction rather than an observation.

Scope of the backbone evidence. The full decomposition—the five arms, the ablations, the

mechanism diagnostics—rests on one 1.5B model. The three further backbones (Section 12.12) carry the three-way contrast at one seed each, which establishes the sign and not its size, and they stop at 7B: we have no evidence about models an order of magnitude larger, where the base policy is sharper and the relative size of the two channels may differ. Family and scale are also not fully crossed—the mathematics backbones span 1.5B to 7B and the others sit at 3B and 7B—so what the latter test is whether the sign survives a different pre-training corpus, not whether it survives one at 8B or beyond. Those rows further differ from the mathematics ones in more than family. They are selected on a different benchmark (Section 11.2), and the Llama row is instruction-tuned where the other three are base models, which was forced rather than chosen: the base checkpoint left almost every training group degenerate. Each difference is what makes the row measurable at all, and each is a reason no claim here compares accuracy across backbones—only the sign of a within-backbone contrast is read.

Ethics Statement

This work studies an optimization-time regularizer for reinforcement learning on mathematics problems with automatically verifiable answers. It introduces no new data collection, no human subjects, and no new model release. The intervention acts on which of several equally-scored continuations receives a slightly smaller gradient weight; it does not alter the reward, and so does not change what behavior is being optimized for. The training data is a public competition-mathematics set and the evaluation set is a public examination.

The main foreseeable risk is indirect. Methods that preserve diversity during RL fine-tuning are not specific to mathematics, and the same mechanism would apply to a reward model encoding any objective, including a harmful one; we note that our method inherits, and does not mitigate, whatever the reward function encodes. The computational cost of the method (a forward pass through a set of small prediction heads at every optimization step) is a real environmental cost, which we quantify in Section 12.13 so that it can be weighed against the benefit.

References

Arash Ahmadian, Chris Cremer, Matthias Gallé, Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ah-

1613	met Üstün, , and Sara Hooker. 2024. Back to basics: Revisiting reinforce-style optimization for learning from human feedback in llms. In Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 12248–12267.	Ganqu Cui, Lifan Yuan, Zefan Wang, Hanbin Wang, Wendi Li, Bingxiang He, Yuchen Fan, Tianyu Yu, Qixin Xu, Weize Chen, , et al. 2025a. Process reinforcement through implicit rewards. ArXiv preprint arXiv:2502.01456.	1667
1614			1668
1615			1669
1616			1670
1617			1671
1618			
1619	Shengnan An, Xunliang Cai, Xuezhi Cao, Xiaoyu Li, Yehao Lin, Junlin Liu, Xinxuan Lv, Dan Ma, Xuanlin Wang, Ziwen Wang, , et al. 2025. Amo-bench: Large language models still struggle in high school math competitions. ArXiv preprint arXiv:2510.26768.	Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, , et al. 2025b. The entropy mechanism of reinforcement learning for reasoning language models. ArXiv preprint arXiv:2505.22617.	1672
1620			1673
1621			1674
1622			1675
1623			1676
1624	Anonymous. 2025. Rethinking entropy regularization in large reasoning models. ArXiv preprint arXiv:2509.25133.	Yu Cui, Feng Liu, Jiawei Chen, Canghong Jin, Xingyu Lou, Changwang Zhang, Jun Wang, Yuegang Sun, , and Can Wang. 2025c. Hatllm: Hierarchical attention masking for enhanced collaborative modeling in llm-based recommendation. ArXiv preprint arXiv:2510.10955.	1677
1625			1678
1626			1679
1627	Anthropic. 2025. Introducing claude opus 4.5. https://www.anthropic.com/news/claude-opus-4-5 .		1680
1628			1681
1629	Sikai Bai, Haoxi Li, Jie Zhang, Yongjiang Liu, , and Song Guo. 2026. Ttvs: Boosting self-exploring reinforcement learning via test-time variational synthesis. ArXiv preprint arXiv:2604.08468.	Jia Deng, Jie Chen, Zhipeng Chen, Wayne Xin Zhao, , and Ji-Rong Wen. 2025. Decomposing the entropy-performance exchange: The missing keys to unlocking effective reinforcement learning. ArXiv preprint arXiv:2508.02260.	1682
1630			1683
1631			1684
1632			1685
1633	Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwei Peng, Jason D. Lee, Deming Chen, and Tri Dao. 2024. Medusa: Simple LLM inference acceleration framework with multiple decoding heads. ArXiv preprint arXiv:2401.10774.	Chongming Gao, Kexin Huang, Jiawei Chen, Yuan Zhang, Biao Li, Peng Jiang, Shiqi Wang, Zhong Zhang, , and Xiangnan He. 2023a. Alleviating matthew effect of offline reinforcement learning in interactive recommendation. In Proceedings of the 46th international ACM SIGIR conference on research and development in information retrieval, pages 238–248.	1686
1634			1687
1635			1688
1636			1689
1637			1690
1638	Yupeng Chang, Yi Chang, , and Yuan Wu. 2026. Balora: Bias-alleviating low-rank adaptation to mitigate catastrophic inheritance in large language models. In The Fourteenth International Conference on Learning Representations.	Chongming Gao, Shiqi Wang, Shijun Li, Jiawei Chen, Xiangnan He, Wenqiang Lei, Biao Li, Yuan Zhang, , and Peng Jiang. 2023b. Cirs: Bursting filter bubbles by counterfactual interactive recommender system. acm transactions on information systems, 42(1):1–27.	1691
1639			1692
1640			1693
1641			1694
1642			1695
1643	Yupeng Chang, Chenlu Guo, Yi Chang, , and Yuan Wu. 2025. Lora-mgpo: Mitigating double descent in low-rank adaptation via momentum-guided perturbation optimization. in findings of the association for computational linguistics: Emnlp 2025, pages 648–659.	Google. 2025. A new era of intelligence with gemini 3. https://blog.google/products/gemini/gemini-3/ .	1696
1644			1697
1645			1698
1646			1699
1647			1700
1648	Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, , et al. 2024. A survey on evaluation of large language models. acm transactions on intelligent systems and technology, 15(3):1–45.	Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, , et al. 2024. The llama 3 herd of models. ArXiv preprint arXiv:2407.21783.	1701
1649			1702
1650			1703
1651			1704
1652			1705
1653			1706
1654	Jiawei Chen, Hande Dong, Xiang Wang, Fuli Feng, Meng Wang, , and Xiangnan He. 2023. Bias and debias in recommender system: A survey and future directions. acm transactions on information systems, 41(3):1–39.	Aaron Grattafiori et al. 2024. The llama 3 herd of models. ArXiv preprint arXiv:2407.21783.	1707
1655			1708
1656			1709
1657			1710
1658			1711
1659	Daixuan Cheng, Shaohan Huang, Xuekai Zhu, Bo Dai, Wayne Xin Zhao, Zhenliang Zhang, , and Furu Wei. 2025. Reasoning with exploration: An entropy perspective. ArXiv preprint arXiv:2506.14758.	Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, , et al. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. ArXiv preprint arXiv:2501.12948.	1712
1660			1713
1661			1714
1662			1715
1663	Xiangxiang Chu, Hailang Huang, Xiao Zhang, Fei Wei, , and Yong Wang. 2025. Gpg: A simple and strong reinforcement learning baseline for model reasoning. ArXiv preprint arXiv:2504.02546.	Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, , and Sergey Levine. 2018. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. in international conference on machine learning, pages 1861–1870. pmlr.	1716
1664			1717
1665			1718
1666			1719
			1720

1721	Yaru Hao, Li Dong, Xun Wu, Shaohan Huang, Zewen Chi, , and Furu Wei. 2025. On-policy rl with optimal reward baseline. ArXiv preprint arXiv:2505.23585.	1776
1722		1777
1723		
1724	Zhezhen Hao, Hong Wang, Haoyang Liu, Jian Luo, Jiarui Yu, Hande Dong, Qiang Lin, Can Wang, and Jiawei Chen. 2026. Rethinking entropy interventions in RLVR: An entropy change perspective. In <i>Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 31105–31133.	1778
1725		1779
1726		1780
1727		1781
1728		1782
1729		1783
1730		1784
1731	Andre He, Daniel Fried, , and Sean Welleck. 2025a. Rewarding the unlikely: Lifting grpo beyond distribution sharpening. ArXiv preprint arXiv:2506.02355.	1785
1732		1786
1733		1787
1734		1788
1735		
1736	Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, , et al. 2024. Olympiad-bench: A challenging benchmark for promoting agi with olympiad-level bilingual multimodal scientific problems. ArXiv preprint arXiv:2402.14008.	1789
1737		1790
1738		1791
1739		1792
1740		1793
1741		1794
1742		1795
1743		
1744	Jujie He, Jiakai Liu, Chris Yuhao Liu, Rui Yan, Chaojie Wang, Peng Cheng, Xiaoyu Zhang, Fuxiang Zhang, Jiacheng Xu, Wei Shen, , et al. 2025b. Skywork open reasoner 1 technical report. ArXiv preprint arXiv:2505.22312.	1796
1745		1797
1746		1798
1747		1799
1748		1800
1749		1801
1750		
1751	Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, , and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the math dataset. ArXiv preprint arXiv:2103.03874.	1802
1752		1803
1753		1804
1754		1805
1755		1806
1756		
1757	Jian Hu, Jason Klein Liu, Haotian Xu, , and Wei Shen. 2025a. Reinforce++: An efficient rlhf algorithm with robustness to both prompt and reward models. ArXiv preprint arXiv:2501.03262, 1(3):5.	1807
1758		1808
1759		1809
1760		1810
1761		1811
1762		
1763	Jingcheng Hu, Yinmin Zhang, Qi Han, Daxin Jiang, Xiangyu Zhang, , and Heung-Yeung Shum. 2025b. Open-reasoner-zero: An open source approach to scaling up reinforcement learning on the base model. ArXiv preprint arXiv:2503.24290.	1812
1764		1813
1765		1814
1766		1815
1767		1816
1768		1817
1769		1818
1770		
1771	Binyuan Hui, Jian Yang, Zeyu Cui, Jiayi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Keming Lu, , et al. 2024. Qwen2.5-coder technical report. ArXiv preprint arXiv:2409.12186.	1819
1772		1820
1773		1821
1774		1822
1775		1823
	Zed Industries. 2025. zeta. https://huggingface.co/datasets/zed-industries/zeta .	1824
		1825
		1826
	Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, , et al. 2024. Openai o1 system card. ArXiv preprint arXiv:2412.16720.	1827
		1828
	Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, , and Ion Stoica. 2024. Live-codebench: Holistic and contamination free evaluation of large language models for code. ArXiv preprint arXiv:2403.07974.	1829
		1830
		1831
	Albert Q. Jiang et al. 2023. Mistral 7b. ArXiv preprint arXiv:2310.06825.	
	Yihang Jiang, Xiaoyang Li, Guangxu Zhu, Hang Li, Jing Deng, Kaifeng Han, Chao Shen, Qingjiang Shi, , and Rui Zhang. 2023. 6g non-terrestrial networks enabled low-altitude economy: Opportunities and challenges. ArXiv preprint arXiv:2311.09047.	
	Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, , et al. 2024. Tulu 3: Pushing frontiers in open language model post-training. ArXiv preprint arXiv:2411.15124.	
	Aitor Lewkowycz, Anders Andreassen, David Do-han, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, , et al. 2022. Solving quantitative reasoning problems with language models. advances in neural information processing systems, 35:3843–3857.	
	Jia Li, Edward Beeching, Lewis Tunstall, Ben Lipkin, Roman Soletskyi, Shengyi Huang, Kashif Rasul, Longhui Yu, Albert Q Jiang, Ziju Shen, , et al. 2024. Numinamath: The largest public dataset in ai4maths with 860k pairs of competition math problems and solutions. hugging face repository, 13(9):9.	
	Mengdi Li, Jiaye Lin, Xufeng Zhao, Wenhao Lu, Peilin Zhao, Stefan Wermter, , and Di Wang. 2025a. Curriculum-rlaif: Curriculum alignment with reinforcement learning from ai feedback. ArXiv preprint arXiv:2505.20075.	
	Qingbin Li, Rongkun Xue, Jie Wang, Ming Zhou, Zhi Li, Xiaofeng Ji, Yongqi Wang, Miao Liu, Zheming Yang, Minghui Qiu, , et al. 2025b. Cure: Critical-token-guided re-concatenation for entropy-collapse prevention. ArXiv preprint arXiv:2508.11016.	
	Siyi Lin, Chongming Gao, Jiawei Chen, Sheng Zhou, Binbin Hu, Yan Feng, Chun Chen, , and Can Wang. 2025. How do recommendation models amplify popularity bias? an analysis from the spectral perspective. In <i>Proceedings of the Eighteenth ACM International Conference on Web Search and Data Mining</i> , pages 659–668.	
	Aixin Liu, Aoxue Mei, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, , et al. 2025a. Deepseek-v3. 2: Pushing the frontier of open large language models. ArXiv preprint arXiv:2512.02556.	
	Mingjie Liu, Shizhe Diao, Ximing Lu, Jian Hu, Xin Dong, Yejin Choi, Jan Kautz, , and Yi Dong. 2025b. Prorl: Prolonged reinforcement learning expands reasoning boundaries in large language models. ArXiv preprint arXiv:2505.24864.	
	Ziyu Liu, Yuhang Zang, Shengyuan Ding, Yuhang Cao, Xiaoyi Dong, Haodong Duan, Dahua Lin, , and Jiaqi Wang. 2025c. Spark: Synergistic policy	

1832	and reward co-evolving framework. ArXiv preprint	Jiakang Wang, Runze Liu, Fuzheng Zhang, Xiu Li, ,	1884
1833	arXiv:2509.22624.	and Guorui Zhou. 2025c. Stabilizing knowledge,	1885
		promoting reasoning: Dual-token constraints for rlvr.	1886
1834	Michael Luo, Sijun Tan, Justin Wong, Xiaoxiang Shi,	ArXiv preprint arXiv:2507.15778.	1887
1835	William Y Tang, Manan Roongta, Colin Cai, Jeffrey		
1836	Luo, Tianjun Zhang, Li Erran Li, , et al. 2025. Deep-	Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shix-	1888
1837	scaler: Surpassing o1-preview with a 1.5 b model by	uan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin	1889
1838	scaling rl. notion blog.	Yang, Zhenru Zhang, , et al. 2025d. Beyond the	1890
		80/20 rule: High-entropy minority tokens drive effective	1891
1839	Shichao Ma, Zhiyuan Ma, Ming Yang, Xiaofan Li, Xing	reinforcement learning for llm reasoning. ArXiv	1892
1840	Wu, Jintao Du, Yu Cheng, Weiqiang Wang, Qiliang	preprint arXiv:2506.01939.	1893
1841	Liu, Zhengyang Zhou, , et al. 2026. Tspo: Break-		
1842	ing the double homogenization dilemma in multi-	Ziqing Wang, Yibo Wen, William Pattie, Xiao Luo,	1894
1843	turn search policy optimization. ArXiv preprint	Weimin Wu, Jerry Yao-Chieh Hu, Abhishek Pandey,	1895
1844	arXiv:2601.22776.	Han Liu, , and Kaize Ding. 2025e. Polo: Pref-	1896
		erenceguided multi-turn reinforcement learning for	1897
1845	Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi	lead optimization. ArXiv preprint arXiv:2509.21737.	1898
1846	Mirza, Alex Graves, Timothy Lillicrap, Tim Harley,		
1847	David Silver, , and Koray Kavukcuoglu. 2016. Asyn-	Fang Wu, Weihao Xuan, Ximing Lu, Zaid Harchaoui, ,	1899
1848	chronous methods for deep reinforcement learning, in	and Yejin Choi. 2025a. The invisible leash: Why	1900
1849	international conference on machine learning, pages	rlvr may not escape its origin. ArXiv preprint	1901
1850	1928–1937. pmlr.	arXiv:2507.14843.	1902
1851	OpenAI. 2025. Introducing gpt-5.2. https://openai.com/index/introducing-gpt-5-2 .	Fei Wu, Zhenrong Zhang, Qikai Chang, Jianshu Zhang,	1903
1852		Quan Liu, , and Jun Du. 2026. Step potential advan-	1904
		tage estimation: Harnessing intermediate confidence	1905
1853	John Schulman, Filip Wolski, Prafulla Dhariwal, Alec	and correctness for efficient mathematical reasoning.	1906
1854	Radford, , and Oleg Klimov. 2017. Proximal	ArXiv preprint arXiv:2601.03823.	1907
1855	policy optimization algorithms. ArXiv preprint		
1856	arXiv:1707.06347.	Junkang Wu, Kexin Huang, Jiancan Wu, An Zhang,	1908
		Xiang Wang, , and Xiangnan He. 2025b. Quantile ad-	1909
1857	Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu,	vantage estimation for entropy-safe reasoning. ArXiv	1910
1858	Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan	preprint arXiv:2509.22611.	1911
1859	Zhang, YK Li, Yang Wu, , et al. 2024. Deepseek-		
1860	math: Pushing the limits of mathematical reason-	Zhenhe Wu, Jian Yang, Jiaheng Liu, Xianjie Wu,	1912
1861	ing in open language models. ArXiv preprint	Changzai Pan, Jie Zhang, Yu Zhao, Shuangyong	1913
1862	arXiv:2402.03300.	Song, Yongxiang Li, , and Zhoujun Li.	1914
		2025c. Tabler1: Region-based reinforcement	1915
1863	Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin	learning for table understanding. ArXiv preprint	1916
1864	Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin	arXiv:2505.12415.	1917
1865	Lin, , and Chuan Wu. 2025. Hybridflow: A flexible		
1866	and efficient rlhf framework. In Proceedings of	Can Xie, Ruotong Pan, Xiangyu Wu, Yunfei Zhang,	1918
1867	the Twentieth European Conference on Computer	Jiayi Fu, Tingting Gao, , and Guorui Zhou. 2025.	1919
1868	Systems, pages 1279–1297.	Unlocking exploration in rlvr: Uncertainty-aware ad-	1920
		vantage shaping for deeper reasoning. ArXiv preprint	1921
1869	Yuda Song, Julia Kempe, , and Remi Munos. 2025.	arXiv:2510.10649.	1922
1870	Outcome-based exploration for llm reasoning. ArXiv		
1871	preprint arXiv:2509.06941.	An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao,	1923
		Bowen Yu, Chengpeng Li, Dayiheng Liu, Jian-	1924
1872	Hongze Tan and Jianfei Pan. 2025. Gtpo and grpo-s:	hong Tu, Jingren Zhou, Junyang Lin, , et al. 2024a.	1925
1873	Token and sequence-level reward shaping with policy	Qwen2.5-math technical report: Toward mathemat-	1926
1874	entropy. ArXiv preprint arXiv:2508.04349.	ical expert model via self-improvement. ArXiv	1927
		preprint arXiv:2409.12122.	1928
1875	Haozhe Wang, Qixin Xu, Che Liu, Junhong Wu,		
1876	Fangzhen Lin, , and Wenhu Chen. 2025a. Emergent	Shihui Yang, Chengfeng Dou, Peidong Guo, Kai	1929
1877	hierarchical reasoning in llms through reinforcement	Lu, Qiang Ju, Fei Deng, , and Rihui Xin. 2025a.	1930
1878	learning. ArXiv preprint arXiv:2509.03646.	Dcpo: Dynamic clipping policy optimization. ArXiv	1931
		preprint arXiv:2509.02333.	1932
1879	Hong Wang, Zhezhen Hao, Jian Luo, Chenxing Wei,		
1880	Yao Shu, Lei Liu, Qiang Lin, Hande Dong, , and	Wei Qin Yang, Jiawei Chen, Xin Xin, Sheng Zhou, Bin-	1933
1881	Jiawei Chen. 2025b. Scheduling your llm reinforce-	bin Hu, Yan Feng, Chun Chen, , and Can Wang.	1934
1882	ment learning with reasoning trees. ArXiv preprint	2024b. Psl: Rethinking and improving softmax	1935
1883	arXiv:2510.24832.	loss from pairwise perspective for recommendation.	1936
		advances in neural information processing systems,	1937
		37:120974–121006.	1938

1939
1940
1941
1942
1943
1944
1945

1946
1947
1948
1949
1950

1951
1952
1953
1954
1955

1956
1957
1958
1959
1960

1961
1962
1963
1964
1965

1966
1967
1968
1969
1970

1971
1972
1973

1974
1975
1976
1977
1978

1979
1980
1981
1982

1983
1984
1985
1986

1987
1988
1989
1990

1991
1992

Weiqin Yang, Jiawei Chen, Shengjia Zhang, Peng Wu, Yuegang Sun, Yan Feng, Chun Chen, , and Can Wang. 2025b. Breaking the top-k barrier: Advancing top-k ranking metrics optimization in recommender systems. In Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2, pages 3542–3552.

Weiqin Yang, Bohao Wang, Zhenxiang Xu, Jiawei Chen, Shengjia Zhang, Jingbang Chen, Canghong Jin, , and Can Wang. 2026. Bear: Towards beam-searchaware optimization for recommendation with large language models. ArXiv preprint arXiv:2601.22925.

Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, , et al. 2025. Dapo: An open-source llm reinforcement learning system at scale. ArXiv preprint arXiv:2503.14476.

Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, , and Junxian He. 2025. Simplerlzo: Investigating and taming zero reinforcement learning for open base models in the wild. ArXiv preprint arXiv:2503.18892.

Kaiyan Zhang, Yuxin Zuo, Bingxiang He, Youbang Sun, Runze Liu, Che Jiang, Yuchen Fan, Kai Tian, Guoli Jia, Pengfei Li, , et al. 2025a. A survey of reinforcement learning for large reasoning models. ArXiv preprint arXiv:2509.08827.

Kangning Zhang, Wenxiang Jiao, Kounianhua Du, Yuan Lu, Weiwen Liu, Weinan Zhang, , and Yong Yu. 2025b. Looptool: Closing the data-training loop for robust llm tool calls. ArXiv preprint arXiv:2511.09148.

Ruipeng Zhang, Ya-Chien Chang, , and Sicun Gao. 2025c. When maximum entropy misleads policy optimization. ArXiv preprint arXiv:2506.05615.

Shengjia Zhang, Weiqin Yang, Jiawei Chen, Peng Wu, Yuegang Sun, Gang Wang, Qihao Shi, , and Can Wang. 2026a. Talos: Optimizing top-k accuracy in recommender systems. ArXiv preprint arXiv:2601.19276.

Wenyuan Zhang, Xinghua Zhang, Haiyang Yu, Shuaiyi Nie, Bingli Wu, Juwei Yue, Tingwen Liu, , and Yongbin Li. 2026b. Expseek: Self-triggered experience seeking for web agents. Preprint, arXiv:2601.08605.

Xingjian Zhang, Siwei Wen, Wenjun Wu, , and Lei Huang. 2025d. Edge-grpo: Entropy-driven grpo with guided error correction for advantage diversity. ArXiv preprint arXiv:2507.21848.

Xinglang Zhang, Yunyao Zhang, ZeLiang Chen, Junqing Yu, Wei Yang, , and Zikai Song. 2026c. Logical phase transitions: Understanding collapse in llm logical reasoning. Preprint, arXiv:2601.02902.

Yunyao Zhang, Xinglang Zhang, Junxi Sheng, Wenbing Li, Junqing Yu, Yi-Ping Phoebe Chen, Wei

Yang, , and Zikai Song. 2026d. Semantic-aware logical reasoning via a semiotic framework. Preprint, arXiv:2509.24765.

Ziqi Zhao, Zhaochun Ren, Jiahong Zou, Liu Yang, Zhiwei Xu, Xuri Ge, Zhumin Chen, Xinyu Ma, Daiting Shi, Shuaiqiang Wang, , et al. 2026. Reinforced efficient reasoning via semantically diverse exploration. ArXiv preprint arXiv:2601.05053.

Yixiao Zhou, Yang Li, Dongzhou Cheng, Hehe Fan, , and Yu Cheng. 2026. Look inward to explore outward: Learning temperature policy from llm internal states via hierarchical rl. ArXiv preprint arXiv:2602.13035.

Xinyu Zhu, Mengzhou Xia, Zhepei Wei, Wei-Lin Chen, Danqi Chen, , and Yu Meng. 2025. The surprising effectiveness of negative reinforcement in llm reasoning. ArXiv preprint arXiv:2506.01347.

Brian D Ziebart, Andrew Maas, J Andrew Bagnell, , and Anind K Dey. 2008. Maximum entropy inverse reinforcement learning. In AAAI, volume 8, pages 1433–1438.

A Closed Forms for Section 4.4

Perturbing one logit of a k -way uniform distribution by δ and writing $S := e^\delta + k - 1$ gives $\pi_1 = e^\delta/S$ and $\pi_j = 1/S$ for $j \neq 1$, so

$$H(\delta) = \log S - \frac{\delta e^\delta}{S},$$

$$H'(\delta) = -\frac{(k-1)\delta e^\delta}{S^2},$$

whence $H'(0) = 0$ and $H''(0) = -(k-1)/k^2$. The sign of H' is $-\text{sign}(\delta)$, so any perturbation lowers the entropy, but only at second order. Table 13 contrasts this with a peaked distribution, where the first-order term does not vanish.

distribution	H_0	Ω	ΔH at $\delta=.1$
uniform, $k=10$	2.303	0	-4.7×10^{-4}
uniform, $k=50$	3.912	0	-1.0×10^{-4}
top .5 + 9 uniform	1.792	-0.549	-5.6×10^{-2}
top .9 + 9 uniform	0.545	-0.396	-3.8×10^{-2}

Table 13: First-order prediction versus actual entropy change. At a uniform branch point the first-order term is exactly zero and the realized damage is two orders of magnitude below that of a peaked state under the same perturbation.

B Implementation Notes

Sibling mask. For a group of G rollouts of length T we form $\text{eq}[i, j, t] = \mathbf{1}[y_t^{(i)} = y_t^{(j)}]$ and take a

cumulative product along t , so that $\text{prefixeq}[i, j, t]$ is true exactly when the two rollouts agree on $y_{\leq t}$. Shifting it one position right gives agreement on $y_{< t}$, which is the sibling relation Eq. (15) requires. Padded positions are treated as disagreement, so a length difference counts as a divergence.

Alignment. $\hat{H}_{\text{togo}}$ is read at position $t + 1$ and moved to position t ; at the final position the value is kept. Omitting this shift scores the state before the branch instead of the branch taken.

Memory. The RL path never materializes head logits. For each head the hidden states are processed in chunks of 4096 rows, mapped through the shared unembedding, reduced to an entropy, and released, so peak memory is $O(\text{chunk} \times V)$ rather than $O(\kappa \times B \times T \times V)$.

Equivalence test. A unit test asserts that with $\lambda = 0$ the token weights are bitwise identical to the base method’s on random inputs; the $\lambda = 0$ early return in the combination step is what makes this hold exactly rather than to floating-point tolerance.

C Pre-Registered Analysis

The analysis plan below was committed to the artifact repository before any run of the multi-seed campaign was launched, so that the statistics could not be chosen to fit the outcome. It is reproduced here in full; the file is `docs/preregistration_5seed.md` in the release.

Unit and primary endpoint. The unit of analysis is the seed, not the validation step. Each run is reduced to one number — its mean acc@32 on AIME24 over steps 40–110 — and arms are paired within a seed.

Contrasts, in the order they are reported. Primary: STEER-F – GRPO. Secondary, in this order: STEER-F – STEER, STEER-F – UNIFORM, STEER-F – PERMUTED. Secondary endpoints: maj@32 , uplift, and the six-benchmark mean. Effect sizes are reported with intervals rather than significance stars, and all five arms are reported whatever the outcome.

Stopping and exclusion. Every run goes to step 110; none is stopped early for looking good or bad. A run is excluded only for mechanical failure — out of memory, a crash, a corrupted checkpoint, or a step-0 validation score outside .030–.055, which

would indicate a different initial checkpoint. Exclusions are listed with their reason. A failed run is relaunched with the same seed, never a new one, and no seed is added after results are seen.

Checkpoint selection. Two rules are computed and both reported: the argmax over a run’s validation points, and a fixed step 110. The six-benchmark evaluation uses the argmax. Selecting a maximum over twelve noisy checkpoints is biased upward — in the single-seed data by $+.011$ to $+.017$ per arm — which is why the primary endpoint is a window mean rather than a maximum.

What would falsify the central claim. Stated in advance so it cannot be renegotiated afterwards. If STEER-F – UNIFORM straddles zero across seeds, the forecast’s *value* does not matter and the apparatus explains the effect; that is the claim of this paper and it would be refuted. If STEER-F – PERMUTED straddles zero, the pairing between a branch and its score does not matter, which the derivation says it must. If aggregate entropy orders the arms the same way accuracy does, the dissociation argument of Section 12.5 fails, since it depends on the two orderings coming apart. And if the mean token weight in the STEER-F arm departs from ≈ 0.999 , then A_H is not sum-zero over sibling sets as derived and the implementation is wrong rather than the theory.

Amendment, recorded. The pre-registration assigned three seeds to UNIFORM and PERMUTED and five to the other three arms. The campaign as executed runs all five arms at five seeds. The change raises the power of the two control contrasts and was made before any result of the campaign was seen; it is recorded here rather than silently absorbed.

Second amendment, added 2026-09-21, after this section’s own entropy condition was retired. Everything above was fixed before the campaign ran. This paragraph was not, and it is marked so rather than folded into the text above. It justifies nothing about the seeds already collected; it binds the seeds not yet in hand.

What it replaces. The third falsification bullet above put a condition on the aggregate entropy ordering. Section 12.5 retires it as a category error — it registers the aggregate as a signal about A_H , which Section 9.3 had already said it cannot be, an hour before this file was committed — and claims nothing from it in either direction. A first attempt

at a replacement, registered earlier on 2026-09-21, named `steerf/branch_entropy`; that registration is withdrawn here rather than deleted, because the monitor it named selects a fixed top decile of A_H rather than the support (Section 12.6).

Registered now, before the runs that can answer it. The quantity is `steerf/support_entropy` and its gap, the entropy split on $A_H \neq 0$ — the same predicate the correction itself applies, so the split covers the positions the method touches and no others. The primary contrast is STEER-F – PERMUTED, paired by seed, because that pair shares λ , $\lambda_{\min}$, the sampler and the magnitudes of A_H , and differs only in whether a branch is paired with its own score; STEER-F – UNIFORM is secondary. The claim under test is that a support-restricted entropy contrast near zero can accompany an accuracy contrast away from zero.

What would falsify it. If the support-restricted contrast tracks the accuracy contrast in sign and magnitude across seeds the way the apparatus contrasts in Table 8 do, then the split is the aggregate statistic in a smaller window and the claim fails.

What is still not instrumented, and we say so rather than substitute. Section 9.3’s own prediction is that the spread *across siblings* widens. Testing it needs the sibling index at the loss site, where only the per-token tensors are currently in scope, so no quantity we log tests it today. We do not offer the support-restricted split as that test: it measures a level, not a spread.

D Within-Run Stability

Earlier drafts of this work paired the arms at each of the eight validation steps in the converged window and reported a paired t over those eight points. That statistic answers a narrower question than the one in Section 12.3: it asks whether a gap is stable along a single trajectory, not whether it reappears when the run is repeated. Its n counts validation points, which are serially correlated within a run, so it is not an estimate of run-to-run uncertainty and cannot be read as one.

We report it here because it is informative about a different thing — whether a gap is an artifact of one lucky validation point — and because the single-seed version of this manuscript quoted it. Where the two disagree, the seed-level statistic in Section 12.8 governs.

The contrast between the two is the point. Ta-

ble 14 separates every A_H comparison at t_7 between +2.56 and +7.18, which reads as decisive and is not: the same gap sits at the same order as the run-to-run spread we measure in Section 12.8. A statistic can be this confident about a trajectory and say nothing about a re-run, which is exactly why we moved it out of the main text.

contrast	$\Delta\text{acc@32}$	paired t_7
STEER-F – GRPO	+0.0145	+3.08
STEER-F – STEER	+0.0165	+6.68
STEER-F – UNIFORM	+0.0182	+7.18
STEER-F – PERMUTED	+0.0125	+2.56

Table 14: Within-run stability over the eight converged-window validation points of seed 1. The points are consecutive validations of one run and are therefore serially correlated, so $\text{df} = 7$ is generous and these t values are not an estimate of run-to-run variation; see Section 12.8 for that. Majority-vote versions of the same four rows separate at +5.09 to +5.50.