

# Entropy Collapse Has a Direction: Correcting the Visitation Channel in Token-Level Entropy Control for RLVR

Anonymous ACL submission

## Abstract

Reinforcement learning with verifiable rewards (RLVR) reliably drives a policy’s entropy toward zero, and a growing family of methods intervenes by re-weighting the policy loss with a per-token estimate of how much each token’s update costs in entropy. We show that every such method estimates only half of the relevant quantity. The gradient of the *trajectory* entropy splits exactly into two terms: a *local* channel, which asks how the conditional distribution at a visited state flattens or sharpens, and a *visitation* channel, which asks how the probability of *reaching* that state changes. Existing token-level estimators capture the first and are structurally blind to the second. We show the blindness is worse than an omission: the local estimator is exactly zero at a uniform branch point, that zero is the correct first-order answer, and the visitation term’s prefactor is simultaneously maximal there — the two channels are *anti-correlated*, so no refinement of the local term can recover the missing signal. We then derive the missing channel into a per-token advantage  $A_H$ . Its reward-to-go is the remaining conditional entropy  $H_{\text{togo}}$ , and the sibling-prefix baseline that keeps the estimator unbiased is supplied by a tree-structured rollout in which several samples share a prefix and diverge at scheduled depths.  $A_H$  sums to zero within a sibling set by construction, so the correction cannot change the average amount of damping — it can only decide which of two siblings is damped. A five-arm design, which runs the full pipeline twice with  $A_H$ ’s value destroyed in two different ways, measures that consequence: switching the damping apparatus on or off moves the aggregate entropy ( $t = -5.87$  and  $+13.09$ ), while changing what  $A_H$  carries with the apparatus held fixed does not move it detectably ( $t = -1.42$ ), and accuracy separates the two groups the other way round. The aggregate and the correction are dissociated, as the construction says they must be. We forecast  $H_{\text{togo}}$  with parallel multi-token-

prediction heads, state the four approximations that estimation requires and what each discards, and report our own measurement that the forecast comes out close to a scalar multiple of a quantity the trainer already computes — so the forecaster is one estimator of this reward-to-go, not a part the decomposition depends on. On accuracy the evidence is weaker and we say so: the gain over GRPO ( $+0.0055$  over three seeds) sits inside GRPO’s seed-to-seed spread and is close to zero on the two seeds where equal steps mean equal compute, and the pre-registered test of  $A_H$ ’s value against a value-free control is met as a refutation. Run under the base method’s own protocol and set beside its published table, two runs of the method average  $-$  on six benchmarks against a reported 38.2 — an uncontrolled comparison, which we report as such.

## 1 Introduction

Policy-gradient fine-tuning with verifiable rewards sharpens a language model’s output distribution. Measured as the mean per-token entropy of the policy, this sharpening is fast, monotone, and, past a point, coincident with the end of useful learning. The standard reading is that the collapse is the problem, and the standard remedy is to spend part of the update budget slowing it down: an entropy bonus, a KL anchor, a clip on the tokens that lose the most entropy, or — in the most recent line of work — a per-token re-weighting of the policy loss by a first-order estimate of the entropy each token’s update destroys.

We argue that the accounting behind that remedy is incomplete, and that the incompleteness is structural rather than a matter of estimator quality.

Write  $\mathcal{H}(\pi_\theta \mid x)$  for the entropy of the distribution the policy induces over whole responses to a prompt  $x$ . The autoregressive factorization turns it into an expectation of a sum of per-token conditional entropies, and the gradient of that expression splits, by the product rule, into exactly two

terms. The first holds the visited states fixed and asks how the conditional distribution at each one changes: this is what a per-token entropy estimate measures. The second holds the per-state entropies fixed and asks how the probability of *visiting* each state changes. An update that concentrates mass on one continuation makes the states downstream of that continuation more likely and the states downstream of its abandoned siblings less likely; the diversity those siblings carried leaves the trajectory distribution even though no conditional distribution anywhere became more peaked.

Our first contribution is to show that the local channel is not merely blind to the second term but *anti-correlated* with it. The local estimator factors into the advantage, a saturation factor, and the deviation of the sampled token’s surprisal from the mean surprisal — and that last factor has mean zero by the definition of entropy. At a branch point where several continuations are close to equiprobable it vanishes identically. The vanishing is not an estimator failure: a uniform distribution maximizes entropy, so the first-order change is genuinely zero and the damage is second order (Section 4.4). Meanwhile the visitation term’s prefactor is  $1 - \pi_\theta(y_t | s_t)$ , which is maximal precisely at those flat states. The local channel therefore spends its budget where the conditional entropy is fragile — rare tokens sampled from confident distributions — and spends nothing where the trajectory entropy is most at risk.

Our second contribution is an estimator for the missing channel that requires no additional rollouts. Eliminating the non-causal cross terms turns the visitation term into a REINFORCE estimator whose reward-to-go is the sum of the *remaining* conditional entropies,  $H_{\text{togo}}$ ; subtracting any function of the state leaves it unbiased. We estimate  $H_{\text{togo}}$  with parallel multi-token-prediction (MTP) heads read off the hidden state the policy already computes, and we obtain the state-only baseline from rollouts that share a prefix. Plain group-relative rollouts almost never share prefixes, so we schedule the sharing: a tree-structured sampler branches a single root at fixed depths, and the resulting sibling sets give the baseline an actual support (Section 9).

The resulting per-token score  $A_H$  is a deviation from a sibling mean and therefore sums to zero over each sibling set. This has a consequence we state up front because it determines what evidence can count: the method *cannot* raise the average

weight at branch points, only spread the weights of siblings apart. “Total entropy went up” is therefore not a prediction of this method, and an experiment looking for it is testing a different hypothesis (Section 9.3).

We make the following contributions.

- An exact two-channel decomposition of the trajectory-entropy gradient, and an argument that the channel captured by current token-level entropy-control methods is anti-correlated with the channel they miss (Sections 4–4.4).
- A derivation of the missing channel into a per-token advantage  $A_H$  with no approximation up to the point of estimating  $H_{\text{togo}}$ , together with an explicit ledger of the approximations that estimation requires, the sign of each one’s error, and the mechanism that absorbs it (Sections 5–10).
- A tree-structured rollout that supplies the sibling-prefix baseline, and a bound showing what fraction of positions can carry a non-zero  $A_H$  under *any* sampler (Section 9).
- A five-arm experiment that separates the contribution of  $A_H$ ’s *value* from that of the apparatus delivering it, by running the full pipeline twice with the value destroyed in two different ways, and the seed-paired contrasts it yields (Sections 11 and 12.4).
- The dissociation that construction predicts, measured rather than argued: the contrasts that switch the damping apparatus move the aggregate entropy ( $t = -5.87$  and  $+13.09$ ) while the contrasts that change only what  $A_H$  carries do not move it detectably ( $t = -1.42$ ), and accuracy separates the same two groups in the opposite direction. The structural claim, not the effect size, is what this campaign establishes (Section 12.3).

We are equally explicit about what the campaign does not establish. Over three seeds STEER-F’s converged AIME24 accuracy is above GRPO’s by  $+0.0055$ , but that gap is smaller than GRPO’s own spread across seeds, it is carried mostly by the one seed on which the treatment also had the larger compute budget (Section 12.11), and the pre-registered test of whether  $A_H$ ’s *value* matters — STEER-F against the same apparatus with that

value discarded — straddles zero across seeds, which the registration defined as a refutation (Section 12.4). We report it as one. This is a paper about where token-level entropy control acts and what an aggregate entropy curve can and cannot tell us about it, not a paper claiming an accuracy gain.

## 2 Related Work

**Entropy collapse in RLVR.** Entropy decay under policy-gradient fine-tuning is well documented (Cui et al., 2025), and a range of interventions add an explicit entropy term to the objective, in the tradition of maximum-entropy reinforcement learning (Ziebart et al., 2008; Haarnoja et al., 2018), or constrain the update so that high-entropy positions are preserved. Recent analyses report that naive entropy regularization does not prevent the collapse and that continuing to train under it degrades accuracy (Anonymous, 2025), which is consistent with our reading: the regularizer acts on the local channel, and the local channel is not where trajectory diversity is being lost.

**Assumption.** Both Hao et al. (2026) and this work rely on one structural assumption, which we restate so that our approximation ledger can refer to it.

**Assumption 1** (parameter-independent softmax; Hao et al., 2026). For any state  $s$ , each action  $a$  carries an independent logit  $z_{s,a}(\theta)$ , so that a gradient step on the sampled token does not substantially move the logits of the other tokens at that state.

It is false for a network with shared parameters. We inherit it rather than introduce it, but our construction adds a second estimator that rests on it.

**Token-level entropy control.** The method we build on re-weights the clipped surrogate loss by a per-token first-order estimate of the entropy change that token’s update induces (Hao et al., 2026). That work observes that earlier interventions are “indirect” and that their effect is therefore limited and can fail. Our decomposition supplies a mechanism for that observation: the estimate is a first-order estimate of the local channel only, and it is exactly zero at the states where the visitation channel is largest.

**Which tokens carry the gradient.** Wang et al. (2025) show that restricting policy-gradient updates to a small minority of high-entropy “forking” tokens matches full-gradient training. That result establishes that the identity of the tokens, not the aggregate entropy, is what matters, and it motivates a per-token treatment; it does not, however, distinguish the two channels, since forking tokens are identified by their local entropy. Our analysis predicts that a local criterion will select a systematically different set from the one the visitation channel would select.

**Local stochasticity without global diversity.** Wu et al. (2025) report that token-level entropy can remain high while the diversity of final answers falls. This is precisely the signature our decomposition predicts when the visitation channel is unmanaged: per-state entropies are preserved — the local channel is doing its job — while the trajectory distribution concentrates.

**Multi-token prediction.** We use Medusa-style parallel prediction heads (Cai et al., 2024) not for their intended purpose (speculative decoding) but as an entropy forecaster: we need the entropy of the distribution over a token  $k$  steps ahead, not a good sample of that token. Section 6 shows that this choice of estimator has a signed, monotone error which we characterize and correct.

## 3 Preliminaries

**Notation.** A prompt is  $x$ , a response is  $y = y_{1:T}$ , and the state before emitting token  $t$  is  $s_t = (x, y_{<t})$ . We write  $H(\pi_\theta(\cdot \mid s_t))$  for the entropy of the conditional distribution at a single position (lowercase, per-token) and  $\mathcal{H}(\pi_\theta \mid x)$  for the entropy of the distribution over whole responses (uppercase, per-trajectory). The surprisal of a token  $a$  is  $I_a := -\log \pi_\theta(a \mid s_t)$ , so that  $H = \mathbb{E}_{a \sim \pi_\theta}[I_a]$ .

**Group-relative policy optimization.** We optimize with GRPO. For each prompt,  $n$  responses are sampled and scored, and each response receives the group-standardized advantage  $A_i = (r_i - \text{mean}_j r_j) / \text{std}_j r_j$ , broadcast to every token of that response. Replacing a learned value function with the group mean is what makes GRPO cheaper than PPO; it also means that a group whose members all receive the same reward contributes no gradient.

**Token-level entropy control.** Hao et al. (2026) leave the advantage untouched and instead multiply each token’s contribution to the clipped surrogate by a weight  $\alpha_t \in [\lambda_{\min}, 1]$  derived from a first-order estimate of the entropy change that token’s update induces. Their Theorem 1 gives that estimate in closed form: under a parameter-independent softmax (their Assumption 1; our 2), the change in the conditional entropy at a *fixed* state  $s$  after one GRPO step is, to first order,

$$\Omega(s) = -\frac{\eta}{L} \mathbb{E}_a \left[ \frac{\mathbb{I}_{\text{clip}}(s, a) A(s, a)}{\pi_{\text{old}}(a | s)} \cdot \pi_{\theta}(a | s) (1 - \pi_{\theta}(a | s)) \cdot (\log \pi_{\theta}(a | s) + H) \right], \quad (1)$$

where  $\eta$  is the learning rate,  $L$  the total decoded length in the update, and  $\mathbb{I}_{\text{clip}}$  the PPO clipping indicator. The implementation evaluates the bracket at the sampled token rather than in expectation. The policy loss becomes  $\sum_u \alpha_u \cdot \text{PPO-clip}(A_u, r_u)$ , with

$$\alpha_u = \exp\left(-\alpha \frac{|\Omega_u|}{\max_v |\Omega_v|}\right), \quad \lambda_{\min} = e^{-\alpha}, \quad (2)$$

a batch-normalized exponential map into  $[\lambda_{\min}, 1]$ , so the weight can only attenuate. The factor  $1/\pi_{\text{old}}$  in Eq. (1) will matter later: it makes  $\Omega$  heavy-tailed, which is why the two channels must be normalized before they are combined (Section 8).

## 4 Two Channels

### 4.1 Step 1: the chain rule

The object we would like to protect is the entropy of the distribution the policy induces over *whole responses*,

$$\mathcal{H}(\pi_{\theta} | x) := -\sum_y \pi_{\theta}(y | x) \log \pi_{\theta}(y | x), \quad (3)$$

a sum over an exponentially large set. The autoregressive factorization  $\pi_{\theta}(y | x) = \prod_t \pi_{\theta}(y_t | s_t)$  turns the logarithm into a sum and lets us exchange the order of summation:

$$\begin{aligned} \mathcal{H} &= -\sum_y \pi_{\theta}(y | x) \sum_{t=1}^T \log \pi_{\theta}(y_t | s_t) \\ &= \sum_{t=1}^T \mathbb{E}_{y \sim \pi_{\theta}} [-\log \pi_{\theta}(y_t | s_t)]. \end{aligned} \quad (4)$$

Each summand is still an expectation over whole responses, but the integrand depends on  $y$  only

through  $y_{\leq t}$ . Conditioning on the prefix and applying the tower rule replaces the inner expectation with a per-position entropy,

$$\begin{aligned} \mathbb{E}_y [-\log \pi_{\theta}(y_t | s_t)] \\ = \mathbb{E}_{y_{<t}} \left[ \underbrace{\mathbb{E}_{y_t \sim \pi_{\theta}(\cdot | s_t)} [-\log \pi_{\theta}(y_t | s_t)]}_{= H(\pi_{\theta}(\cdot | s_t))} \right], \end{aligned} \quad (5)$$

because the inner expectation is, by definition, the entropy of the conditional distribution at  $s_t$ . Collecting the terms back under a single expectation gives

$$\mathcal{H}(\pi_{\theta} | x) = \mathbb{E}_{y \sim \pi_{\theta}} \left[ \sum_{t=1}^T H(\pi_{\theta}(\cdot | s_t)) \right]. \quad (5)$$

Equation (5) is an identity, not an approximation. Its left-hand side is what we want to protect; its right-hand side is what instrumentation can measure, since  $H(\pi_{\theta}(\cdot | s_t))$  is read directly off the logits at each position.

The entire argument of this paper follows from one feature of Eq. (5) that is easy to pass over: *the measure  $\mathbb{E}_{y \sim \pi_{\theta}}$  depends on  $\theta$  as well*. The right-hand side is not a sum of  $\theta$ -dependent terms over a fixed set of positions — it is a  $\theta$ -dependent average of them over a  $\theta$ -dependent set of trajectories.

### 4.2 Step 2: the product rule

Write the per-trajectory total

$$f_{\theta}(y) := \sum_{t=1}^T H(\pi_{\theta}(\cdot | s_t(y))), \quad (6)$$

so that Eq. (5) reads  $\mathcal{H} = \sum_y \pi_{\theta}(y | x) f_{\theta}(y)$ . Both factors of every summand carry  $\theta$ , so differentiating requires the product rule:

$$\begin{aligned} \nabla_{\theta} \mathcal{H} &= \sum_y [\nabla_{\theta} \pi_{\theta}(y | x)] f_{\theta}(y) \\ &\quad + \sum_y \pi_{\theta}(y | x) \nabla_{\theta} f_{\theta}(y). \end{aligned} \quad (7)$$

The second sum is already an expectation. The first is not, because  $\nabla_{\theta} \pi_{\theta}$  is not a probability distribution; the score-function identity  $\nabla_{\theta} \pi_{\theta} = \pi_{\theta} \nabla_{\theta} \log \pi_{\theta}$  restores it to one. Substituting and naming the two terms,

$$\begin{aligned} \nabla_{\theta} \mathcal{H} &= \underbrace{\mathbb{E}[\nabla_{\theta} f_{\theta}(y)]}_{\text{(L) local}} \\ &\quad + \underbrace{\mathbb{E}[f_{\theta}(y) \nabla_{\theta} \log \pi_{\theta}(y | x)]}_{\text{(V) visitation}}. \end{aligned} \quad (8)$$

Equations (3)–(8) contain no approximation of any kind. The split into two channels is forced by the product rule, not chosen.

### 4.3 Reading the two terms

**(L), the local channel**, holds the visited states fixed and asks how the conditional distribution at each one changes. Equation (1) is a first-order estimate of exactly this.

**(V), the visitation channel**, holds the per-state entropies fixed and asks how the probability of reaching each state changes. An update that concentrates mass on one continuation makes the subtree below it more likely and the sub-trees below its siblings less likely. Nothing in Eq. (1) can see this:  $\Omega$  reads the conditional distribution at a position and never reads the probability of arriving there.

The two terms can have opposite signs. A method that drives (L) to zero has not thereby bounded  $\nabla_\theta \mathcal{H}$ .

### 4.4 Why the local channel cannot be repaired

A natural objection is that the two channels might coincide in practice: branch points are positions where the policy is uncertain, so they are high-entropy positions, so  $\Omega$  should be large there and the local term should protect them automatically. The premise is correct — in warm-up rollouts on our model, branch diversity correlates with local entropy at Spearman  $\rho = 0.65$ , and with the entropy-to-go at  $\rho = 0.82$ . The inference from it fails, in three steps.

**Step 1:  $\Omega$  reads a deviation, not a level.** Collecting the constants of Eq. (1) into  $c := \eta \mathbb{I}_{\text{clip}} A / (L \pi_{\text{old}})$  and rewriting in terms of surprisal gives, for the contribution of a single action  $a$ ,

$$\Omega \propto c \pi_a (1 - \pi_a) (H - I_a), \quad (9)$$

and  $H = \mathbb{E}_{a \sim \pi} [I_a]$  by definition. The factor  $H - I_a$  is therefore a mean-zero fluctuation: what sets the magnitude of  $|\Omega|$  is the *spread* of surprisal, not the *height* of the entropy. Nothing in  $c$  depends on the shape of  $\pi(\cdot | s)$ , so the argument below is unaffected by the clipping indicator or the importance weight.

**Step 2: at a uniform branch point  $\Omega \equiv 0$ .** If the distribution is uniform over  $k$  candidates then  $I_a = \log k = H$  for every  $a$ , so  $\Omega$  vanishes identically. The positions that are most branch-like — several continuations at comparable probability —

are exactly where the local term contributes nothing.

**Step 3: and that zero is the correct answer.**

This is the step that closes the argument, and it follows from Hao et al.’s own derivation. Differentiating the conditional entropy with respect to a single logit gives (Hao et al., 2026, Eq. 20)

$$\frac{\partial H(\pi_\theta | s)}{\partial z_{s,a}} = -\pi_\theta(a | s) (\log \pi_\theta(a | s) + H(\pi_\theta | s)), \quad (10)$$

and at a uniform distribution  $\log \pi_\theta(a | s) = -\log k = -H$  makes the bracket vanish for every  $a$ . So  $\partial H / \partial z_a = 0$  identically:  $\Omega$  is not wrong, it is right. Perturbing one logit of a uniform distribution by  $\delta$  and writing  $S := e^\delta + k - 1$  gives the closed form

$$H(\delta) = \log S - \frac{\delta e^\delta}{S}, \quad H''(0) = -\frac{k-1}{k^2}, \quad (11)$$

with  $H'(\delta)$  of sign  $-\text{sign}(\delta)$ : any push in any direction lowers the entropy, but the loss is entirely second order. For a 10-way uniform distribution and  $\delta = 0.1$  the actual entropy loss is  $4.7 \times 10^{-4}$ , roughly two orders of magnitude below that of a peaked distribution under the same perturbation. *As far as local entropy is concerned, uniform branch points are genuinely safe*, and a local estimator that allocates them no protection is allocating correctly for its own objective. No refinement of the local term recovers the missing signal, because the missing signal is not local.

**The two channels are anti-correlated.** The visitation term’s prefactor, derived in Section 7, is  $1 - \pi_a$ , which approaches 1 as the distribution flattens — and the probability that a group of  $n$  samples actually diverges at a position approaches 1 in the same limit. Table 1 evaluates both channels on a family of distributions that interpolates between confident and flat.  $|\Omega|$  is dominated by positions where the model is fairly confident but a tail token was sampled; at those positions siblings diverge only about half the time. Where divergence is certain, the local term is identically zero and the visitation prefactor is maximal.

The statement we can therefore make is stronger than “the local term is incomplete”: *the local term is anti-correlated with the term it omits*, and the omission cannot be fixed within the local channel.

| distribution         | $H$  | $P(\text{split})$ | $\mathbb{E} \Omega / A $ | $\mathbb{E}[1-\pi_a]$ |
|----------------------|------|-------------------|--------------------------|-----------------------|
| $q=.99$              | 0.10 | 0.08              | 0.092                    | 0.020                 |
| $q=.9, \text{ tail}$ | 1.25 | 0.57              | <b>1.129</b>             | 0.190                 |
| $q=.6, m=3$          | 1.11 | 0.98              | 0.457                    | 0.587                 |
| unif. $k=10$         | 2.30 | <b>1.00</b>       | <b>0.000</b>             | <b>0.900</b>          |
| unif. $k=100$        | 4.61 | <b>1.00</b>       | <b>0.000</b>             | <b>0.990</b>          |

Table 1: The local channel and the visitation channel are ordered oppositely. Distribution family: top token has mass  $q$ , the remaining  $1 - q$  is spread uniformly over  $m$  tokens.  $P(\text{split})$  is the probability that  $n=8$  samples do not all agree.  $\mathbb{E}|\Omega|$  is taken over the token actually sampled at that position.

## 5 Estimating the Visitation Channel

### 5.1 Causal elimination gives a reward-to-go

The visitation term of Eq. (8) is not yet in a usable form: the factor  $f_\theta(y)$  multiplying the score is a property of the *whole* trajectory, so a naive estimator would credit every token with entropy that appeared before it. Causality removes those terms exactly.

Expanding both factors —  $f_\theta$  by Eq. (6) and the score by  $\nabla_\theta \log \pi_\theta(y \mid x) = \sum_u \nabla_\theta \log \pi_\theta(y_u \mid s_u)$  — turns (V) into a double sum:

$$(V) = \sum_t \sum_u \mathbb{E}[H(\pi_\theta(\cdot \mid s_t)) \cdot \nabla_\theta \log \pi_\theta(y_u \mid s_u)]. \quad (12)$$

Consider a term with  $u \geq t$ . The entropy  $H(\pi_\theta(\cdot \mid s_t))$  depends on  $y$  only through  $y_{<t}$ , and  $y_{<t} \subseteq y_{<u}$ , so it is measurable with respect to  $y_{<u}$  and may be pulled out of a conditional expectation given  $y_{<u}$ . What remains is the expected score at a single position, which vanishes for any distribution:

$$\begin{aligned} & \mathbb{E}[\nabla_\theta \log \pi_\theta(y_u \mid s_u) \mid y_{<u}] \\ &= \sum_a \pi_\theta(a \mid s_u) \frac{\nabla_\theta \pi_\theta(a \mid s_u)}{\pi_\theta(a \mid s_u)} \\ &= \nabla_\theta \sum_a \pi_\theta(a \mid s_u) = \nabla_\theta 1 = 0. \end{aligned} \quad (13)$$

By the tower rule every  $(t, u)$  term with  $u \geq t$  is therefore *exactly* zero — not small, zero. This is the same identity that licenses a baseline in policy gradients, used here to enforce causality rather than to reduce variance.

Only  $u < t$  survives. Swapping the order of summation collects, for each token  $u$ , the entropies

at all later positions:

$$(V) = \mathbb{E}\left[\sum_u \nabla_\theta \log \pi_\theta(y_u \mid s_u) H_{\text{togo}}(s_u \oplus y_u)\right], \quad (14)$$

with  $H_{\text{togo}}(s_u \oplus y_u) := \sum_{t>u} H(\pi_\theta(\cdot \mid s_t))$ .

We stress that  $H_{\text{togo}}$  was not posited. Equation (14) is a REINFORCE estimator whose reward is “the conditional entropy remaining after this token”, and  $H_{\text{togo}}$  is that reward-to-go.

### 5.2 The sibling-prefix baseline

Any function  $b(s_u)$  of the state alone leaves Eq. (14) unbiased, by the same argument. Subtracting one gives the per-token advantage

$$A_H(s_u, y_u) := H_{\text{togo}}(s_u \oplus y_u) - \bar{H}_{\text{togo}}(s_u), \quad (15)$$

and up to this point *no approximation has been made*.

The baseline we use is the mean of  $H_{\text{togo}}$  over the rollouts in the same prompt group whose prefixes  $y_{<u}$  agree exactly. Those rollouts share the state  $s_u$  by definition, so their mean is a function of  $s_u$  alone and the baseline is admissible. The condition is agreement on  $y_{<u}$  and not on  $y_{\leq u}$ : the rollouts must diverge *at*  $u$  for the difference to be a branch score.

Two properties follow immediately and matter later. First, the implementation includes a rollout in its own sibling set, so  $b$  depends weakly on  $y_u$  and the estimator is not strictly leave-one-out; the resulting bias is  $O(1/m)$  in the number  $m$  of surviving siblings, the same property GRPO’s group-mean advantage has. Second, when  $m = 1$  the baseline equals the rollout’s own value and  $A_H \equiv 0$ , so a position with no siblings is silent rather than noisy.

## 6 Four Approximations

Equation (15) is exact but not computable:  $H_{\text{togo}}$  sums conditional entropies at states the update has not visited. Three approximations make it computable and a fourth bounds its influence.

### 6.1 A1: horizon truncation and discounting

We truncate the sum to  $\kappa$  steps and discount:

$$H_{\text{togo}}(s_u \oplus y_u) \approx \sum_{k=1}^{\kappa} \gamma_H^k H(\pi(\cdot \mid s_{u+k})). \quad (16)$$

The discount serves two purposes: distant terms carry more forecast error (A2) and therefore deserve less trust, and the sum stays bounded independently of  $\kappa$ .

What this discards — the tail  $\sum_{k>\kappa}$  and a fraction of the retained terms — is largely harmless for a reason specific to our use. Equation (15) consumes only the *difference* between siblings. Any bias common to rollouts sharing  $s_u$  cancels in that difference, and truncation bias is common to first order.

## 6.2 A2: forecasting the unrealized future

Even Eq. (16) requires the future states  $s_{u+k}$ . At update time they do not exist: rollouts came from  $\pi_{\text{old}}$ , each state has exactly one sampled continuation, and re-rolling is prohibitive. We replace the realized future with a parallel forecast from the current hidden state  $h_u$ :

$$H(\pi(\cdot | s_{u+k})) \rightarrow H(p_{\text{MTP}}^{(k)}(\cdot | h_u)). \quad (17)$$

Head  $k$  is a two-layer residual MLP applied to  $h_u$ , followed by the policy’s own unembedding; the heads own no vocabulary parameters. Their last layer is zero-initialized, so before any training every head reproduces the policy’s own next-token distribution and the forecast starts from a sane value. In the RL path the logits are materialized in chunks and reduced to an entropy immediately, so the  $[\kappa, B, T, V]$  tensor never exists.

**The error of this substitution has a known sign.** Conditioning on the intervening tokens is exactly what the forecast omits. Fixing  $S_u$ , the  $k$ -th term of Eq. (16) has expectation  $H(Y_{u+k} | Y_{u+1:u+k-1}, S_u)$ , whereas the MTP head models the *marginal*  $p(Y_{u+k} | S_u)$ . Since conditioning reduces entropy,

$$\begin{aligned} H(Y_{u+k} | S_u) - \mathbb{E}[H(\pi(\cdot | S_{u+k-1}))] \\ = I(Y_{u+k}; Y_{u+1:u+k-1} | S_u) \\ \geq 0. \end{aligned} \quad (18)$$

The forecast *always over-estimates*, and the excess is precisely the mutual information with the intervening tokens, which grows monotonically with  $k$ .

Equation (18) is not merely a bound: it is the reason the marginal is the quantity we want. It is the fan-out of the continuation — how many ways the next  $k$  tokens can unfold — rather than the residual uncertainty along the one path that happened to be sampled.

## 6.3 A3: per-head affine calibration

Equation (18) predicts a per-head, monotone inflation, so we absorb it with a per-head affine correction fitted by least squares against oracle entropies on warm-up rollouts:

$$\hat{H}_k = a_k H(p_{\text{MTP}}^{(k)}(\cdot | h_u)) + b_k. \quad (19)$$

The fitted coefficients confirm the prediction. On our model the first head has  $a_1 = 1.03$  — a one-step forecast has no intervening tokens, so the mutual information in Eq. (18) is zero and the approximation is exact — while  $a_k$  drops to 0.10–0.19 for  $k \geq 2$ . Calibration switches the unreliable heads off on its own.

Combining A1–A3 gives the forecast that is actually computed:

$$\begin{aligned} \hat{H}_{\text{togo}}(s_u \oplus y_u) = \\ \max\left(0, \sum_{k=1}^{\kappa} \gamma_H^k (a_k H_k(h_u) + b_k)\right), \end{aligned} \quad (20)$$

the clamp encoding the fact that an entropy cannot be negative. With  $\kappa = 2$  and  $\gamma_H = 0.7$  this instantiates to  $\hat{H}_{\text{togo}} = 0.72 H_1 + 0.06 H_2 + 0.11$ : the second head contributes about 8%, so the effective horizon is close to one step. We regard this as a consequence of Eq. (18) rather than a tuning artifact, and note it as a limitation — a forecaster whose error did not grow with  $k$  would leave a longer horizon usable.

**Alignment.** One detail decides whether the signal is a branch score at all. The hidden state at position  $u$  is conditioned on  $s_u$  and predicts  $y_u$ ; it does not know  $y_u$ . The quantity we need,  $\hat{H}_{\text{togo}}(s_u \oplus y_u)$ , must therefore be read from position  $u + 1$  and shifted back by one. Omitting the shift measures the value of the state *before* the branch rather than the value of the branch taken.

## 7 From Gradient to Token Weights

Equation (15) is a term to add to a gradient, but methods in this family do not add gradients — they re-weight a loss. What is needed is a first-order prediction of how much one update, acting through token  $u$ , changes the trajectory entropy.

**How far one step moves a logit.** We reuse the logit-space update of Hao et al. (2026, Eq. 23) rather than rederiving it. With  $w_u = \text{clip}(r_u, 1 - \epsilon_{\text{lo}}, 1 + \epsilon_{\text{hi}})A_u$  the coefficient the policy loss

actually applies, Assumption 2 makes a single step move only the sampled token’s logit,  $\Delta z_{u,a} = \eta w_u \delta_{a,y_u} + O(\eta^2)$ . The softmax Jacobian  $\partial \log \pi_a / \partial z_a = 1 - \pi_a$  then gives

$$\widehat{\Delta \log \pi_u} = \eta w_u (1 - \pi_u), \quad (21)$$

which vanishes as  $\pi_u \rightarrow 1$ : a token that already holds the mass cannot be given more. This is the prefactor whose anti-correlation with  $|\Omega|$  Table 1 records.

**The visitation-channel prediction.** Because  $\nabla_{\theta} \log \pi_u \cdot \Delta \theta = \Delta \log \pi_u$ , substituting into Eq. (15) gives the per-token quantity

$$\text{visit}_u := \widehat{\Delta \log \pi_u} \cdot \text{clip}(A_H(s_u, y_u), -c, +c). \quad (22)$$

**A4** is the clip, and it is applied to  $A_H$  *before* the product rather than after, so that one token’s influence is bounded by  $|\eta w_u| \cdot c$ . Clipping the product instead would let forecast error be amplified without bound at tokens with large  $w_u$ .

**Sign quadrants.** Equation (22) defines the collapse channel precisely. When  $\widehat{\Delta \log \pi_u} > 0$  and  $A_H > 0$ , mass moves toward a branch richer than its siblings and trajectory entropy rises. When  $\widehat{\Delta \log \pi_u} > 0$  and  $A_H < 0$ , mass moves into a dead end and the diversity carried by the abandoned siblings is destroyed — this is the collapse path that (L) is structurally unable to see. Negative  $\widehat{\Delta \log \pi_u}$  reverses both readings.

## 8 Combining the Channels

The visitation term could be added to  $\Omega$  before the weight map, but measured on real rollouts that map destroys the signal:  $\Omega$  carries a factor  $1/\pi_{\text{old}}$ , so  $\max |\Omega|/\text{median}|\Omega|$  reaches  $3 \times 10^4$  and the band is spent on one or two outliers while the bulk lands on a point. We therefore apply the correction *after* the map, as a bounded adjustment of the weights the base method has already produced.

Let  $\mathcal{S} = \{u : \text{visit}_u \neq 0\}$  be the support — by Eq. (24), the branch points — and let  $b = \lambda_{\max} - \lambda_{\min}$  be the width of the weight band. For  $u \in \mathcal{S}$ ,

$$\alpha_u^f = \text{clamp}\left(\alpha_u + \lambda b \tanh(\text{visit}_u/\rho), \lambda_{\min}, \lambda_{\max}\right), \quad (23)$$

with  $\rho = \text{rms}\{\text{visit}_v : v \in \mathcal{S}\}$ ;  $\alpha_u^f = \alpha_u$  elsewhere.

Three properties matter. The tanh is odd, monotone and bounded, so it preserves the sign of the visitation term, caps one token’s movement at  $\lambda b$ , and does not depend on the absolute scale of visit. The RMS is taken over  $\mathcal{S}$  alone rather than over the whole response: normalizing by the whole response would shrink the correction in proportion to how sparse branch points are, which is a property of the sampler and not of the signal. And at  $\lambda = 0$  the correction vanishes identically, so the pipeline is bitwise identical to the base method — a regression test enforces this.

Because the map runs first and the correction second, the step size  $\eta$  of Eq. (21) enters only through  $\text{visit}_u/\rho$ , where it cancels. We fix  $\eta = 1$  and sweep  $\lambda$  alone.

**The support is part of the treatment.**  $\mathcal{S}$  is defined by  $\text{visit}_u \neq 0$  rather than by a threshold, deliberately: it selects both where the correction lands and, through  $\rho$ , how large every correction is, so a tighter test would change what is trained rather than what is logged. One consequence is that  $\mathcal{S}$  admits float32 round-off —  $A_H$  is a deviation from a mean over token-identical siblings, and  $\sum_j g_j/m$  does not reproduce  $g_i$  bitwise, leaving  $|A_H| \sim 6 \times 10^{-8}$  where it is interpretively zero. Those residuals are  $\sim 10^{-6}$  relative to a real branch score and cannot move a weight, but they are counted in the support statistics of Section 12.6, which should be read as an upper bound on the branch-point fraction.

## 9 Tree-Structured Rollouts

### 9.1 Why plain rollouts cannot support $A_H$

$\hat{H}_{\text{togo}}$  is a deterministic function of the causal prefix, so two rollouts sharing  $y_{\leq u}$  produce identical forecasts. Combined with Eq. (15) this gives a sharp characterization:

$$A_H[i, u] \neq 0 \iff \text{siblings share } y_{<u} \text{ but diverge at } y_u, \quad (24)$$

that is,  $u$  is a branch point of the group. Under independent sampling from a sharpened policy, group members diverge at their first token and never share a prefix again, so  $m = 1$  almost everywhere and  $A_H \equiv 0$ : the visitation channel is not merely weak, it is identically absent.

We therefore schedule the sharing. A tree sampler expands a single root and cuts it at fixed depths  $d_1 < d_2 < d_3$  with branching factor 2, yielding

$n = 8$  leaves whose sibling structure is known in advance. The cuts create the only positions at which Eq. (24) can fire.

## 9.2 A ceiling on the support

It is tempting to treat the fraction of positions with  $A_H \neq 0$  as a quantity to be maximized. It is not, and the reason is combinatorial rather than empirical. Two quantities behave very differently under the tree:

- the fraction of positions where *some* other rollout shares  $y_{<u}$  is moved by the cuts, and equals  $d_L/T$  by construction;
- the fraction where  $A_H \neq 0$  is *not*, because by Eq. (24) it counts branch points, and  $n$  sequences form a trie with at most  $n - 1$  internal branching nodes.

Since the one-position alignment shift widens each branch node into two columns, the fraction of positions carrying a non-zero  $A_H$  is bounded by  $2(n - 1)/T$  for *any* sampler whatsoever. With  $n = 8$  and responses of  $\sim 10^3$  tokens this ceiling is on the order of a percent. The tree’s role is to make the branch points exist and to place them at controlled depths, not to make them numerous.

## 9.3 The correction cannot move the mean

$A_H$  is a deviation from a sibling mean, so on each sibling set

$$\sum_{j \in \text{sib}(s_u)} A_H(s_u, y_u^{(j)}) = 0. \quad (25)$$

The correction in Eq. (23) therefore cannot change the average weight assigned to branch tokens. It can only spread siblings apart. This is not a defect to be engineered away: an unbiased estimator requires a baseline, and a baseline forces the mean to zero.

The consequence for evaluation is direct, and we state it before presenting any experiment. “Mean entropy at branch points increased” is *not* a prediction of this method and cannot serve as evidence that it works; the corresponding prediction is that the *variance across siblings* widens. A method designed to raise the mean is a different and weaker hypothesis — it does not need a forecast at all, only a support — and we include it below as a control arm.

## 10 The Approximation Ledger

Table 2 collects every departure from Eq. (8).

The strongest assumptions are A2 and B2. A2 has a known sign and a known monotonicity (Eq. 18), which is what makes the affine correction in A3 an appropriate remedy rather than a curve fit. B2 — that updating one token does not change the conditional distribution at other positions — is false for a network with shared parameters and has no corrective mechanism. It is not a risk this work introduces, since the base method already relies on it, but adding a second channel adds a second copy of its error.

## 11 Experimental Setup

### 11.1 Model, data, and optimization

All runs fine-tune the same base model, Qwen2.5-Math-1.5B (Yang et al., 2024), on the same data with the same optimizer settings, and differ only in the intervention; the five-arm decomposition and every ablation are run on it. We run no second backbone, and the Limitations say what that leaves open. All runs train on the same 17k-problem competition-mathematics set, DAPO-Math-17k (Yu et al., 2025). With a training batch of 512 prompts, 110 optimization steps correspond to about 3.2 epochs. Two further runs of the full method at 200 steps, under the base method’s own protocol, are described separately in Section 13.

### 11.2 Validation set

Every run validates on AIME 2024 (30 problems, 32 samples each) every 10 steps. AIME 2024 is small: with a per-problem accuracy spread of .1547 measured across our runs, the standard error of its mean is  $.1547/\sqrt{30} \approx .0282$ . That figure governs two things. It bounds what any single validation point can resolve, which is why the primary statistic below is a window mean rather than one point; and it bounds checkpoint selection, since a benchmark this noisy chooses between nearby checkpoints partly by chance. Where a checkpoint is selected (the six-benchmark evaluation of Section 12.9), the AIME24 column is therefore marked in-sample.

Optimization uses GRPO (Shao et al., 2024) with  $n = 8$  responses per prompt, a constant learning rate of  $1e-6$  with no warm-up (so that the schedule is independent of the declared total step count), a maximum response length of 3072 tokens, no entropy bonus, and no KL loss. The clipped

|    | approximation                              | what it discards                                                  | sign / bound                                     | mitigation                                             |
|----|--------------------------------------------|-------------------------------------------------------------------|--------------------------------------------------|--------------------------------------------------------|
| —  | Eqs. (5)–(15)                              | nothing — identities                                              | —                                                | —                                                      |
| A1 | horizon $\kappa$ , discount $\gamma_H$     | the tail $\sum_{k>\kappa}$ , and a fraction of the retained terms | under-estimate                                   | the common component cancels in the sibling difference |
| A2 | realized future $\rightarrow$ MTP marginal | mutual information with intervening tokens, Eq. (18)              | <b>over-estimate, monotone in <math>k</math></b> | A3                                                     |
| A3 | per-head affine calibration                | non-affine residual of A2                                         | least-squares residual                           | oracle-forecast control arm                            |
| A4 | clip $A_H$ to $\pm c$                      | magnitude of extreme branch scores                                | bounded by $c$                                   | saturation rate is logged                              |
| B1 | first order in $\eta$                      | $O(\eta^2)$                                                       | —                                                | small learning rate                                    |
| B2 | only the sampled logit moves               | cross-position effects of weight sharing                          | —                                                | <i>inherited from the base method</i>                  |
| B3 | rollouts come from $\pi_{\text{old}}$      | mismatch with the $\pi_\theta$ expectation                        | —                                                | importance ratio restores $r$                          |
| B4 | self included in sibling set               | $O(1/m)$ relative to leave-one-out                                | —                                                | same property as the GRPO advantage                    |

Table 2: Every approximation between the exact decomposition and the implemented weight. A1–A4 are introduced by this work; B1–B4 are shared with the token-level entropy-control method we extend.

surrogate (Schulman et al., 2017) is asymmetric,  $\varepsilon_{\text{low}} = 0.2$  and  $\varepsilon_{\text{high}} = 0.28$  as in Yu et al. (2025), with the dual-clip lower bound at  $c = 10$ ; we state this because “GRPO” with a symmetric  $\varepsilon = 0.2$  and a KL term is a different baseline, and because the follow-up arm of Section 12.10 is defined by moving these two numbers. Every arm, the baseline included, uses the same three values: the only configuration difference between the GRPO arm and STEER-F is the loss mode and the method hyperparameters listed next. Rollouts are sampled at temperature 1.0 with no nucleus truncation; validation uses temperature 1.0 and top- $p = 0.7$ . Every arm uses the same data seed and the same initial checkpoint, which we verify by requiring the step-0 validation scores to agree.

### 11.3 Method hyperparameters

The forecaster uses  $\kappa = 2$  heads with discount  $\gamma_H = 0.7$  and the per-head calibration of Eq. (19). The combination uses  $\lambda = 0.25$ , RMS normalization, clip  $c = 1.0$ , and the sibling-prefix baseline. The weight map is the base method’s exponential map of Eq. (2), taken un-rescaled, at  $\lambda_{\min} = 0.7$ , so weights land in  $[0.7, 1]$  and can only attenuate. This is the value the base method’s paper sets, inside the range  $[0.6, 0.8]$  over which it reports the method performing consistently (Hao et al., 2026, App. F.1); it was taken from there and not tuned on these runs. The base method’s released launcher uses  $\lambda_{\min} = 0.8$ , and every other value in that launcher matches ours.<sup>1</sup> The tree sampler uses

| arm      | $\lambda$ | rollout | signal in the weight |
|----------|-----------|---------|----------------------|
| GRPO     | —         | plain   | none                 |
| STEER    | 0         | plain   | local $\Omega$ only  |
| uniform  | .25       | tree    | support only         |
| permuted | .25       | tree    | $A_H$ shuffled       |
| STEER-F  | .25       | tree    | $A_H$ as derived     |

Table 3: The five arms. UNIFORM and PERMUTED run the full pipeline with the forecast’s value destroyed in two different ways, so any gain they show is attributable to the apparatus rather than to the visitation signal.

a single root, cuts at depths  $\{64, 192, 384\}$ , and branching factor 2 at each cut, giving  $n = 8$  leaves.

### 11.4 Five arms

The design isolates one question: does the *value* of the forecast matter, or only the apparatus that carries it? A comparison between the base method and the full method changes two things at once ( $\lambda$  and the sampler), so it cannot answer this. We therefore add two arms that keep the entire apparatus — the tree, the forecaster, the support, the weight mapping — and destroy only the information in  $A_H$ .

UNIFORM applies the same attenuation to every position in the support, discarding  $A_H$ ’s value but keeping its support — this is the implementation of the “raise the mean” hypothesis that Section 9.3 distinguishes from ours. PERMUTED keeps the multiset of  $A_H$  values within each sibling set and shuffles the assignment, so the support, the magnitudes, and Eq. (25) all survive while the pairing between

<sup>1</sup>Our implementation exposes an optional reshaping of  $\tilde{\Omega}$  before this map (winsorizing or rank-transforming it); every

run reported here uses the identity setting, so the base method’s map sees exactly what it would have seen.

a branch and its score is destroyed. Together they bracket the claim: if the visitation signal is doing the work, both controls should behave like the no-forecast arms.

## 11.5 Metrics

We report  $\text{acc@32}$  (mean accuracy over 32 samples) and  $\text{maj@32}$  (accuracy of the majority vote over the same 32 samples), and their difference

$$\text{uplift} := \text{maj@32} - \text{acc@32},$$

which measures how much better a set of 32 samples is than its average member — that is, whether correct answers concentrate while incorrect ones scatter. We use uplift rather than  $\text{pass@32}$  because the latter does not separate the arms.

Because single validation points are noisy on a 30-problem set, the primary statistic is the mean over the converged window, steps 40–110 (8 validation points per arm). We quote differences in units of the across-problem standard error,  $\text{SE} = \text{std@32}/\sqrt{30} \approx 0.027$ , which is the correct scale for generalization to new problems. We separately note that the pure re-sampling noise of the protocol — estimated from two validations of identical weights — is about 0.004 in accuracy, an order of magnitude smaller; the two scales answer different questions and we are explicit about which is in use.

One instrumentation caveat propagates into every entropy number we report. The logged `actor/entropy` is aggregated with a token-mean, so it is  $\mathcal{H}$  divided by the mean response length, not  $\mathcal{H}$ . Since response lengths differ by under 5% across arms, we report the raw quantity and, where the distinction could matter, the length-corrected product as well.

## 12 Results

All numbers below come from the runs described in Section 11: 110 optimization steps per arm from the same initial checkpoint, validated on AIME24. Every figure is recomputed from a training log released with this paper by one script; we report our own measurements only, and no figure in this section is taken from prior work. Section 13, which follows it, is the one place that does.

Every statistic is over the same 3 seeds, 1, 3, 4 — the seeds on which all five arms finished — so a difference of two rows of Table 5 is the corresponding paired entry of Table 7. Seed 2 is excluded because its STEER-F run did not complete; its other four

arms did, and every run, excluded or not, is in the released `per_seed.tsv`. Section 12.8 reports what the repeats measure — the run-to-run spread, and the machine each seed ran on.

**Which result carries which weight.** What follows is not one kind of result, and the order is the order of how much each one bears. Section 12.3 comes first. It tests a prediction the construction makes before any run — that a correction this sparse and this mean-preserving cannot move an average taken over every response token — and it tests it as a contrast *between groups of contrasts*, on 3 seeds, with the two groups’  $|t|$  not overlapping. That is the structural claim, and it is what this campaign establishes. The accuracy contrasts after it rest on 3 paired seeds each, and they do not all survive: the gap against STEER clears a conventional threshold, the gap against PERMUTED is marginal, and the headline gap against GRPO is smaller than GRPO’s own spread across its seeds (Section 12.8), does not, and is close to zero on the seeds where equal steps mean equal compute (Section 12.11). One pre-registered falsification condition is met at this seed count, on the contrast that isolates  $A_H$ ’s value; Section 12.4 states which and what it leaves standing. A reader who wants one result from this section should take the dissociation, not the gap.

### 12.1 Validation trajectory

Table 4 gives AIME24 accuracy and majority-vote accuracy at every validation point of seed 1, the one seed every arm ran first; the other seeds’ trajectories are in the release. Four of the five arms start from an identical step-0 score; STEER starts at .043 rather than .039, a gap we attribute to re-sampling noise rather than to a different checkpoint, since two validations of the *same* base weights in our logs differ by .004 in accuracy.

### 12.2 Converged window

Table 5 averages each arm over steps 40–110.

Two orderings are worth stating before anything is subtracted. On accuracy all three arms that do not carry  $A_H$ ’s value — STEER, UNIFORM, PERMUTED — sit *below* plain GRPO, and STEER-F is the only one above it. It is above by less than GRPO’s own standard error across seeds, which is why the ordering and not the size is what we take from this table. On entropy GRPO holds the most by a clear margin and the four intervened arms sit close to one another, STEER-F among them, so

| step | GRPO      | STEER     | UNIF      | PERM      | STEER-F   |
|------|-----------|-----------|-----------|-----------|-----------|
| 0    | .039/.035 | .043/.040 | .039/.035 | .039/.035 | .039/.035 |
| 10   | .075/.127 | .064/.109 | .070/.124 | .062/.126 | .078/.148 |
| 20   | .120/.176 | .114/.195 | .102/.171 | .087/.155 | .103/.181 |
| 30   | .108/.195 | .135/.215 | .121/.223 | .099/.178 | .113/.211 |
| 40   | .144/.236 | .115/.187 | .119/.201 | .116/.197 | .141/.230 |
| 50   | .120/.204 | .125/.196 | .121/.216 | .153/.224 | .145/.258 |
| 60   | .126/.178 | .139/.202 | .125/.203 | .135/.190 | .145/.237 |
| 70   | .130/.200 | .128/.200 | .129/.197 | .132/.227 | .147/.257 |
| 80   | .135/.182 | .135/.202 | .140/.204 | .134/.197 | .144/.206 |
| 90   | .136/.191 | .144/.212 | .142/.178 | .130/.210 | .166/.239 |
| 100  | .132/.171 | .146/.207 | .134/.192 | .144/.212 | .157/.222 |
| 110  | .157/.208 | .132/.181 | .140/.171 | .152/.218 | .151/.230 |

Table 4: AIME24 acc@32/maj@32 at every validation point, seed 1. Rows below the rule are the converged window used for all subsequent statistics. One seed is one trajectory: the across-seed statistics are in Tables 5 and 7.

the two columns are neither aligned nor reversed. Neither ordering is the instrument for either question at this seed count, and the next section says why: read as contrasts rather than as a ranking, the entropy column tracks the damping apparatus and the accuracy column tracks  $A_H$ .

### 12.3 Entropy: the aggregate moves with the apparatus, not with $A_H$

#### What the method predicts here, fixed in advance.

This method does not set out to hold the aggregate entropy up, and two facts already established above say it cannot. Section 9 bounds the positions at which  $A_H$  can be non-zero at  $2(n-1)/T$ , about one in eighty; Section 9.3 shows  $A_H$  sums to zero over a sibling set, so at the positions it does reach it redistributes weight among siblings rather than adding any. An intervention that is this sparse and mean-preserving by construction should leave an average over every response token close to where it found it. Both statements were written into the manuscript before the campaign was launched, Section 9.3 explicitly so (“we state it before presenting any experiment”), and the artifact history dates them an hour ahead of the pre-registration discussed below.

**What the seeds show.** Table 6 groups the contrasts by what each one changes, and the grouping separates them cleanly. The two contrasts that change the apparatus move the aggregate: turning the damping on costs  $-.0419$  at  $t = -5.87$ , and adding the tree with damping applied uniformly returns  $+.0184$  at  $t = +13.09$  — the latter reproducing across all 3 seeds to within a thou-

| arm      | $n$ | acc@32            | maj@32            | uplift | entropy |
|----------|-----|-------------------|-------------------|--------|---------|
| GRPO     | 3   | .1436 $\pm$ .0064 | .2145 $\pm$ .0100 | .0710  | .1575   |
| STEER    | 3   | .1371 $\pm$ .0021 | .2090 $\pm$ .0054 | .0719  | .1155   |
| UNIFORM  | 3   | .1422 $\pm$ .0057 | .2114 $\pm$ .0093 | .0692  | .1340   |
| PERMUTED | 3   | .1405 $\pm$ .0030 | .2129 $\pm$ .0023 | .0724  | .1219   |
| STEER-F  | 3   | .1491 $\pm$ .0005 | .2301 $\pm$ .0042 | .0810  | .1263   |

Table 5: Means over the converged window (steps 40–110, eight validation points per run), every arm averaged over the same 3 seeds — 1, 3, 4, the ones all five arms have finished — with the standard error across them. Every number in this paper is over that set and no other. Because it is shared, a difference of two rows here is the paired contrast of Table 7, for all six: STEER-F – GRPO reads  $+.0055$  either way, and the one row where the two disagree in print does so by a single unit in the last place, because each mean is rounded before the subtraction. Using one seed set also equalizes the machine: each arm here is one two-GPU seed and two four-GPU seeds, so a row difference no longer carries a box mix (Section 12.8). The  $\pm$  is a standard error over 3 draws, not a precision — STEER-F’s three land within .0016 of one another, which is agreement worth noting and still three draws; read it against the run-to-run ranges in Section 12.8. Every seed, including the ones outside the shared set and the one held out for not finishing, is in the released `per_seed.tsv`. *uplift* is  $\text{maj@32} - \text{acc@32}$ . Entropy is the token-mean quantity discussed in Section 11, i.e.  $\mathcal{H}$  divided by mean length.

sandth. The contrasts that hold the apparatus fixed and change only what  $A_H$  contains move it by  $-.0077$  ( $t = -1.42$ ) when its value is discarded and  $+.0044$  when its pairing is destroyed. Accuracy runs the other way: the apparatus contrasts give  $+.0051$  and  $-.0065$ , while the two  $A_H$  contrasts give  $+.0069$  and  $+.0087$ . The quantity that moves the aggregate entropy is the damping apparatus; the quantity that moves accuracy is  $A_H$ . That is the predicted pattern, and it is the sense in which this paper is not a method for suppressing entropy.

#### A pre-registered condition we retire rather than score.

The pre-registration (Section C) also fixed a condition on this section: if aggregate entropy orders the arms the same way accuracy does, a dissociation argument fails. We retire that condition instead of reporting it, and say plainly why, since retiring a registered test is exactly the move a reader should be suspicious of. First, it tests the aggregate as a signal about  $A_H$ , which the two paragraphs above — and Section 9.3, committed an hour before the registration was — say it cannot be. Second, an ordering over five arm means is not stable at this seed count: the rank correlation between

| contrast                                                | $n$ | $\Delta\text{ent.}$ | $t$    | $\Delta\text{acc}$ | $t$   |
|---------------------------------------------------------|-----|---------------------|--------|--------------------|-------|
| <i>the apparatus changes, <math>A_H</math> does not</i> |     |                     |        |                    |       |
| STEER – GRPO                                            | 3   | -.0419              | -5.87  | -.0065             | -1.24 |
| UNIF – STEER                                            | 3   | +.0184              | +13.09 | +.0051             | +1.37 |
| <i><math>A_H</math> changes, the apparatus does not</i> |     |                     |        |                    |       |
| STEER-F – UNIF                                          | 3   | -.0077              | -1.42  | +.0069             | +1.16 |
| STEER-F – PERM                                          | 3   | +.0044              | +2.15  | +.0087             | +3.11 |
| <i>both change</i>                                      |     |                     |        |                    |       |
| STEER-F – STEER                                         | 3   | +.0108              | +1.76  | +.0120             | +5.24 |

Table 6: Converged aggregate entropy and accuracy, paired within a seed and grouped by what the contrast changes. Turning the damping apparatus on or off moves the aggregate entropy at  $t = -5.87$  and  $+13.09$ ; changing what  $A_H$  contains, with the apparatus held fixed, does not move it detectably ( $t = -1.42$ ). Accuracy runs the other way. A row with  $n = 1$  has no  $t$ .

converged entropy and converged accuracy comes out with a different sign depending on whether each arm is averaged over the seed common to all five (Table 5) or over its own, which differ in number while the campaign runs. Both seed sets are in the released `per_seed.tsv`. Third, and because of the first two, *we claim nothing from it in either direction*: the condition is not reported as passed, and the theory is not reported as damaged by it. It was the wrong instrument to have registered, and the honest correction is to say so and to register a better one.

**The measurement that is actually owed.** Section 9.3 states the prediction that a sum-zero correction does make: not that entropy at branch points rises, but that the *spread across siblings* widens. No quantity we currently log tests it. The trainer does log an entropy split (`steerf/branch_entropy`), and an earlier draft of this section read it as entropy at branch points; it is not. That monitor takes a fixed top decile of  $A_H$  — a tenth of all response tokens — while the positions  $A_H$  can reach are the one-in-eighty that Section 12.6 measures, an eighth as many. Since  $A_H$  is exactly zero at almost every position, the remainder of that decile is filled by an arbitrary tie-break rather than by branching. Section 12.6 reports what that monitor does and does not select. The replacement restricts the split to  $A_H \neq 0$ , the same predicate the correction itself uses, and is instrumented in the release; no run reported here carries it, so this test is registered and not performed. That test is registered in Section C, before those runs, and we state here that it is reg-

| contrast        | $n$ | $\Delta\text{acc}$ | $t$   | $\Delta\text{maj}$ | $t$   |
|-----------------|-----|--------------------|-------|--------------------|-------|
| STEER-F – GRPO  | 3   | +.0055             | +0.81 | +.0155             | +1.34 |
| STEER-F – STEER | 3   | +.0120             | +5.24 | +.0211             | +2.37 |
| STEER-F – UNIF  | 3   | +.0069             | +1.16 | +.0187             | +1.77 |
| STEER-F – PERM  | 3   | +.0087             | +3.11 | +.0172             | +2.67 |
| STEER – GRPO    | 3   | -.0065             | -1.24 | -.0055             | -0.87 |
| UNIF – STEER    | 3   | +.0051             | +1.37 | +.0024             | +0.39 |

Table 7: Contrasts paired by seed: each arm is averaged over its converged window, arms are paired within a seed, and  $t$  is the paired  $t$  on the  $n$  seed-level differences ( $\text{df} = n - 1$ ). Pairing within a seed cancels any per-seed effect common to all five arms. A contrast’s  $n$  is the smaller of its two arms’ seed counts.

istered after this section’s original condition was retired rather than before.

**What the numbers above cannot carry.** The claim is about magnitudes, not about one row failing a threshold. Changing the apparatus moves the aggregate by  $+.0184$  and  $-.0419$ ; changing what  $A_H$  contains moves it by  $+.0044$  and  $-.0077$ , and the smaller of the two apparatus effects is still more than twice the larger of the two  $A_H$  effects. The  $|t|$  bands separate the same way and do not overlap —  $-5.87$  and  $+13.09$  against  $-1.42$  and  $+2.15$  — which is why the argument is the contrast *between groups* rather than any single row. Two cautions travel with it. A small difference that reproduces across seeds can still reach a large  $t$ : the  $A_H$  contrast on PERMUTED does, at  $+.0044$ , and that is a statement about consistency, not about size. And 3 seeds entitle us to say the aggregate does not move *much* when  $A_H$  changes, not that it is unchanged.

#### 12.4 Does the forecast’s value matter?

Table 7 averages each arm over its converged window, pairs the arms within a seed, and reports the mean difference together with a paired  $t$  on those seed-level differences. The step-by-step version of this statistic, which the single-seed draft reported, is in Appendix D; it answers a question about one trajectory rather than about re-running.

These contrasts carry less than the dissociation above, and we read them for direction rather than magnitude. They pair on the same 3 seeds as that section, and the quantity they measure — an accuracy gap — is the one the seeds in hand do not separate from run-to-run spread (Section 12.8).

The block structure is what the design was built to expose. Every contrast that changes the *value* of  $A_H$  is positive; the one that leaves it alone is not distinguishable from zero. UNIFORM runs the

entire apparatus — tree sampler, forecaster, support, weight map — and differs from STEER by  $+0.0051$  ( $t = +1.37$  over 3 seeds): turning the machinery on and feeding it a constant buys nothing. PERMUTED keeps the machinery *and* the magnitudes of  $A_H$  and destroys only the pairing between a branch and its score, and STEER-F still exceeds it by  $+0.0087$ .

Every contrast here now rests on 3 paired seeds, and they do not all say the same thing. Against STEER the gap is  $+0.0120$  at  $p = 0.035$ ; against PERMUTED it is  $+0.0087$  at  $p = 0.090$ . Against UNIFORM and against GRPO it is  $0.366$  and  $0.503$  — the headline contrast is the weakest row in the table, for the reason Section 12.8 gives: it is the one that changes three things at once, and GRPO’s own spread across seeds is wider than the effect. The signs agree in the mean everywhere  $A_H$  carries information, but the per-seed signs do not, and the next paragraph is about the one place they part.

**A pre-registered falsification condition is met.** The pre-registration (Appendix C) named four things that would count against this paper, before any of these runs existed. One of them has happened, and it is the one that clause calls the claim of the paper: *if STEER-F – UNIFORM straddles zero across seeds, the forecast’s value does not matter and the apparatus explains the effect*. On accuracy it straddles. Across the 3 paired seeds the difference runs from  $-0.0020$  to  $+0.0182$ , with 1 of them negative. We report that as a refutation of the registered claim rather than as noise, and we do not rescue it with the mean or with the majority-vote column, where the same contrast happens not to straddle (0 of 3 negative) — the registered primary endpoint is converged-window accuracy.

The other three conditions did not trigger, and the contrast between the first and the second is the informative part. STEER-F – PERMUTED, which holds the magnitudes of  $A_H$  and destroys only the pairing, does *not* straddle:  $+0.0033$  to  $+0.0125$ , 0 negative. Aggregate entropy does not order the arms the way accuracy does — STEER-F is third from the top on one column and first on the other in Table 5 — so the dissociation argument of Section 12.3 stands. And the mean token weight stays at  $.9990$ , so  $A_H$  is sum-zero over sibling sets as derived.

Read together: the *pairing* of  $A_H$  to a branch survives its registered test and its *value* does not. The two arms were built to separate exactly those

two things, so this is the distinction the design can draw, not an evasion of the result. It leaves the paper with a tension we do not resolve here. The pre-registration calls the UNIFORM contrast the central claim; the argument this paper now leads with is the dissociation of Section 12.3, which on these same 3 seed pairs separates entropy from accuracy at  $p = 0.005$  against  $p = 0.503$ . Both are on the record; a reader who weights the pre-registration above the reframing should read the accuracy claim as refuted at this seed count.

**The same contrast on questions the checkpoints were not chosen on.** The registered secondary check of this reading is the six-benchmark evaluation of Section 12.9. On the evaluated seed (1) it counts STEER-F above UNIFORM on 3 of six benchmarks, which the pre-registered thresholds call *equivocal*, and on 2 of the 5 that played no part in selection; paired over 1445 problems, the difference is  $-0.0090 \pm 0.0099$  ( $t = -0.91$ ). It does not change the condition above, which is defined on converged-window accuracy across the training seeds. It is the registered check on other questions, reported because it was registered and not as a rescue.

This is the contrast the design was built for. A comparison of STEER-F against STEER alone changes two things at once, and cannot separate the forecast from the sampler that delivers it; the two destroyed-forecast arms do separate them.

## 12.5 Where the gain comes from

The majority-vote gain decomposes into a part explained by per-sample accuracy and a part explained by the sample set’s structure:

$$\underbrace{+0.0155}_{\Delta_{\text{maj}}} = \underbrace{+0.0055}_{\Delta_{\text{acc}}} + \underbrace{+0.0100}_{\Delta_{\text{uplift}}},$$

against GRPO; against STEER the uplift term is  $+0.0091$  ( $t = +1.24$  over 3 seeds). Against GRPO roughly two thirds of the improvement is not that individual samples are more often right, but that the right ones agree with each other more and the wrong ones agree less. Both terms are means over the same 3 seeds, and the split inherits the weakness of the gap it splits: against STEER the uplift term is smaller and well inside its own spread, and against GRPO the whole gap is inside GRPO’s.

Where that structural change comes from is a separate question, and the arm that answers it finished after this section was first written. UNIFORM

runs the sampler at  $\lambda = .25$  with the forecast’s value destroyed and shows no uplift gain, which reads as though the structure belonged to  $A_H$ . It does not: the  $\lambda = 0$  tree arm — the sampler with the future channel off, STEER’s local damping still in place — reaches an uplift of .0870, against .0719 for STEER and .0810 for STEER-F. On this evidence most of the structural gain is the sampler’s, and Section 12.10 reports it that way. The decomposition above is unaffected — it is measured against GRPO, which has no tree — but the attribution is not what the five-arm design alone could see.

## 12.6 Mechanism diagnostics

Section 9 bounds the fraction of positions that can carry a non-zero  $A_H$  at  $2(n - 1)/T$  for any sampler. With  $n = 8$  and the tree arms’ mean response length of 977 tokens over the converged window the bound is .0143. Measured over the same window and seeds, the three tree arms sit at .0131 (UNIFORM), .0123 (PERMUTED) and .0125 (STEER-F), i.e. at 86–91% of a ceiling that no sampler can exceed. A plain-rollout pilot with the same  $\lambda$  (seed 1, released with the logs) reaches .003. As Section 8 notes, this count admits float32 round-off and is an upper bound; the runs reported here predate the round-off-free counter, so we quote the loose figure and label it as such.

Two readings follow. The tree sampler is close to saturating what the combinatorics allow, so there is little headroom in scheduling more cuts; and the intervention is extremely sparse — about one position in eighty — which makes the size of the effect in Table 7 the surprising part rather than the small support.

That sparsity also disqualifies one of our own monitors, which we record here rather than leave for a reader to find. The logged entropy split (`steerf/branch_entropy` and its gap) selects the top decile of  $A_H$  by rank, a tenth of all response tokens, against a support an eighth that size. A rank-based cut was chosen deliberately —  $A_H$  is exactly zero wherever a rollout has become its own only sibling, so a value-based threshold selects almost everything — but with a support this small the decile is filled out by the tie-break among those zeros, and the tie-break is not a random sample of them. A top-decile cut does recall about half of the true divergence points, five times what selecting at random would give, so the bucket is genuinely enriched; it is nonetheless mostly not branch positions, and the mean entropy it reports

cannot be attributed to them. Section 12.3 states what replaces it.

## 12.7 Is the future term a local term?

Section 6 reports a calibration that suppresses every head past the first, leaving  $\hat{H}_{\text{togo}}$  dominated by one step ahead. That invites the sharpest objection to the paper: if the forecast reaches one step, is  $A_H$  anything a purely local quantity could not have produced? The question is answerable without training anything. We apply the *same* operator — the same sibling baseline, mask and grouping — to two inputs: the calibrated forecast, and the policy’s own next-token entropy read at the same offset, the state after the branch token. Any difference between the two is the forecaster’s contribution and nothing else. We report it on the sibling support, where  $A_H$  can be non-zero at all; outside it both are identically zero and a statistic over all positions measures the support rather than the forecast.

The answer is that the forecaster is not earning its cost. Over 464 branch points from 4 validation prompts, the two advantages correlate at  $r = +0.917$  (Spearman  $+0.937$ ), the local twin explains  $R^2 = 0.840$  of the treatment, and the two agree in sign 0.921 of the time — and the sign is what the token weights are monotone in. The levels correlate less than the advantages ( $+0.881$ ), which is the wrong way round for an artefact of scale.

The mechanism is one the paper had already measured without connecting it. With  $\kappa = 2$  and  $\gamma_H = 0.7$  the calibration instantiates to  $\hat{H}_{\text{togo}} = 0.72 H_1 + 0.06 H_2 + 0.11$ , the sibling baseline removes the constant, and the first head is built to reproduce the policy’s own next-token distribution — its final layer is zero-initialised and its fitted  $a_1 = 1.03$  says it stayed there. So  $A_H$  should be close to  $\gamma_H a_1 = 0.72$  times the local twin, and the measured magnitudes are 0.3445 against 0.4647, a ratio of 0.741.  $A_H$  is very nearly a scalar multiple of a quantity the trainer already computes, and Eq. (23) divides the scale out through  $\rho = \text{rms}\{\text{visit}\}$ .

Two things this does *not* say. It does not touch Eq. (8), which is an identity, nor Section 4.4: the local twin is not  $\Omega$ .  $\Omega$  carries no sibling baseline and no tree, vanishes at a uniform branch point, and the twin exists only because the tree supplies siblings to difference against. What the measurement indicts is the *forecasting apparatus* — the MTP heads, their warm-up and their calibration — not the visitation channel. Read alongside Sec-

tion 12.10, where the sampler alone recovers .1439 against STEER’s .1371 while the forecast adds the remainder, two independent measurements now point the same way.

The caveats are the ones a single diagnostic run carries: 4 prompts at one validation step of the base policy, on independently sampled rollouts, whose sibling support therefore sits early in the response rather than at the scheduled branch depths. The head-free arm of Table 9 is the direct test, and it has now been run. ORACLE reads  $H_{\text{togo}}$  from the policy’s realised entropy, which verl (Sheng et al., 2025) computes in the same pass, so it carries none of the MTP overhead of Section 12.12 — 416 s per optimization step against STEER-F’s 516 at the same seed on the four-GPU box, 19.3% less. It is also behind it on accuracy. Section 12.10 reports both numbers and what one seed of each can carry.

## 12.8 Seeds, not steps

Every statistic above treats the seed as the unit of analysis: each run is first collapsed to one number by averaging its converged window, the arms are then paired within a seed, and the test is a one-sample  $t$  on the paired differences with  $\text{df} = n - 1$ . Pairing within a seed matters because a seed carries effects common to all five arms — the data order, and on a split campaign the machine a run landed on — and those cancel exactly in a difference. A contrast’s  $n$  is the smaller of its two arms’ seed counts, which is why the table reports  $n$  per row rather than once.

The alternative statistic, a paired  $t$  over the eight window steps *within* one run, measures whether a gap is stable along a single trajectory. That is a different question from whether it survives re-running, and we report it only in Appendix D so the two are not confused.

### Mechanical failures, and what replaced them.

The pre-registration (Appendix C) excludes a run for mechanical failure and relaunched it with the same seed. STEER-F at seed 4 stopped at training step 91 when its machine lost a GPU; it was relaunched at seed 4 and the relaunch, which reached step 110, is the run reported. STEER-F at seed 2 failed at start-up on a missing module and was not relaunched within this campaign’s budget, which is why seed 2 is outside the seed set. Both failed runs are in the release, flagged rather than deleted.

**A seed index also names a machine.** A seed of this campaign ran on a two-GPU box or a four-GPU

one, uniformly across arms, and the seed set this paper uses — 1, 3, 4 — holds 1 of the first and 2 of the second for every arm. The effect is measurable on *GRPO*, which has both: .1350 on the smaller box against .1479 on the larger, a difference of +.0129, larger than the gap between the top two rows of Table 5.

Read that difference as an upper bound and not as an estimate. With 1 run on the smaller box it is one run against 2, and a single run carries this arm’s own run-to-run spread (Section 12.8 below) inside it. We report it at that strength because the alternative was to keep a seed outside 1, 3, 4 alive for this one number while every other number in the paper excluded it, and a nuisance estimate is not worth a second seed set.

What the figure does not reach is the comparisons, for two reasons that do not depend on how well it is estimated. The contrasts pair inside a seed and therefore inside a machine, so whatever a box is worth cancels in the pair. And every arm in Table 5 carries the same 1-to-2 mix, so a row difference does not carry a box mix either. The release prints the per-arm machine mix on every run of the analysis.

### Run-to-run spread, and it is the number that governs this section.

Every arm has now finished more than once. GRPO’s converged-window accuracy moves from .1350 to .1560 across its 3 seeds, a spread of .0210; STEER-F moves from .1481 to .1497 across its 3 seeds, a spread of .0016; PERMUTED moves only .0095 across 3. Both extremes are measured on the same benchmark with the same protocol, and they differ by more than an order of magnitude.

Set the headline gap of +.0055 against those. **It is a fraction of GRPO’s own spread across seeds and several times STEER-F’s**, so the two arms it compares do not even agree on how large a seed-to-seed difference is, and each is observed a handful of times. That is the regime in which a single paired difference cannot decide the question, and we do not claim it does. A range between three draws is also not a standard error: the largest of these ranges is more than ten times the smallest, so they do not bound each other and no contrast here should be read against a noise floor quoted to one significant figure. For scale at the other end, two validations of the *same* weights on these 30 questions differ by about .004, so seed variation is the larger term by roughly an order of magnitude and any comparison

quoted at single-run precision is quoting noise.

We therefore state what this experiment establishes and what it does not. The mechanism contrasts — what changes when  $A_H$ 's value is discarded or its pairing destroyed, with the apparatus held fixed — are the claim. The effect size would need more seeds than this campaign has, and on the seeds it has the headline gap is not separated from the spread.

## 12.9 Six benchmarks

A single 30-question benchmark cannot separate a method from a lucky draw, so each run's selected checkpoint is re-evaluated on six mathematics benchmarks (Table 8) under the protocol of Section 11: AIME24, AIME25 and AMC23 at avg@32, MATH500 (Hendrycks et al., 2021), Minerva (Lewkowycz et al., 2022) and OlympiadBench (He et al., 2024) at avg@1.

An evaluation enters the table only if its log agrees with itself. The trainer prints each benchmark twice on the same line, as accuracy and as the mean of a reward that is +1 for a correct sample and −1 otherwise, so on every benchmark the accuracy must equal half of one plus the reward; the release checks this for every evaluation log (scripts/check\_eval\_logs.py). It is a test of coherence, not of provenance — an edit that moved both numbers together would pass it. Absent from the table, for want of an evaluation that passes it: STEER. Each row averages the checkpoints evaluated at the seeds in its last column.

One of the six is not held out. The checkpoint behind each row is the argmax of that run's AIME24 validation accuracy (Section 11.2), so AIME24 is the set the selection was made on and its column is in-sample: the pre-registration (Appendix C) records that this argmax lifts an arm by about a hundredth in absolute accuracy, and by an amount that differs between arms, so the bias does not cancel cleanly inside a contrast. The column is marked † and kept — dropping a registered endpoint after seeing it is the worse of the two errors — and the five remaining benchmarks, which played no part in selection, carry no such bias and only more variance.

The quantity fixed in advance as the headline here is not any single cell but *direction consistency*: in how many of the six benchmarks the ordering observed on AIME24 is preserved. Four or better of six we declared support, three equivocal, two or fewer evidence against. The observed count is 3 of

| arm      | AIME24 <sup>†</sup> | AIME25 | AMC23 | MATH500 | Minerva | Olympiad | mean | seed |
|----------|---------------------|--------|-------|---------|---------|----------|------|------|
| GRPO     | 15.8                | 8.2    | 54.8  | 71.4    | 26.5    | 33.9     | 35.1 | 3    |
| UNIFORM  | 15.7                | 6.7    | 56.0  | 72.2    | 25.7    | 33.6     | 35.0 | 4    |
| PERMUTED | 15.2                | 6.4    | 55.8  | 69.0    | 23.9    | 31.2     | 33.6 | 3    |
| STEER-F  | 17.1                | 7.3    | 52.5  | 69.4    | 23.9    | 34.5     | 34.1 | 4    |
| ORACLE   | 14.2                | 8.2    | 58.2  | 70.8    | 24.3    | 33.7     | 34.9 | 4    |

Table 8: Six-benchmark evaluation of each arm's selected checkpoint, averaged over the seeds in the last column. The mean column is the unweighted mean over the six. †: AIME24 is the set each checkpoint was selected on (Section 11.2), so that column is in-sample; the other five are held out. Rows at different seeds are not meant to be subtracted — a seed index also names a training run and the machine it ran on — and the contrasts below are counted within a seed.

six for STEER-F > UNIFORM, which those thresholds call *equivocal*, over the seeds the two arms share (1). The STEER-F > PERMUTED count is not reported: a count needs the two arms evaluated at a common seed, and the evaluations that passed the coherence check do not include one. Counting the benchmarks on the 5 that played no part in selection gives 2: the registered six-count is the decision rule, and the held-out count is printed beside it so a reader can see how much of it survives without the set the checkpoints were chosen on. This is the test that distinguishes a method from an artifact of one question set, and it is reported whatever it returns.

These two contrasts carry the attribution because each changes one thing: UNIFORM has the apparatus without  $A_H$ 's value, PERMUTED has its value without its pairing. The differences against GRPO and STEER are registered as well and appear in Table 7, but they change three things and two respectively, so a benchmark difference there names no cause. Arms are paired inside a seed throughout, for the reason Table 5's caption gives; the arm-mean rows of Table 8 are not what is being counted.

**The paired error bar.** A single measurement of one arm on AIME24's 30 questions carries an across-problem standard error of .0282, which is why the gaps in Table 7 are a fraction of one problem in absolute terms. That figure is the wrong one for a contrast: the two arms answer the same questions, so per-problem difficulty cancels between them and what matters is the standard error of the *paired* difference. This evaluation records one row per generated sample, so on the three avg@1 benchmarks with enough problems to support it — MATH500, OlympiadBench and Minerva, 1445

problems together, over the seeds the arms share (1) — the paired quantity is computed rather than asserted. STEER-F–UNIFORM is  $-0.0090 \pm 0.0099$  ( $t = -0.91$ ) against an unpaired error of 0.0185; the PERMUTED contrast has no common seed and so no paired figure. How much pairing buys is decided by how often the two arms disagree on a problem, which is 0.142 here; the measurement is what we report, not the reduction we wanted. Two cautions travel with it. At avg@1 each problem’s score is a single Bernoulli draw, so these errors carry the sampling noise of one generation per problem on top of the difficulty spread; and pooling problems across seeds treats a (seed, problem) pair as one observation, which ignores the clustering by seed that Section 12.8 shows is real. GSM8K is computed the same way and left out of the pool: at avg@1 it sits near saturation, where a paired difference measures the ceiling rather than the method.

## 12.10 How much of the design is load-bearing

Table 9 varies one element at a time, at a single seed per row. Every comparison in it is therefore  $n = 1$ : these rows describe a mechanism, and none of them carries an error bar or licenses an effect size. A row is read against the main arm at *its own* seed, which is not the same number as Table 5’s across-seed mean.

**The coefficient.**  $\lambda$  scales the visitation term before the two channels are combined in Eq. (23). The value used throughout was not tuned on these runs, and the sweep exists to say whether the result rests on a knife edge. A flat response across 0.1–0.5 would say the method is insensitive to it; a peak at 0.25 with both neighbours below would say it is a tuned quantity and must be reported as one.

**The base method against GRPO, and what it does and does not reach.** We measure STEER below GRPO by  $-.0065$  over 3 seeds, where the base method’s own results place it above. We read the released weighting function against ours and it agrees line for line, as does every other value in the two launchers except  $\lambda_{\min}$ , which we set to the paper’s 0.7 where the released script uses 0.8. That difference does not explain the gap on the base method’s own evidence: both values lie in the range its sensitivity analysis reports as performing consistently. The settings that do differ from the base method’s main experiments are the

| arm                                                                                           | what changes | acc@32 | maj@32 |
|-----------------------------------------------------------------------------------------------|--------------|--------|--------|
| <i>the damping coefficient</i>                                                                |              |        |        |
| $\lambda = 0.1$                                                                               | weaker       | .1365  | .2360  |
| $\lambda = 0.25$                                                                              | as used      | .1491  | .2301  |
| $\lambda = 0.5$                                                                               | stronger     | .1444  | .2265  |
| <i>the sampler, with the forecast switched off</i>                                            |              |        |        |
| $\lambda = 0$ , tree                                                                          | tree only    | .1439  | .2309  |
| STEER                                                                                         | neither      | .1371  | .2090  |
| <i>the forecaster, replaced by what the trainer already has</i>                               |              |        |        |
| ORACLE                                                                                        | no heads     | .1369  | .2234  |
| <i>relaxed ratio clipping (<math>\epsilon_{hi}=5</math>, <math>\epsilon_{lo}=0.99</math>)</i> |              |        |        |
| STEER-F                                                                                       | —            | .1516  | .2304  |

Table 9: Follow-up ablations, converged window, one seed per row and no error bar. The seed is whichever one the arm was run at, and the release emits it with the row; the ORACLE arm is seed 4, run there so that it shares a seed with the STEER-F and UNIFORM runs the six-benchmark evaluation pairs. The  $\lambda = 0.25$  line is the reference and is Table 5’s across-seed mean rather than one run. Every row but ORACLE ran on the 4-GPU box while the seed-1 STEER-F reference ran on the smaller one, so those rows carry the machine effect of Section 12.8 (+.0129, in the direction that flatters them) on top of whatever the ablation does. The release emits a same-seed contrast for ORACLE alone, because it is the one row whose box matches its reference’s. Each block isolates one element of the design. The relaxed-clipping block has no STEER control: that run stopped at training step 62, before the converged window closes, and a mean over the part of the window it reached is not the quantity the other rows report. Arms that were planned and did not produce a finished run are listed in Appendix C.

schedule (110 steps against its 200), the checkpoint rule (a fixed window against its best checkpoint) and the number of runs, and we have no measurement that separates them. We did not run the base method at the released 0.8, so the question of why it lands below GRPO here is left open rather than attributed.

What this discrepancy can reach is bounded, and by the design rather than by luck. All four damped arms share  $\lambda_{\min} = 0.7$  exactly, so the setting cancels in any contrast between two of them. STEER-F – UNIFORM and STEER-F – PERMUTED — the pair that carries this paper’s central claim, that the forecast’s *value* is what matters — hold  $\lambda_{\min}$ ,  $\lambda$ , the sampler and the weight map all fixed and vary only what  $A_H$  contains. Those two do not depend on whatever puts the base method below GRPO here. What it does reach is how two other contrasts read. STEER – GRPO is the discrepancy itself. STEER-F – STEER is measured against a base method that is not showing its reported gain over GRPO, so it is not a measurement of improve-

ment on top of that gain. We report both with this stated rather than resolved.

**The sampler.** This is the cell no other arm supplies. Tree rollouts make a group share prefixes, which changes the GRPO advantage estimate itself, quite apart from any damping. UNIFORM and PERMUTED both run at  $\lambda = .25$  and so cannot separate the sampler from the correction. The  $\lambda = 0$  tree arm switches the forecast off while leaving the sampler in place, and is the direct answer to “is tree rollout simply a better sampler”. Read against STEER, which has neither, its difference is the sampler’s own contribution; read against STEER-F, what remains is the forecast’s.

The answer is not the one the rest of this paper was written expecting. The sampler alone reaches .1439 against STEER’s .1371, and .2309 against .2090 — a little over half of STEER-F’s accuracy margin over STEER and all of its majority-vote margin. Its uplift, .0870, is not below STEER-F’s .0810. What remains for the forecast is the accuracy gap between .1439 and .1491.

Three things keep this from being decisive, and none of them rescues the original reading. Every cell here is one run, and the spread across this block is about twice the run-to-run spread measured in Section 12.8 — suggestive, not settled. The  $\lambda$  sweep is not monotone:  $\lambda = .1$  sits below  $\lambda = 0$  on accuracy, which is what noise looks like at this sample size. And the row ran on the 4-GPU box while the seed-1 references it is read against did not, so the +.0129 of Section 12.8 sits inside this comparison and points the same way as the conclusion drawn from it: the sampler’s share is an upper bound, not an estimate. We report the block as measured and do not reconcile it. Whether the sampler or the forecast carries the structural term needs these arms at more seeds, which this paper does not have. Nothing here establishes that the tree does not carry it, and Section 12.4 is explicit that the forecast’s *value* does not survive its own registered test.

**The forecaster, replaced by what the trainer already has.** Section 12.7 builds the same sibling-differenced quantity from the policy’s own next-token entropy and finds it close to  $A_H$ . That measurement is taken on a frozen policy at one validation step, and it cannot say what the substitution does *inside* the loop, where the weights it produces move the trajectory the next measurement would be taken on. The ORACLE arm makes the substi-

tution there: same tree, same  $\lambda$ , same weight map, with  $H_{\text{togo}}$  read from the realised entropy of the rollout verl has already computed. A2 and A3 — the head substitution and the affine calibration — are the only things it removes, and removing them is the whole of its cost difference.

It is cheaper, and by the amount the forward pass costs: 416 s per optimization step against STEER-F’s 516 at the same seed, a saving of 19.3%. Both are the window’s training steps on the four-GPU box, where every arm is substantially faster than on the two-GPU seeds, so neither is comparable to a step time quoted for a different box.

It is also behind on accuracy. Over the 8 validation points of the converged window, STEER-F at seed 4 leads it by +.0112 on the endpoint this campaign reports and by +.0104 under majority vote. We therefore keep A2 and A3 rather than replacing the forecaster with the realised value, and we state the weight of that decision exactly: one run on each side, one seed. The gap is larger than STEER-F’s own range across seeds, .0016, and smaller than GRPO’s, .0210. It orders the two arms on this seed; it does not measure the difference between them, and the measurement that would be this arm at more seeds.

**Clipping.** The relaxed-clipping setting is where entropy interventions are supposed to matter most, because ratio clipping otherwise removes much of the update a token weight would have scaled. The row is one STEER-F run with no finished control under the same clipping, and it ran on a different machine from the seed-1 STEER-F reference, so it carries no contrast in either direction. It shows that the method trains to the end of the schedule with the clip relaxed; it does not say whether relaxing the clip changes what the method does.

## 12.11 Compute-matched comparison

An equal-step comparison is compute-fair only if the two arms cost the same per step, and here that depends on the machine. On the 4-GPU box (seeds 3, 4) STEER-F’s step costs 0.99 times GRPO’s, so equal steps is equal compute to within the seed spread of either arm (Section 12.12). On the 2-GPU box (seed 1) it costs 1.49 times as much, and a GRPO run given the same wall clock would reach about step 164. We did not train GRPO past step 110, so that point is not measured, and we do not extrapolate a trajectory to it.

What the split does measure is where the head-

line gap comes from, and it is not evenly spread. STEER-F – GRPO is  $+0.0145$  on the seed where STEER-F also had the larger compute budget, and  $+0.0011$  averaged over the 2 seeds where the comparison is compute-fair. The across-seed mean of  $+0.0055$  in Table 7 is therefore carried mostly by a seed on which the treatment was also granted about half as much again in wall clock, and on the compute-fair seeds the gap is not distinguishable from zero. We report this as the reading the evidence supports: this paper does not show that STEER-F buys accuracy over GRPO at equal compute. None of this touches the contrasts that carry the structural claim — STEER-F – UNIFORM and STEER-F – PERMUTED — whose arms share the tree sampler and cost 493, 500 and 538 s per step on the larger box.

## 12.12 Cost

Step times are read inside a box. A seed index also names a machine (Section 12.8), and the two boxes are further apart than any two arms are, so a mean over both measures the campaign’s scheduling.

On the 4-GPU box, where every arm has finished 2 seeds, one optimization step takes 449 s for STEER, 493 for STEER-F, 500 for UNIFORM, 538 for PERMUTED and 500 for GRPO. STEER-F’s overhead over STEER there is  $+10\%$ . The three tree arms take 493, 500 and 538 s while STEER, which has neither the tree nor the heads, sits below all three: the cost is the rollout and the forward pass, not the  $A_H$  computation, which is a subtraction over a mask.

On the 2-GPU box the same overhead reads  $+86\%$ . That is not a disagreement to average away. The tree rollout and the MTP forward are a larger share of a step when fewer GPUs serve it, so the overhead is a property of the deployment as much as of the method, and a single figure quoted without its box — as earlier drafts of this section quoted the smaller one — reports the machine.

Neither figure is resolved to better than the spread underneath it. On the larger box STEER’s 2 seeds differ by 54 s and STEER-F’s by 46, against a gap between their means of the same order. On the smaller box every arm has 1 seed inside 1, 3, 4, so there is no spread to put beside it at all. We therefore report these as the cost of the runs we made and not as a coefficient of the method. What does isolate the heads from the tree is the ORACLE arm of Section 12.10, which removes the heads at fixed tree and fixed seed.

## 12.13 A note on how these numbers were selected

We report window means rather than the best checkpoint, and the difference is not cosmetic. Selecting the maximum over the twelve validation points inflates every arm, averaged over the 3 seeds: by  $.0161$  for GRPO,  $.0172$  for STEER,  $.0128$  for UNIFORM,  $.0182$  for PERMUTED and  $.0189$  for STEER-F, relative to the window mean. Since a single AIME24 measurement carries a standard error of  $.0282$  across problems, a maximum over a dozen noisy points is a biased statistic whose bias grows with the number of checkpoints compared, and it is not the same for every arm: the lifts differ by more than the gap between UNIFORM and PERMUTED, so ranking by the maximum would swap those two. We regard the window mean as the defensible statistic and use it throughout. The one exception is Section 13, which uses the maximum on purpose, because that is the rule of the table it is compared against.

## 13 Against the Reported Numbers

Every number before this section is our own measurement under the schedule of Section 11, which differs from the base method’s in length, checkpoint rule and number of runs (Section 12.10). This section is the one place a figure is taken from prior work. It runs the full method under the base method’s own protocol and sets the result beside that method’s published 1.5B table (Hao et al., 2026, Table 12), which also reports GRPO and the untrained model. It is a second kind of evidence, not a sixth arm, and we say below what it cannot show.

**Protocol.** Two independent runs, with data seeds 1 and 2, of 200 optimization steps each; AIME24 validation every 10 steps; in each run the checkpoint with the highest AIME24 accuracy is selected and evaluated on the six benchmarks of Section 12.9, and a table entry is the mean over the two runs. Each of these items is the base method’s (Hao et al., 2026, App. E). Everything else is the configuration of Section 11, with one change: the tree sampler uses two roots, each cut at depths  $\{64, 192\}$  with branching factor 2, so the eight responses to a prompt form two independent groups of four instead of one group that shares its first 64 tokens. We chose this shape after the five-arm campaign, from short runs measuring how often the eight rewards of a group tie, which removes signal

|                                                  | AIME24 <sup>†</sup> | AIME25 | AMC23 | MATH500 | Minerva | Olympiad | mean |
|--------------------------------------------------|---------------------|--------|-------|---------|---------|----------|------|
| <i>reported by Hao et al. (2026), not re-run</i> |                     |        |       |         |         |          |      |
| base model                                       | 4.1                 | 2.1    | 24.7  | 29.0    | 9.2     | 20.5     | 14.9 |
| GRPO                                             | 16.2                | 7.6    | 56.0  | 74.4    | 26.1    | 34.6     | 35.8 |
| STEER                                            | 17.4                | 9.7    | 61.6  | 75.5    | 28.2    | 36.6     | 38.2 |
| <i>ours, mean of the two runs</i>                |                     |        |       |         |         |          |      |
| STEER-F, selected                                | —                   | —      | —     | —       | —       | —        | —    |
| STEER-F, step 200                                | 16.9                | —      | —     | —       | —       | —        | —    |
| <i>ours, each run's selected checkpoint</i>      |                     |        |       |         |         |          |      |
| seed 1                                           | 18.0                | 9.2    | 57.1  | 74.0    | 24.6    | 34.8     | 36.3 |
| seed 2                                           | —                   | —      | —     | —       | —       | —        | —    |

Table 10: Six benchmarks under the base method’s protocol, accuracy in percent; the mean column is un-weighted over the six. The upper block is quoted from the base method’s paper and was not re-run. In each of our runs the selected checkpoint is the argmax of its AIME24 validation accuracy. <sup>†</sup>: in our rows AIME24 is the training-time validation value at the selected step — the value the selection itself used — and is biased upward by that selection; the other five columns re-evaluate the selected checkpoint. The base method’s paper selects the same way and evaluates the selected checkpoint, but does not say whether its AIME24 column is the selection value or a re-measurement; if it is a re-measurement, this column favours our rows.

from a GRPO update. The choice is exploratory, it was made on training prompts and not on any benchmark, and these runs are not one of the five arms. The selected steps are 150 and 190.

Each AIME24 cell of ours is read from the training log and used only if it agrees with the reward printed beside it and with the trainer’s own record of the best value; the other columns pass the coherence check of Section 12.9 (`scripts/emit_r2_numbers.py`). The two-run AIME24 mean is averaged over correct-sample counts rather than over two rounded fractions.

**Reading the table.** Averaged over the two runs, STEER-F’s six-benchmark mean is —, against 38.2 for the base method and 35.8 for GRPO as reported; on AIME24 the same three figures are —, 17.4 and 16.2. One AIME24 evaluation carries an across-problem standard error of .0282 as a fraction (Section 11.2); a mean of two runs halves its variance and no more, and the quoted rows carry a noise of their own that their source does not state. We therefore read the table as placing the method at the level of the published numbers, and not as ordering it against them. The step-200 row is printed because the selected checkpoint is chosen on AIME24 and so flatters its own row; the last checkpoint is chosen by nobody.

**What this comparison is not.** It is not controlled. The quoted rows ran with the authors’ code on different hardware (a cluster of H20 GPUs, against our four H100s), and the base method under this protocol in our environment is a run this paper does not have. The controlled comparison with the base method is the five-arm campaign, where STEER-F — STEER is +.0120 over 3 paired seeds at 110 steps (Table 7), with the caveats of Section 12.10. Nor does the table speak to the visitation channel: like every comparison between the full and the base method (Section 11), it changes the sampler and the correction together, and the attribution rests on the controls of Section 12 alone.

## 14 Conclusion

The gradient of trajectory entropy has two terms, and the token-level entropy controls now in use estimate one of them. We showed that the omission is structural: the estimated term vanishes exactly where the omitted term is largest, and it vanishes for a correct reason, so the gap cannot be closed by improving the estimator. We then showed that the omitted term is a REINFORCE estimator whose reward-to-go is the remaining conditional entropy, and that a tree-structured rollout supplies the sibling baseline it requires — which is what makes the correction estimable at all. The resulting correction is mean-preserving by construction: it does not slow entropy collapse, it redirects it. The campaign measures that consequence and finds it. The aggregate entropy moves when the damping apparatus is switched and does not move detectably when only  $A_H$ ’s contents change, while accuracy separates the same two groups in the opposite direction. On the seeds in hand that is a statement about what is detectable, not a proof that the aggregate is untouched.

The accuracy evidence is weaker, and we leave it that way rather than dress it. The gain over GRPO is inside GRPO’s own seed spread and close to zero where equal steps mean equal compute; the registered test of  $A_H$ ’s value against a value-free control straddles zero and is reported as a refutation; the test of its pairing survives. A practitioner should read this paper as an account of where token-level entropy control acts — and of why an aggregate entropy curve cannot certify a correction of this kind — not as a recommendation to adopt the method for accuracy. Under the base method’s own protocol the method’s two-run average sits beside

the published numbers rather than above or below them by more than the measurement resolves (Section 13); that comparison is uncontrolled and does not change this reading.

Estimating  $H_{\text{togo}}$  is a separate problem, and the decomposition does not depend on how it is solved. Ours is a parallel multi-token-prediction forecaster whose error has a known sign and monotonicity in closed form; we also measure it to reproduce a quantity the trainer already computes at  $R^2 = 0.840$ , and to cost the step time reported in Section 12.12 to do so. That is a reason to look for a cheaper estimator of this reward-to-go, not a reason to doubt the channel: the head-free ORACLE arm of Table 9 is the control that separates those two readings, and it has been run. At the one seed it carries, the realised value costs 19.3% less per step and reads  $+0.0112$  below the forecast on the primary endpoint. That is a reason to keep the estimator rather than a licence to declare it necessary: a single draw on each side orders the two and does not size the difference.

## Limitations

**A pre-registered falsification condition is met.** Appendix C fixed, before these runs, that if STEER-F – UNIFORM straddled zero across seeds then the forecast’s *value* does not matter and the apparatus explains the effect. On the registered primary endpoint it straddles:  $-0.0020$  to  $+0.0182$  over 3 paired seeds, with 1 negative. That clause calls this the claim of the paper, and we do not renegotiate it here. What it does not reach is the companion condition: STEER-F – PERMUTED, which destroys only the pairing of  $A_H$  to a branch and keeps its magnitudes, does not straddle (0 of 3 negative), nor do the two structural conditions on the entropy ordering and on the mean token weight. The honest summary is that the pairing survives its registered test and the value does not, at this seed count, on this benchmark. Section 12.4 states the tension that leaves between the pre-registration’s choice of central claim and the one this paper leads with.

**The forecaster does not earn its cost.** The calibration in Eq. (19) suppresses heads beyond the first, so  $\hat{H}_{\text{togo}}$  is dominated by a single step ahead. This is the predicted behavior of Eq. (18) rather than a bug, but Section 12.7 measures the consequence and it is worse than a short horizon: put through the same sibling baseline, the policy’s own next-token entropy reproduces  $A_H$  at  $r = +0.917$

and  $R^2 = 0.840$ , with the magnitudes in the ratio 0.741 that  $\gamma_H a_1$  predicts. On that evidence the MTP heads, their warm-up and their calibration are machinery around a quantity the trainer already computes, and they carry the overhead of Section 12.12 to do it. The ORACLE arm of Table 9 is the direct test and it does not agree: made inside the training loop, the substitution saves 19.3% of the step time and gives up  $+0.0112$  on the primary endpoint at the one seed it was run at. Both readings stand as measured. A correlation on a frozen policy does not predict a training loop, and one draw on each side does not overturn a correlation over 4 prompts; whether the heads earn their cost is left open here, and the arm that closes it is the same arm at more seeds. What the measurement does not reach is the visitation channel itself: the local twin is not  $\Omega$ , and it exists only because the tree supplies siblings to difference against. A forecaster whose error did not grow with  $k$  would test the idea more sharply; a sequential forecaster is the obvious comparison and we have not run it.

## Few seeds, and a headline gap inside the spread.

Every contrast here carries a seed-level test, but on 3 paired seeds, 1, 3, 4, for every row, which is few; and one of the three ran on a different machine from the other two (Section 12.8). This bounds the effect sizes rather than the structural result: the dissociation of Section 12.3 is a comparison between two groups of contrasts whose  $|t|$  do not overlap, and no single row carries it. The accuracy gap is the claim that does rest on one row, and Section 12.8 shows the floor under it is not measured — GRPO’s converged-window accuracy moves  $.0210$  across its 3 seeds, STEER-F’s  $.0016$  across 3, and PERMUTED’s only  $.0095$ , three estimates of the same noise differing by more than an order of magnitude, so none of them bounds the others. Against the largest, the headline  $+0.0055$  is the smaller number. Rollout sampling is not seeded across processes, so re-running a configuration is a second draw rather than a restoration — which is what makes the repeat measurable and what makes it necessary. The pre-registration planned five seeds per arm; the campaign completed three for every arm (Appendix C lists the deviations), which settles the structural result and not the size of the accuracy gap. At equal compute the gap is close to zero on the seeds where it can be measured (Section 12.11).

**A single validation benchmark, and what our  $t$ -statistics do not cover.** Every claim rests on

AIME 2024, which has 30 problems. A single measurement of one arm carries an across-problem standard error of .0282, so the gaps in Table 7 are a fraction of one problem in absolute terms. The paired  $t$  in that table is taken over seeds, so it speaks to re-running a configuration and not to a different set of questions; the within-run version, over the eight validation steps of a single trajectory, is reported separately in Appendix D so the two are not read as one. The statistic that would speak to generalization — the across-problem standard error of the *paired* difference, in which per-problem difficulty cancels between two arms evaluated on the same questions — is not recoverable from the *training* logs, which record per-benchmark means and standard deviations rather than per-problem scores. The six-benchmark evaluation records the per-problem scores and so does carry it, on the three benchmarks large enough to support one; those figures are in Section 12.9. What remains missing is the same statistic for AIME24 itself, which is where every trajectory-level number in this paper is measured, and which is also the set the checkpoints were selected on. The gap this paragraph names is narrowed, not closed. The figure that narrows it rests on the seeds evaluated so far (1) and on one sample per problem at the validation temperature, so it carries the sampling noise of a single generation as well as the spread in difficulty; the evaluation covers one seed per arm, pairs PERMUTED with nothing, and does not include STEER.

**One prediction of the analysis is unmeasured.** Section 4.4 argues that the two channels are anti-correlated. The prefactor half of that claim is analytic and the  $A_H$  half is inferred from a correlation with branch diversity; we have not measured their *product* at real sibling-divergence positions in a live rollout. Until that measurement exists, Table 1 is a prediction rather than an observation.

A second one is owed by the result we lead with. Section 9.3 says a sum-zero correction should widen the spread *across siblings* at the positions it reaches, which is the positive claim behind the dissociation of Section 12.3; the release instruments the support-restricted split that tests it, and no run reported here carries it. The dissociation we do measure is the negative half — that the aggregate does not move with  $A_H$  — and the two are not the same statement.

**The support is small by construction.** Section 9 bounds the fraction of positions carrying a non-zero

$A_H$  at  $2(n - 1)/T$ , on the order of a percent for our settings. Whether the effect scales with  $n$ , or with a branching schedule chosen adaptively rather than at fixed depths, is untested.

**The base method does not reproduce its reported gain.** We run the base method with  $\lambda_{\min} = 0.7$ , the value its paper sets; its released launcher uses 0.8, and both lie in the range its paper reports as performing consistently. Under our schedule it lands below GRPO, where its authors report it above, and Section 12.10 lists the setting differences we cannot separate: schedule length, checkpoint rule and number of runs. We did not run it at the released 0.8 or at its full schedule, and we did not tune  $\lambda_{\min}$  per arm.

What this does not reach is the claim the paper makes. All four damped arms carry the same  $\lambda_{\min}$ , so it cancels between any two of them, and the contrasts that separate the forecast’s value from its apparatus — STEER-F against UNIFORM and against PERMUTED — hold it fixed along with everything else. The two contrasts the discrepancy does reach, STEER-F — STEER and STEER — GRPO, are reported with it stated.

**The comparison with published numbers is not controlled.** Section 13 sets two 200-step runs beside the base method’s published table. The quoted rows were not re-run, they come from different hardware and code, and their source does not state a run-to-run spread; the base method was not run under that protocol here. The sampler in those two runs was chosen after the five-arm campaign from measurements on training prompts, so it is an exploratory configuration and not a registered one. Their AIME24 column is the value the checkpoint selection used, biased upward by it, and whether the published column is the same kind of value is not stated by its source. Two runs bound none of this.

**An assumption with no remedy.** Approximation B2 — that a token’s update does not change the conditional distribution at other positions — is false in a network with shared parameters. It is inherited from the base method rather than introduced here, but our construction adds a second estimator that relies on it.

**One model.** Every result here is measured on one 1.5B mathematics-specialized model. We did not run a second backbone, so nothing in this paper says whether the dissociation, the ordering of the

arms or the ablations carry over to a different scale or pre-training family. The derivation does not depend on the model; the measurements do, and a backbone on which the apparatus contrasts stopped moving the aggregate entropy, or the  $A_H$  contrasts started to, would bear directly on the claim this paper leads with.

## Ethics Statement

This work studies an optimization-time regularizer for reinforcement learning on mathematics problems with automatically verifiable answers. It introduces no new data collection, no human subjects, and no new model release. The intervention acts on which of several equally-scored continuations receives a slightly smaller gradient weight; it does not alter the reward, and so does not change what behavior is being optimized for. The training data is a public competition-mathematics set and the evaluation set is a public examination.

The main foreseeable risk is indirect. Methods that preserve diversity during RL fine-tuning are not specific to mathematics, and the same mechanism would apply to a reward model encoding any objective, including a harmful one; we note that our method inherits, and does not mitigate, whatever the reward function encodes. The computational cost of the method (a forward pass through a set of small prediction heads at every optimization step) is a real environmental cost, which we quantify in Section 12.12 so that it can be weighed against the benefit.

## References

- Anonymous. 2025. Rethinking entropy regularization in large reasoning models. ArXiv preprint arXiv:2509.25133.
- Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwei Peng, Jason D. Lee, Deming Chen, and Tri Dao. 2024. Medusa: Simple LLM inference acceleration framework with multiple decoding heads. ArXiv preprint arXiv:2401.10774.
- Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, et al. 2025. The entropy mechanism of reinforcement learning for reasoning language models. ArXiv preprint arXiv:2505.22617.
- Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. 2018. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International Conference on Machine Learning*, pages 1861–1870. PMLR.
- Zhezhen Hao, Hong Wang, Haoyang Liu, Jian Luo, Jiarui Yu, Hande Dong, Qiang Lin, Can Wang, and Jiawei Chen. 2026. Rethinking entropy interventions in RLVR: An entropy change perspective. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 31105–31133.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, et al. 2024. Olympiad-Bench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. ArXiv preprint arXiv:2402.14008.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the MATH dataset. ArXiv preprint arXiv:2103.03874.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, et al. 2022. Solving quantitative reasoning problems with language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 3843–3857.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. ArXiv preprint arXiv:1707.06347.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. 2024. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. ArXiv preprint arXiv:2402.03300.
- Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. 2025. HybridFlow: A flexible and efficient RLHF framework. In *Proceedings of the Twentieth European Conference on Computer Systems*, pages 1279–1297.
- Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, et al. 2025. Beyond the 80/20 rule: High-entropy minority tokens drive effective reinforcement learning for LLM reasoning. ArXiv preprint arXiv:2506.01939.
- Fang Wu, Weihao Xuan, Ximing Lu, Zaid Harchaoui, and Yejin Choi. 2025. The invisible leash: Why RLVR may not escape its origin. ArXiv preprint arXiv:2507.14843.
- An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, et al. 2024. Qwen2.5-Math technical report: Toward mathematical expert model via self-improvement. ArXiv preprint arXiv:2409.12122.

Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. 2025. DAPO: An open-source LLM reinforcement learning system at scale. ArXiv preprint arXiv:2503.14476.

Brian D Ziebart, Andrew Maas, J Andrew Bagnell, and Anind K Dey. 2008. Maximum entropy inverse reinforcement learning. In *AAAI*, volume 8, pages 1433–1438.

## A Closed Forms for Section 4.4

Perturbing one logit of a  $k$ -way uniform distribution by  $\delta$  and writing  $S := e^\delta + k - 1$  gives  $\pi_1 = e^\delta/S$  and  $\pi_j = 1/S$  for  $j \neq 1$ , so

$$H(\delta) = \log S - \frac{\delta e^\delta}{S},$$

$$H'(\delta) = -\frac{(k-1)\delta e^\delta}{S^2},$$

whence  $H'(0) = 0$  and  $H''(0) = -(k-1)/k^2$ . The sign of  $H'$  is  $-\text{sign}(\delta)$ , so any perturbation lowers the entropy, but only at second order. Table 11 contrasts this with a peaked distribution, where the first-order term does not vanish.

| distribution       | $H_0$ | $\Omega$ | $\Delta H$ at $\delta=.1$ |
|--------------------|-------|----------|---------------------------|
| uniform, $k=10$    | 2.303 | 0        | $-4.7 \times 10^{-4}$     |
| uniform, $k=50$    | 3.912 | 0        | $-1.0 \times 10^{-4}$     |
| top .5 + 9 uniform | 1.792 | -0.549   | $-5.6 \times 10^{-2}$     |
| top .9 + 9 uniform | 0.545 | -0.396   | $-3.8 \times 10^{-2}$     |

Table 11: First-order prediction versus actual entropy change. At a uniform branch point the first-order term is exactly zero and the realized damage is two orders of magnitude below that of a peaked state under the same perturbation.

## B Implementation Notes

**Sibling mask.** For a group of  $G$  rollouts of length  $T$  we form  $\text{eq}[i, j, t] = \mathbf{1}[y_t^{(i)} = y_t^{(j)}]$  and take a cumulative product along  $t$ , so that  $\text{prefixeq}[i, j, t]$  is true exactly when the two rollouts agree on  $y_{\leq t}$ . Shifting it one position right gives agreement on  $y_{< t}$ , which is the sibling relation Eq. (15) requires. Padded positions are treated as disagreement, so a length difference counts as a divergence.

**Alignment.**  $\hat{H}_{\text{togo}}$  is read at position  $t+1$  and moved to position  $t$ ; at the final position the value is kept. Omitting this shift scores the state before the branch instead of the branch taken.

**Memory.** The RL path never materializes head logits. For each head the hidden states are processed in chunks of 4096 rows, mapped through the shared unembedding, reduced to an entropy, and released, so peak memory is  $O(\text{chunk} \times V)$  rather than  $O(\kappa \times B \times T \times V)$ .

**Equivalence test.** A unit test asserts that with  $\lambda = 0$  the token weights are bitwise identical to the base method’s on random inputs; the  $\lambda = 0$  early return in the combination step is what makes this hold exactly rather than to floating-point tolerance.

## C Pre-Registered Analysis

The analysis plan below was committed to the artifact repository before any run of the multi-seed campaign was launched, so that the statistics could not be chosen to fit the outcome. It is reproduced here in full; the file is `docs/preregistration_5seed.md` in the release.

**Unit and primary endpoint.** The unit of analysis is the seed, not the validation step. Each run is reduced to one number — its mean `acc@32` on AIME24 over steps 40–110 — and arms are paired within a seed.

**Contrasts, in the order they are reported.** Primary: STEER-F – GRPO. Secondary, in this order: STEER-F – STEER, STEER-F – UNIFORM, STEER-F – PERMUTED. Secondary endpoints: `maj@32`, uplift, and the six-benchmark mean. Effect sizes are reported with intervals rather than significance stars, and all five arms are reported whatever the outcome.

**Stopping and exclusion.** Every run goes to step 110; none is stopped early for looking good or bad. A run is excluded only for mechanical failure — out of memory, a crash, a corrupted checkpoint, or a step-0 validation score outside .030–.055, which would indicate a different initial checkpoint. Exclusions are listed with their reason. A failed run is relaunched with the same seed, never a new one, and no seed is added after results are seen.

**Checkpoint selection.** Two rules are computed and both reported: the argmax over a run’s validation points, and a fixed step 110. The six-benchmark evaluation uses the argmax. Selecting a maximum over twelve noisy checkpoints is biased upward — in the single-seed data by +.011

to  $+0.017$  per arm — which is why the primary end-point is a window mean rather than a maximum.

**What would falsify the central claim.** Stated in advance so it cannot be renegotiated afterwards. If STEER-F — UNIFORM straddles zero across seeds, the forecast’s *value* does not matter and the apparatus explains the effect; that is the claim of this paper and it would be refuted. If STEER-F — PERMUTED straddles zero, the pairing between a branch and its score does not matter, which the derivation says it must. If aggregate entropy orders the arms the same way accuracy does, the dissociation argument of Section 12.3 fails, since it depends on the two orderings coming apart. And if the mean token weight in the STEER-F arm departs from  $\approx 0.999$ , then  $A_H$  is not sum-zero over sibling sets as derived and the implementation is wrong rather than the theory.

**Amendment, recorded.** The pre-registration assigned three seeds to UNIFORM and PERMUTED and five to the other three arms. The campaign as executed runs all five arms at five seeds. The change raises the power of the two control contrasts and was made before any result of the campaign was seen; it is recorded here rather than silently absorbed.

**Deviations as executed.** Recorded here because the plan above was not carried out in full. (i) Seeds: every arm completed seeds 1–4 except STEER-F at seed 2, which failed at start-up and was not relaunched; the seed-5 runs were started and not completed. The analysis therefore uses the seeds on which all five arms finished, 1, 3, 4, and the other runs are released but enter no statistic. (ii) Machines: seed 1 ran on a two-GPU box and seeds 3 and 4 on a four-GPU box, for every arm (Section 12.8). (iii) The six-benchmark evaluation covers one seed per arm and no STEER checkpoint (Section 12.9). (iv) The long GRPO run for the compute-matched comparison was not completed; Section 12.11 reports what the step costs support instead. (v) The planned second backbones were not run. (vi) The support-restricted split registered below is instrumented but carried by no run reported here. (vii) Follow-up arms: the four advantage-estimator runs (RLOO and OPO, each with STEER-F and STEER) failed at start-up on infrastructure errors and completed no training step; the base method at the released  $\lambda_{\min} = 0.8$  was not launched; the relaxed-clipping STEER control

stopped at step 62. None of them enters Table 9, and their logs are released. The choice to analyse on the common seed set was made after per-seed results were visible; it is the only set on which all five arms exist, and every run outside it is released.

**Second amendment, added 2026-09-21, after this section’s own entropy condition was retired.** Everything above was fixed before the campaign ran. This paragraph was not, and it is marked so rather than folded into the text above. It justifies nothing about the seeds already collected; it binds the seeds not yet in hand.

*What it replaces.* The third falsification bullet above put a condition on the aggregate entropy ordering. Section 12.3 retires it as a category error — it registers the aggregate as a signal about  $A_H$ , which Section 9.3 had already said it cannot be, an hour before this file was committed — and claims nothing from it in either direction. A first attempt at a replacement, registered earlier on 2026-09-21, named `steerf/branch_entropy`; that registration is withdrawn here rather than deleted, because the monitor it named selects a fixed top decile of  $A_H$  rather than the support (Section 12.6).

*Registered now, before the runs that can answer it.* The quantity is `steerf/support_entropy` and its gap, the entropy split on  $A_H \neq 0$  — the same predicate the correction itself applies, so the split covers the positions the method touches and no others. The primary contrast is STEER-F — PERMUTED, paired by seed, because that pair shares  $\lambda$ ,  $\lambda_{\min}$ , the sampler and the magnitudes of  $A_H$ , and differs only in whether a branch is paired with its own score; STEER-F — UNIFORM is secondary. The claim under test is that a support-restricted entropy contrast near zero can accompany an accuracy contrast away from zero.

*What would falsify it.* If the support-restricted contrast tracks the accuracy contrast in sign and magnitude across seeds the way the apparatus contrasts in Table 6 do, then the split is the aggregate statistic in a smaller window and the claim fails.

*What is still not instrumented, and we say so rather than substitute.* Section 9.3’s own prediction is that the spread *across siblings* widens. Testing it needs the sibling index at the loss site, where only the per-token tensors are currently in scope, so no quantity we log tests it today. We do not offer the support-restricted split as that test: it measures a level, not a spread.

## D Within-Run Stability

Earlier drafts of this work paired the arms at each of the eight validation steps in the converged window and reported a paired  $t$  over those eight points. That statistic answers a narrower question than the one in Section 12.4: it asks whether a gap is stable along a single trajectory, not whether it reappears when the run is repeated. Its  $n$  counts validation points, which are serially correlated within a run, so it is not an estimate of run-to-run uncertainty and cannot be read as one.

We report it here because it is informative about a different thing — whether a gap is an artifact of one lucky validation point — and because the single-seed version of this manuscript quoted it. Where the two disagree, the seed-level statistic in Section 12.8 governs.

The contrast between the two is the point. Table 12 separates every  $A_H$  comparison at  $t_7$  between +2.56 and +7.18, which reads as decisive and is not: the same gap sits at the same order as the run-to-run spread we measure in Section 12.8. A statistic can be this confident about a trajectory and say nothing about a re-run, which is exactly why we moved it out of the main text.

| contrast           | $\Delta\text{acc@32}$ | paired $t_7$ |
|--------------------|-----------------------|--------------|
| STEER-F – GRPO     | +0.0145               | +3.08        |
| STEER-F – STEER    | +0.0165               | +6.68        |
| STEER-F – UNIFORM  | +0.0182               | +7.18        |
| STEER-F – PERMUTED | +0.0125               | +2.56        |

Table 12: Within-run stability over the eight converged-window validation points of seed 1. The points are consecutive validations of one run and are therefore serially correlated, so  $\text{df} = 7$  is generous and these  $t$  values are not an estimate of run-to-run variation; see Section 12.8 for that. Majority-vote versions of the same four rows separate at +5.09 to +5.50.